

Capability or Optimizer? What Lets an LLM Proposer Leave Its Template Family on Circle Packing

Anonymous authors

Paper under double-blind review

Abstract

Zero-shot LLM proposers on circle packing concentrate on a grid-plus-filler family whose best value follows in closed form from N . We cross two proposer tiers with three output channels and ask what lets a proposer leave that family. A Sonnet-tier program with `numpy` and `scipy` beats the family in 23 of 25 valid cases; without them, on the same path and budget, it clears 0 of 35; a weak-tier program with them clears 0 of 19, and 4 of 41 under a lenient reparse. A fixed SLSQP program with no model in the loop clears 45 of 45 through the same pipeline and beats the best Sonnet program at every cell: the optimizer clears the family, and the tiers differ in whether their programs drive it there. Handed the better construction, the weak tier reaches for it in 36 of 45 invocations and builds it in 6 of 14: execution, not preference, is the bottleneck. The zero-shot characterization survives neither a conditioning step, nor five loop generations, nor a change of serving path, where the reasoning budget is what restores it. Twenty-six registered arms are reported, whichever way each went.

1 Introduction

Ask a language model to place 21 circles in a unit square so as to maximize the sum of their radii, and it will not search. It reaches for the nearest square grid — five by five, radius $1/(2k)$ — then truncates it to 21 circles, leaving four cells empty and their area unclaimed. A better construction is one parameter away: drop to a four-by-four grid and add interior-vertex fillers of radius $(\sqrt{2} - 1)/(2k)$, one in each gap between four grid circles, and the sum rises. The model does neither, and more often than not the next sample repeats the same move. The family of grid-plus-filler constructions has a best expressible value at every N , computable in closed form (§2.2); the question this paper asks is what it takes for a proposer to leave that family.

Circle packing is the showcase benchmark of the LLM-driven discovery literature, the lineage FunSearch (Romera-Paredes et al., 2024) opened. AlphaEvolve (Novikov et al., 2025) reported 2.63586276 at $n = 26$, and later systems report a narrow band — ShinkaEvolve (Lange et al., 2025) 2.6359831 (2.6359777 strict), HELIX 2.63598308 (Su et al., 2026), GigaEvo 2.63598 (Khrulkov et al., 2025), AdaEvolve 2.636 (Cemri et al., 2026) — with ThetaEvolve (Wang et al., 2025) claiming new best-known bounds. Every one of those systems pairs a capable model with an executed optimizer, and two critiques from outside ask what each factor contributes: Gideon et al. (2026) show simple baselines recover much of the reported advantage, and Berthold et al. (2026) show classical solvers do too. We ask the same question from inside, by varying the two factors separately.

The result is a table. Two proposer tiers (a weak tier and the Sonnet tier of one vendor) cross three output channels (direct coordinate emission; a program restricted to the standard-library `math` module; a program allowed `numpy` and `scipy`), every cell scored under one clearance rule against the family’s best expressible value (Table 2). One cell clears at rate: a Sonnet-tier program with the libraries clears the family in 23 of 25 valid cases, at every cell, median margin 3.6%, best per cell 2.5%, 3.6% and 4.2%. A weak-tier program with the same libraries clears it in 0 of 19 valid over two registered arms (90 invocations, Wilson upper bound 16.8%, under the registered 20% bar), and in 4 of 41 under a post-hoc reparse of the malformed outputs, which print `numpy` scalars; the claim that survives either reading is the gap to the Sonnet cell, not the zero (§3.5). A Sonnet-tier program without them — simulated annealing, multi-restarts — clears it in 0 of 37 on the agent runtime and 0 of 35 on the library arm’s own path and budget, so the gap between the two Sonnet

program cells is the library (arm CCP, §3.5). Every weak-tier cell is zero at the trap cells under the registered scoring. Two cells show escapes at low rate, both out-of-family constructions rather than in-family search: one Sonnet direct emission in 30, and two weak-tier emissions in 15 at one extend cell. A fixed random-restart SLSQP program with no model in the loop, committed before its run and executed through the library arm’s own pipeline, clears 45 of 45 and its best packing beats the best Sonnet program at every cell (arm B, §3.5). On this benchmark, at these cells, the optimizer clears the family: the Sonnet tier’s programs clear because they invoke it, and land below the fixed reference; the weak tier’s programs invoke it and do not clear.

Why: execution, not preference. Handed the provably better in-family construction with its score, the weak tier reaches for it in 36 of 45 invocations and builds it correctly in 6 of 14 valid outputs (arm CH, §4.1). Delegated to a math-only program it reaches the argmax construction exactly twice in 78 valid programs and never passes it (§3.4). A second weak-tier family emits the argmax construction at 27 of 43 valid outputs at the cells where template and argmax differ, where the original emits it in 2 of 57 (§5.4), and at a held-out cell the original selects the right template and mis-places its one filler (§4.2). The template is not what the model prefers; it is what the model builds reliably, and the library optimizer is what supplies the execution the weak tier lacks.

And a methods result the table needed. The weak tier’s template is a closed form: $k^*(N) = \text{round}(\sqrt{N})$ names the grid, and its value is the empirical modal output at all seven tested N and at four of five held-out N (Appendix E). That characterization was built from zero-shot calls, and it does not survive a change of setting. One parent-conditioning step dissolves it (arm MU, §5.1); five generations of a registered loop do not restore it (arm L, §5.2); the same prompt sent to the same vendor’s pinned API yields 0 valid packings in 105 where the agent runtime yields the template in 4 of 6, and turning the vendor’s extended thinking on restores them (arms P and P-D, §5.3). A study that characterizes an LLM proposer zero-shot and reasons about its behaviour in a loop owes that step a measurement; here it cost one arm each time and the characterization survived none of them.

Contributions.

1. *A tier $\times$ channel table, and the reference program that reads it.* Across two tiers, three channels and two serving paths, only a Sonnet-tier program with `numpy/scipy` clears the family’s best expressible value at rate (23 of 25, 3 of 3 cells); the same-path, same-budget isolating arm attributes the gap to the library, and a fixed reference program with no model in the loop clears 45 of 45 through the same pipeline and beats the best Sonnet program at every cell. The library optimizer clears the family; what the tiers differ in is whether their programs drive it there: the weak tier’s programs with the same libraries clear 0 of 19 under the registered parser and 4 of 41 under a lenient one, against 23 of 25. Every weak-tier cell is 0 at the trap cells under the registered scoring, which gives a discovery claim from those proposers a computable value to beat.
2. *The bottleneck is execution.* Handed the argmax, the weak tier reaches for it in 36 of 45 invocations and builds it in 6 of 14 valid outputs; the same asymmetry appears in code (argmax reached, never passed), across a vendor (the family that can execute the harder construction emits it) and at a held-out cell (right template, wrong filler).
3. *Zero-shot characterization does not transfer.* Three changes of setting each carried a registered prediction that the anchor persists, and that prediction failed in each: under conditioning the falsifier fired; under iteration two of four predictions were not met; on the pinned path every prediction was unscorable because validity collapsed, and the component that restores it is the reasoning budget.

Scope. Two tiers of one vendor (a third alias reached 13% validity and is reported in Appendix C); three trap cells ($N = 13, 21, 31$) carry every channel, and the direct weak-tier row adds $N = 43$. One clearance rule, §2.4: valid at 10^{-9} and above the family argmax by more than 10^{-6} ; nothing else carries a clearance verdict. Agent-runtime rows and pinned-API rows are never pooled; the Sonnet contrast is same-instrument and same-budget since arm CCP, while the weak tier’s math-only cells ran at 10 seconds and its library cells at 120, a difference §7 discloses. The direct rows are unconditioned calls, and the anchor result is scoped to them (§5.1).

What is not claimed. Nothing about mechanism inside the weights: the model *emits*, the distribution *concentrates*, the closed form *identifies the modal output*, and §4.1 reports the direct probe of the leading mechanism-free alternative. Nothing about what any named system does at generation zero, and nothing about whether its diversity machinery (Mouret & Clune, 2015; Lehman & Stanley, 2011) repairs or explores. Preregistration is an adopted standard (Thomas et al., 2026; Williams et al., 2026; Vaccaro, 2026; Zhang, 2026), not a contribution; behavioral anchoring is established territory (Huang et al., 2026; Lou & Sun, 2024; Malberg et al., 2025), and the anchor here is a construction template, not a number.

2 The family and its argmax

2.1 The benchmark

Maximize the sum of radii of N non-overlapping circles in the unit square. Feasibility is decidable exactly from emitted coordinates, so proposals score with no human in the loop. The direct-emission prompt (Appendix A) fixes the count, states containment and non-overlap, forbids writing or executing code, and demands a raw Python list of `[x, y, r]` triples. The two program channels replace the code-free clause with a program contract — standard-library `math` only, or `math` with `numpy` and `scipy` — and score the program’s printed output by the same conventions (§3.4, §3.5).

2.2 The recipe family

Three terms, used consistently: the **recipe family** is the parametric set of grid-plus-filler constructions; a **template** is one member; the **anchor** is the template the distribution concentrates on at a given N .

A $k \times k$ grid places one circle per cell center with $r_{\text{grid}} = 1/(2k)$; m fillers sit on interior grid vertices, tangent to their four surrounding grid circles, with $r_{\text{filler}} = (\sqrt{2} - 1)/(2k)$, $0 \leq m \leq (k - 1)^2$. Hence $V(k, m) = k/2 + m(\sqrt{2} - 1)/(2k)$. When $N < k^2$ the observed behavior is not to drop to a smaller grid and fill it but to *truncate* — lay out the $k \times k$ lattice, occupy N cells, leave the radius unchanged — giving $T(k, N) = N/(2k)$. These reproduce every previously reported anchor with no fitting: $V(5, 1) = 2.5414214$, $V(5, 2) = 2.5828427$, $V(5, 0) = 2.500$, $T(5, 23) = 2.300$, $V(4, 7) = 2.3624369$.

The **family argmax** at N is the largest value the family can express there, recomputed in closed form over every admissible (k, m) ; inside a trap zone it is $V(k^* - 1, m)$, the drop-and-fill construction the model does not make. Every closed-form value in the paper is recomputed over the constructed coordinates by an independent linear program that knows nothing about the recipe and aborts on disagreement: 83 configurations at $k = 2 \dots 7$ with drift below 10^{-9} , and a further 130 at $k = 8$ and 9 with worst drift 1.8×10^{-14} (Appendix E.1). Across the 155 weak-tier rows collected before arm M only two exceeded the family best, by 3.0×10^{-7} and 3.5×10^{-9} , both below the 10^{-6} bar of §2.4.

2.3 The order rule is a definition; the branch rule is the hypothesis

We write $k^*(N) = \text{round}(\sqrt{N})$ for the nearest-square order. This is a *definition*: given some $k \times k$ grid from this family, “the nearest square lattice” yields $\text{round}(\sqrt{N})$ by construction, and two formalizations of “nearest” agree on the integers. The falsifiable content is the **branch rule**: $k^{*2} \leq N \rightarrow \text{extend}$ with $m = N - k^{*2}$ fillers (converge); $k^{*2} > N \rightarrow \text{truncate the } k^*\text{-grid (trap)}$. The extend arm survived testing only at $m \leq 1$; a preregistered extension disconfirmed it at $m \geq 4$, where the model prefers the $(k^* + 1)$ rectangle grid (Appendix E.4). Substituting k^* into the branch condition gives the *trap zones* $N \in [k^2 - k + 1, k^2 - 1]$: [13, 15], [21, 24], [31, 35], [43, 48], [57, 63].

A cell is *discriminating* when the prediction differs from the family argmax. At a converge cell the extend prediction *is* the argmax — we exhausted $N = 10 \dots 80$ and found no exception (`extend_branch_scope.py`: 38 converge cells, none discriminating; 27 of 33 trap cells discriminating) — so every discriminating cell in this paper is a trap cell, and the in-family predictive content that survives is the truncate arm $T(k^*, N) = N/(2k^*)$. Nothing about “nearest” forces truncation: a proposer with any local search would drop to $k^* - 1$ and fill, which is why the trap zones exist and why a higher-valued rival is reachable inside them. Truncate-versus-

drop-and-fill is the dichotomy every arm tests. The governing quantity is the signed distance $N - k^{*2}$; primality is not the variable.

2.4 Notation and conventions

Terms used throughout, gathered here so no verdict table depends on a definition made elsewhere:

Term	Meaning
$k^*(N)$	nearest-square grid order, $\text{round}(\sqrt{N})$ (a definition, §2.3)
$V(k, m), T(k, N)$	extend-with- m -fillers value $k/2 + m(\sqrt{2} - 1)/(2k)$; truncated-grid value $N/(2k)$
anchor	the template the proposal distribution concentrates on at a given N
on-prediction	emitted sum within 2×10^{-3} of the registered predicted value
rival (argmax)	the highest-valued member of the recipe family at that N when it differs from the prediction — inside a trap zone, $V(k^* - 1, m)$
discriminating cell	N where prediction $\neq$ family argmax; only these can distinguish anchoring from family search
validity	non-overlap and containment at 10^{-6} (primary) and 10^{-9} (logged), both always reported
trap / converge	branch by sign of $N - k^{*2}$: truncate when $k^{*2} > N$, extend when $k^{*2} \leq N$
tier	nominal capability class of the serving alias addressed (weak/mid/top); opus_alias is a bare alias with no attestable weights binding
instrument	the serving path: the agent runtime (decoding parameters not exposed, conditioning context locked only in part) or the vendor’s pinned chat-completions API (temperature 1.0, no system prompt); never pooled
UNSCOREABLE	cell with fewer valid samples than its arm’s registered floor; claims nothing
evaluability floor	the minimum valid-sample count a cell needs before its verdict counts, fixed per arm at that arm’s registration and <i>not</i> common across arms: 3 for the GM chain, 5 for arms CN, CP, P, CL, CL-W, CCP and P-D, none for arm CH. Where the loosest one carries a verdict we report the stricter reading beside it
family argmax	the highest value the recipe family can express at N , recomputed in closed form; every clearance verdict is measured against it, and it is the number a discovery claim from a measured proposer has to beat. Table 2 says which proposers clear it
clears at rate	a cell clears <i>at rate</i> when its clearance count meets the arm’s registered rate bar (20% of valid outputs for arms CL and CCS) at every scoreable cell; a cell with one or two clearances out of dozens shows an <i>escape</i> , counted and named, not a clearance at rate

The clearance rule, stated once. Five thresholds appear in this paper — 10^{-6} and 10^{-9} for validity, 2×10^{-3} for value matching, 10^{-6} again as a minimum excess, and 5×10^{-8} as the granularity at which sums are recorded — and a reader could not otherwise tell which governs a claim that an output *left the recipe family*. One rule governs every such claim in this paper:

An output leaves the family iff it is valid at 10^{-9} *and* its sum exceeds the family argmax by more than 10^{-6} .

Both conditions are load-bearing and each excludes a real set of rows. The validity condition excludes arm MU’s eighteen rows that sit above the argmax at the primary tolerance by at most 1.5×10^{-6} and fail the tighter gate — their excess is their overlap (§5.1). The excess condition excludes two square-arm rows above the family best by 3.0×10^{-7} and 3.5×10^{-9} (§2.2), which are at the argmax as recorded. Applied across the arms swept for it before arm CL, exactly three rows leave the family: two converge-cell rows in arm M (Appendix E.4) and one Sonnet sample at $N = 31$ (§3.3), all three out-of-family constructions rather than in-family rivals; arm CL’s Sonnet tier adds twenty-three more, all library-optimized programs, and a post-hoc lenient reading of the weak tier’s library programs four more that the registered parser does not count (§3.5). The 2×10^{-3} value window never carries a clearance verdict; §3.3 records the one case where it was mistakenly allowed to.

3 Main result: who clears the family argmax

Twenty-six registered arms appear in this paper. Each was registered before its sampling, each is reported regardless of outcome, and the falsifiers that fired are flagged in place (Table 1). Ten arms carry the table (F, M, S, CC, CC2, CCS, P, CL, CCP, CL-W); the rest establish the mechanism (§4), the transfer boundaries (§5) or the mode prediction (Appendix E); arms T, G and O carry no verdict the table needs and are reported in Appendix C and in the supplement. Arms whose verdicts rest on non-discriminating cells, on fewer valid samples than their floor, or on a null that does not separate are reported in full in the supplement that accompanies this submission, with a row here and a pointer (the criterion is stated in §8).

Table 1: The twenty-six registered arms: section (a letter is an appendix; S_n is a section of the supplement), proposer, design, registered test, outcome.

Arm	§	Proposer(s)	Design	Registered test	Outcome
F (square, bare)	3.3, E	weak tier	7 $N \times 5$ -20	P1–P5 exact-value modes	modal at 7/7 N
S (Sonnet tier)	3.3	Sonnet tier	3 $N \times 10$	P-S1–P-S4 tier contrasts	1/30 on-prediction, 100%/90% valid ($10^{-6}/10^{-9}$); one out-of-family escape; 5 of 20 samples at $N = 13/21$ seen before registration, disclosed 13% valid; 3 predictions not evaluable (deviation, Table 3)
O (top alias)	C, S3	opus_alias	3 $N \times 10$	P-O1–P-O4	falsifier triggered (tie reading); P-T3 fragile
T (trace)	S6	weak tier	3 $N \times 20 \times 2$ arms	P-T1–P-T4 + falsifier	partial, 5/11; formula negative result kept
G (rectangle)	S5	16 invocations	2 cells	out-of-sample transfer	F-M1 triggered 3/3 : extend branch dies at $m \geq 4$; fifth k confirms; two out-of-family escapes at $N = 20$
M (branch stress)	3.3, E.4, S4	weak tier	4 $N \times 15$	P-M1–P-M4 + F-M1/F-M2	F-MU1 triggered : anchor dissolves under conditioning
MU (mutation)	5.1	weak tier	3 parents $\times$ 3 $N \times 15$	P-MU1–P-MU4 + F-MU1	P-CH1 holds 2/3, P-CH2 1/3 (registered); post-hoc attempt reading: 36/45 reach for the argmax, 30/30 evaluable invalid rows at its order
CH (choice)	4.1	weak tier	3 $N \times 15$	P-CH1 vs P-CH2	UNDERPOWERED (20 valid / 140)
GM (gemini-flash-lite)	7	2nd vendor	7 $N \times 20$	GM-chain bars	0/140 parseable — budget confound, diagnosed by GM3
GM2 (gemma, 4k budget)	7	2nd vendor	7 $N \times 20$	GM-chain bars	pooled bar met (57.5%) but discriminating cells 11/43 = 26%, below it; misses are the rival construction
GM3 (gemma, 16k budget)	5.4, S2	gemma-4-26b	7 $N \times 20$	P-GM3.1–.3 + falsifier	2 vendors TRANSFER, gemma-4-31b DOES-NOT-TRANSFER, V2 HOLDS
V (cross-vendor zero-shot)	5.4, S2	13 free-tier aliases	5 $N \times 5$	floor + V1/V2	P-CC1 holds: anchor stays modal; 0/37 above the family argmax
CC (code channel)	3.4	weak tier	3 $N \times 15$	P-CC1 vs P-CC2	REPLICATED : 0/41 above argmax, anchor modal 3/3
CC2 (fresh-draw replication)	3.4	weak tier	3 $N \times 15$	R1 + R2 + F-CC2.1	P-CCS2 holds: 0/37 above argmax — the channel, not the tier
CCS (tier probe)	3.4	Sonnet tier	3 $N \times 15$	P-CCS1 vs P-CCS2	

Arm	§	Proposer(s)	Design	Registered test	Outcome
CN (held-out N)	4.2, E.5	weak tier	5 $N \times 15$	P-CN1 vs P-CN2 + S-CN1-3	P-CN1 holds: mode at 4/5 held-out cells, 3/3 discriminating; structure 5/5
CP (perturbed container)	E.6	weak tier	5 $N \times 15$	P-CP1 vs P-CP2 + S/F	P-CP1 holds: mapped prediction modal 4/5; S-CP1 not met (4/61 rival)
RP (direct re-call)	E.6	weak tier	6 $N \times 15$	P-RP1 vs P-RP2 + F-RP1	P-RP1 holds: 0/44 recalls; 87/87 scored responses UNKNOWN
PP (paraphrase)	E.6	weak tier	3 rewrites $\times$ 2 $N \times 15$	P-PP1 vs P-PP2 + S-PP1/2	P-PP1 holds 6/6: modal under every paraphrase; S-PP1 not met (6/75 rival)
L (iterated loop)	5.2	weak tier	2 $N \times$ 2 archives $\times$ 6 rounds $\times$ 5	L1-L4 + F-L1	L2 holds; L1 and L3 not met (L3 reversed); L4 open: 0/49 conditioned outputs clear the family
P (pinned path)	5.3	weak tier, bare API	15 cells $\times$ 15	P-P1-P-P4 + F-P1	validity collapses: 0/105 valid at the square cells, 4/75 held-out; every direct-emission prediction UNSCOREABLE; P-P3 holds (0 of 12 valid programs clear); same-day agent-runtime control 6/6 valid, anchor 4/6
CL (library code)	3.5	weak Sonnet, numpy/scipy	+ 2 tiers $\times$ 3 $N \times 15$	P-CL1 vs P-CL2 + F-CL1	weak: 0/11 clear, P-CL1 holds; Sonnet: 23/25 clear at 3/3 cells , P-CL2 holds, F-CL1 fires
CCP (isolating arm)	3.5	Sonnet tier, bare API, math only, 120 s	3 $N \times 15$	P-CCP1 vs P-CCP2 + F-CCP1	P-CCP1 holds: 0/35 clear, none reaches the argmax; best per cell 1.625, 2.1, 2.673 — the library, not the path or the budget
CL-W (weak top-up)	3.5	weak tier, numpy/scipy	3 $N \times$ 15, pooled with CL	P-CLW1 vs P-CLW2 + F-CLW1; lenient reading secondary	P-CLW1 holds: 0/19 pooled clear, upper bound 16.8%; lenient reading 4/41, dead zone
B (optimizer alone)	3.5, S7	no model in the loop; a fixed SLSQP program	3 $N \times$ 15 seeds, 120 s	P-B1 vs P-B2 + S-B1	P-B1 holds: 45/45 clear , 3/3 cells; S-B1: beats the Sonnet best at 3/3 cells; amendment 1 (version 1 blocked by the allowlist, kept)
P-D (runtime component)	5.3	weak tier, bare API, $N = 13$	2 conditions $\times$ 15	P-PD1 vs P-PD2 vs P-PD3 + S-PD1	P-PD1 holds (closed at 27/30 rows by amendment, verdict unchangeable): thinking on 12/14 valid, template modal; system line only 0/13 — the reasoning budget is the component

3.1 Registration protocol

For each cell the prompt is pinned and its SHA-256 digest written to disk first — the $N = 13$ square prompt hashes to **32db485b**. . . Predictions are registered in the header of each arm’s analysis script or in a standalone registration file, with competing predictions, numeric thresholds, tie conventions and a falsifier fixed before the first invocation. Raw outputs are stored verbatim, failures included; scoring is deterministic

and local. Which of the square arm’s seven cells carry an individually registered point prediction, and which follow from the same closed form without one, is recorded in Appendix E.2.

Four properties are disclosed, not repaired. Temperature, top-p and top-k are not exposed by the agent runtime. The alias-to-weights binding is a promise, not a hash. The subagent inherits instruction files outside the task prompt, so the digests lock a *fragment* of the conditioning context — a replicator cannot reconstruct the inherited files from what we release. Hash coverage is incomplete (§8). Arms P and CL exist because of the first and third of these: they send registered prompts through a serving path that exposes and pins temperature and carries no inherited context (§5.3).

Two scoring conventions were fixed in advance. Validity is reported at 10^{-9} and 10^{-6} , both logged, 10^{-6} primary: proposers print six to eight decimals, so an eight-decimal tangency misses contact by $\sim 5 \times 10^{-9}$, while 10^{-6} sits far below the $\sim 10^{-2}$ gap between rival constructions. Value matching uses a 2×10^{-3} window, which never carries a clearance verdict (§2.4). Every on-prediction rate is computed over *valid* outputs and so conditions on execution success — the variable §4.1 shows can separate what a model attempts from what it lands; per-cell validity denominators are reported beside every rate so the reader can bound the exposure.

3.2 The table

Table 2 is the paper. Each row is one arm, or one registered pair, scored under the clearance rule of §2.4; “clear” counts valid outputs whose sum exceeds the family argmax at that N by more than 10^{-6} while valid at 10^{-9} . Five readings.

One cell clears at rate. A Sonnet-tier program allowed `numpy` and `scipy` clears the family in 23 of 25 valid cases, at 3 of 3 cells, its best program per cell by 2.5%, 3.6% and 4.2% of the argmax and its median clearance by 3.6% (§3.5). Against the published best-known sums (Appendix E.1), the best programs sit 0.5%, 0.9% and 0.8% below at $N = 13, 21$ and 31 : above the grid family, not at the frontier. A fixed reference program with no model in the loop, run through the same pipeline, clears 45 of 45 and beats the best Sonnet program at every cell (arm B, §3.5): the clearing is the optimizer’s, and the row’s programs clear because they invoke it.

Neither tier without the optimizer; the optimizer without any tier. A weak-tier program with the same libraries clears 0 of 11 in arm CL and, pooled with its registered top-up (arm CL-W), 0 of 19 under the registered parser, a Wilson upper bound of 16.8% that excludes the registered 20% bar, though its $N = 21$ cell holds only 3 valid outputs and stays under arm CL’s five-valid floor; under the post-hoc lenient reparse it is 4 of 41, a point estimate of 9.8% whose upper bound of 22.5% does not exclude the bar. The weak tier clears at rate under neither reading, and the exclusion belongs to the registered one alone. A Sonnet-tier program without the libraries — simulated annealing over hand-rolled generators, Halton multi-restarts, force-directed relaxation — clears 0 of 37 (§3.4), and on the library row’s own serving path and 120-second budget clears 0 of 35 without reaching the argmax (arm CCP, §3.5). Every weak-tier cell is 0 at the trap cells under the registered scoring, and the reference program that clears at every cell calls the same library the weak tier’s 19 valid programs call.

Escapes exist and are counted. Two cells show low-rate escapes, both out-of-family constructions found by sampling rather than in-family search: one Sonnet direct emission in 30, at $N = 31$, and two weak-tier emissions in 15 at the extend cell $N = 20$ (§3.3). They are in the table because a zero that hides them would be false; they do not change which cell clears at rate.

The instrument column is not decoration. Arm P (§5.3) shows the same weak-tier prompt yields 78% valid packings through the agent runtime and 0 of 105 through the pinned API, so rows are never pooled across the two paths. The clearing cell ran on the pinned path; the Sonnet math-only cell ran on the runtime, and under a 10-second execution budget against the library arm’s 120 seconds. Within the weak tier the library comparison is same-instrument (0 of 12 math-only and 0 of 11 with libraries, both pinned; each has one cell above the five-valid floor, $N = 21$ for the math-only arm and $N = 31$ for the library arm, 0 of 6 at both; with the top-up 0 of 19 pooled, and a post-hoc reparse reads 4 of 41, §3.5) but still crosses the budget (10 against 120 seconds). Within the Sonnet tier the isolating arm has run: arm CC’s math-only prompt on the pinned

Tier	Channel	Instrument	Cells	Clear	Note
weak	direct emission	agent runtime	13, 21, 31, 43	0 of 57	trap cells; at the extend cell $N = 20$ (arm M) 2 of 15 clear, both out-of-family
weak	direct emission	pinned API	7 square cells	0 valid of 105	unscorable (arm P)
Sonnet	direct emission	agent runtime	13, 21, 31	1 of 30	one out-of-family packing at $N = 31$; 27 of 30 valid at 10^{-9} , the clearance among them (arm S)
weak	math-only program	agent runtime	13, 21, 31	0 of 78	reaches the argmax exactly 2 times, never passes it (CC, CC2)
weak	math-only program	pinned API	13, 21, 31	0 of 12	two cells under the five-valid floor; 0 of 6 at the scoreable cell (arm P code cells)
Sonnet	math-only program	agent runtime	13, 21, 31	0 of 37	annealing and multi-restart search, still 0 (CCS)
Sonnet	math-only program	pinned API	13, 21, 31	0 of 35	the library row's path and 120 s budget without the libraries; best per cell 1.625, 2.1, 2.673, none at the argmax (CCP)
weak	numpy/scipy program	pinned API	13, 21, 31	0 of 11; 0 of 19 pooled	arm CL, then pooled with its registered top-up CL-W (90 invocations; the $N = 21$ cell stays under the five-valid floor): Wilson upper bound 16.8%, under the 20% bar; a post-hoc reparse of the <code>numpy</code> -scalar stdouts reads 4 of 41 (CL, CL-W)
Sonnet	numpy/scipy program	pinned API	13, 21, 31	23 of 25	3 of 3 cells; excess over the argmax 0.67–4.24%, median 3.62% (CL)
none	numpy/scipy program, fixed	local, 120 s	13, 21, 31	45 of 45	a fixed random-restart SLSQP reference program, no model, through the library row's pipeline; 3 of 3 cells at the 20% bar; best per cell exceeds the Sonnet best at 3 of 3 cells (arm B)

Table 2: **Who clears the family argmax.** Each row is one arm or one registered pair of arms, scored under §2.4's clearance rule: valid at 10^{-9} and above the family argmax by more than 10^{-6} , the argmax recomputed in closed form. "Clear" counts valid outputs. Instrument names the serving path, which arm P shows is not interchangeable: the same weak-tier prompt yields 78% valid packings through the agent runtime and 0 of 105 through the pinned API. Rows are not pooled across instruments.

path under the library arm's 120 seconds clears 0 of 35 and never reaches the argmax (arm CCP, §3.5), so the 0 of 37 to 23 of 25 gap is attributed to the library, on one path and one budget.

Outside the table. Three arms sample regimes no row covers — five held-out cells absent from every scoreboard (arm CN), one parent-conditioning step (arm MU) and five generations of a registered loop (arm L) — and none of their valid outputs clears the family argmax either. Pooled with the three math-only program rows, 0 of 290 valid outputs across six arms clear, a pool that counts arm MU at 10^{-9} and the other five at the primary tolerance; held at 10^{-9} throughout it reads 0 of 281, and held at 10^{-6} throughout 315 with 18 above the argmax, every one of the 18 an arm-MU overlap the tighter gate rejects. The pool, its membership rule and its tolerance accounting are Supplement S1. The figure describes those arms and is not a bound: the escapes above sit outside it by named regime.




Figure 1: *Three attractor families at a single cell, $N = 31$. One representative packing per tier, ledger sample id in each panel title (weak tier **bare** id 1, on-prediction truncation 2.5833333; Sonnet tier **sonnet_bare** id 1, mixed-radius rival 2.7485281; opus_alias id 1, invalid, with only the offending circles highlighted). Source: arm_f_candidates_v2.jsonl via fig_scripts.py, deterministic.*

3.3 Direct emission

The weak tier emits a computable template. The weak-tier proposer produces uniform grids of radius $1/(2k)$ truncated to N circles, with interior-vertex fillers only when the grid underfills. Its modal output is the closed form’s prediction at all seven tested N , and at the four discriminating cells — $N = 13, 21, 31, 43$, where prediction and family argmax differ — 41 of 57 valid emissions land on the prediction and 2 reach the rival argmax, both at $N = 21$. No valid emission at those cells clears the family. The per-cell table, the template-shape null it survives, the held-out replication at five N absent from every scoreboard, and the perturbed-container and paraphrase probes are Appendix E.

Two weak-tier emissions do clear the family, and they are the reason the direct row carries a second entry. At the extend cell $N = 20$ (arm M, Appendix E.4) the registered prediction $V(4, 4) = 2.2071068$ was disconfirmed — the model emits the 5×4 rectangle, $T(5, 20) = 2.0$, in 12 of 15 — and two of the fifteen valid rows are staggered rows of four at $r = 1/9$, sum 2.2222222, a hexagonal layout the family does not contain, valid at 10^{-9} . They are not the modal output, they were found by a post-hoc sweep disclosed as such, and they are counted.

The Sonnet tier perturbs and mixes. The Sonnet-tier arm was preregistered in arm_s_preregistration.txt, disclosing that 5 of 20 samples at $N = 13/21$ had been seen before registration and that $N = 31$ was fully blind. It was valid 30/30 at 10^{-6} and 27/30 at 10^{-9} , and almost never on-prediction: 0/10, 0/10, 1/10, against the weak tier’s 12/16 at the same cells. Instead of truncating a uniform grid it perturbs one, with enlarged edge rows, hexagonal interior rows and two or three distinct radii in the same packing (29/30 multi-radius; Figure 1). One sample at $N = 31$ emitted 27 circles at $r = 1/12$ with 4 corner circles at $r = 1/8$, summing to 2.7499999991, above the family’s 2.7485281; recomputed from its coordinates the construction is tangency-tight, with minimum pairwise and wall slack exactly zero, and it is among the 27 samples valid at 10^{-9} . It is the direct-emission row’s one clearance: 1 of 30 at the primary tolerance, 1 of 27 at 10^{-9} .

That sample also exposed a scoring defect, and it is why §2.4 states the clearance rule once. At $N = 31$ the family rival (2.7485281) and this out-of-family value (2.75) differ by 1.47×10^{-3} , inside the registered 2×10^{-3} window, so the value classifier counted the escape as a hit on the construction it exceeded. The window was chosen for the anchor comparison, where the nearest competing value is ~ 0.15 away. Classified by structure instead, Sonnet’s rival count at $N = 31$ is 2/10 and pooled 5/30. A third alias, addressed through a bare tier name we call opus_alias, was valid in 4 of 30 and is reported in Appendix C with the ladder it completes.

3.4 Math-only programs: an executed program does not escape the family

The code-free design rested on a rationale the first draft asserted and did not measure: that a code channel routes construction to an executed optimizer, so the emission distribution would reflect the optimizer rather than the model. Deployed discovery loops elicit *programs*, so that gap was the standing objection to loop relevance. Arm CC measures it. The design was registered and pushed before sampling: the three trap cells $N = 13, 21, 31$, $n = 15$ weak-tier invocations per cell, prompt equal to the bare stem with exactly two registered line replacements — the code-free clause becomes a program contract (standard-library `math` only, stated in the prompt) and the output clause asks for program source. Programs run sandboxed (`python -I`, 10-second timeout, registered AST allowlist; the executor and its four-bin smoke test were committed with the registration), and stdout is scored by arm-F conventions unchanged. Two competing predictions: P-CC1, the modal valid executed output remains the $T(k^*, N)$ bucket at ≥ 2 of 3 cells; P-CC2, it sits strictly above the anchor at ≥ 2 of 3; ties or mixed modes count as neither.

N	Valid	Modal bucket (count, margin)	Anchor-rate	Above anchor
13	12/15	anchor 1.6250000 (2/12, one sample)	2/12	2 — 1.6614054, and the family argmax exactly
21	13/15	anchor 2.1000000 (9/13, margin 8)	9/13	0
31	12/15	anchor 2.5833333 (3/12, one sample)	3/12	0

P-CC1 holds as registered (3 of 3 cells), and its margins belong next to it. At $N = 21$ the mode is heavy — nine exact 2.1000000 emissions, most programs writing the 5×5 grid truncated to 21 circles — but at $N = 13$ and $N = 31$ the modal bucket is the anchor value with a **one-sample margin** over dispersed sub-anchor values, so the modal-identity claim is thin at two of three cells. The robust result is the registered decomposition row: **no valid program output exceeds the family argmax anywhere** (0 of 37 above the rival), and only two exceed the anchor at all, both at $N = 13$ — one at 1.6614054, and one program that computes the family argmax construction $V(3, 4) = 1.7761424$ *exactly*. Failure taxonomy: 7 of 45 programs produce geometrically invalid packings, 1 imports `random` against the stated contract (caught by the registered gate), none time out. What the channel changes relative to direct emission is dispersion, not escape: code-free anchor concentration at these cells runs 56–80%, the code channel’s pooled anchor-rate is 14/37, and the lost mass sits entirely *below* the anchor — greedy radius-growth heuristics converging under the lattice.

Both registered follow-ups hold, and they harden the result. A fresh-draw replication wave (arm CC2) and a Sonnet-tier probe (arm CCS) were registered together, after CC’s results were known and disclosed as such, 45 invocations each on the byte-identical prompts. **CC2 returned REPLICATED** under its pre-stated verdict map: 0 of 41 valid outputs above the rival (the registered zero-count prediction, exactly), and the anchor bucket modal at 3 of 3 cells — including $N = 31$, where the mode that was one-sample-thin in CC arrived at 10 of 15 valid with a margin of 8. Pooled over both weak-tier waves, 0 of 78 valid program outputs exceed the family argmax and exactly two reach it, one per wave, both at $N = 13$. **CCS returned P-CCS2:** against competing registered predictions (escape at $\geq 20\%$ above-rival versus a zero-count prediction), **0 of 37 valid Sonnet-tier outputs exceed the argmax**. The Sonnet programs are qualitatively different search — simulated annealing over hand-rolled linear-congruential generators, Halton-sequence multi-restarts, force-directed relaxation, against the weak tier’s greedy radius growth — and the tier barely anchors (5 of 37 on the anchor value), yet its executed math-only search never passes the family argmax at any cell. Taxonomy: CC2 41 valid, 3 geometry-invalid, 1 blocked import; CCS 37 valid, 7 geometry-invalid, 1 timeout. One protocol note, disclosed: a single CC2 invocation ($N = 13$, slot 2) used runtime tools despite the no-tools dispatch wrapper; per the arm-M convention its verbatim final message is scored like every other row. Ledgers and scripts for these arms are listed in Appendix D.

3.5 Library programs: the channel with numpy and scipy

Arm CC’s registration named the arm it was not: a library-enabled code channel, the one deployed loops permit. Arm CL is arm CC’s registered prompt with a single clause changed — the import allowlist widens from `math` to `math`, `numpy` and `scipy` — run at two tiers (the vendor’s `claude-haiku-4.5` and `claude-sonnet-4.5` aliases, temperature 1.0, the same pinned path as arm P) at the three trap cells, $n = 15$

per tier and cell, 90 invocations of which 4 (all Sonnet-tier) were refused by the serving path with an HTTP 402 and are recorded as such, 85 returned a parseable program and one Sonnet response failed to parse, each program executed in a fresh `python -I -S` subprocess with a 120-second wall clock and scored by arm F’s conventions with clearance decided only by §2.4’s rule. One execution-mechanism correction is disclosed in the registration itself: an earlier draft of it said the libraries would reach the child through `PYTHONPATH`, which `-I` ignores, so a fixed driver committed with the analysis script inserts the interpreter’s site-packages into `sys.path` and runs the model’s source unmodified via `runpy`; the correction was made before any sampling, in the same commit as the analysis script, and the driver is identical for every row and not part of the scored source. Competing predictions were registered per tier: P-CL1, no valid output clears the family; P-CL2, at least 20% do; and a falsifier F-CL1, the Sonnet tier clears at two or more of the three cells, with the integration rule for that outcome written down before it was known.

Weak tier: none clears, underpowered. 11 valid programs of 45 (3, 2 and 6 by cell), every one calling `scipy.optimize`, and none clears the family. The 34 failures are 22 programs whose stdout the registered parser could not read, 7 geometrically invalid packings, 4 execution errors and 1 blocked import; none timed out. P-CL1 holds as registered. Under the arm’s own five-valid floor only $N = 31$ is scoreable, so the row is 0 of 6 at one cell, with 0 of 11 pooled beside it; the Wilson 95% upper bounds are 39.0% and 25.9%, both above the 20% bar P-CL2 set. The arm can say the weak tier did not clear, not that it cannot at the registered rate; arm CL-W, below, powers the cell.

Post hoc: what the 22 unreadable outputs contain. Every one of the 22 parse failures prints its packing as `numpy 2.x` scalar reprs, `np.float64(0.49...)`, which the registered parser’s `literal_eval` rejects; no Sonnet-tier program did this. Stripping that wrapper, and nothing else, parses all 22 (`cl_lenient_reparse.py`, over a re-execution that reproduced every registered bin; 11 of the 22 call `numpy.random` unseeded, so their sums are this execution’s): 9 are valid at both tolerances (3, 4 and 2 by cell), 13 are invalid packings (7 overlaps, 4 outside the square, 2 non-positive radii), and **one clears**, a deterministic $N = 13$ program with 13 distinct radii at 1.812531608, 2.0% above the argmax. Under this reading the weak tier is 1 of 20 pooled and 1 of 6 at $N = 13$, a rate inside the registered dead zone (above 0, below 20%). The registered count stands as the arm’s result and the table’s entry; the lenient count is reported beside it because a zero that depends on a print format would be the wrong thing to lean on.

Arm CL-W: the top-up. Registered after the reparse (`arm_clw_preregistration.txt`, commit `f73222a`, pushed before the first call) with the pooling rule, both readings and the predictions fixed in advance: arm CL’s weak-tier prompt byte-identical, same path and decoding, 45 more invocations (15 per cell; the draft said 90 and the credit allowed 45, disclosed), pooled with arm CL’s 45 into one 90-row unit scored in one pass by arm CL’s pipeline. Of the 45 new rows 8 are valid (3, 1 and 4 by cell); the failures are 23 unreadable stdouts (again `numpy` scalars), 4 invalid packings, 3 execution errors, 3 blocked imports, 3 timeouts and 1 with no program. Pooled, under the registered parser: **0 of 19 valid clear** (6, 3 and 10 by cell, all valid at 10^{-9} ; the $N = 21$ cell stays under the five-valid floor and claims nothing on its own), Wilson 95% upper bound 16.8% on the pooled count, so the weak library cell now excludes the registered 20% bar; P-CLW1 holds and F-CLW1 does not fire. Under the lenient reading: 41 valid and **4 clear**, 9.8% (upper bound 22.5%), inside the registered dead zone — at $N = 13$ 1.784872559 and 1.812531608, at $N = 21$ 2.331990574, at $N = 31$ 2.851728045, that last a sum arm CL’s Sonnet tier also reached. An independent re-scoring with its own parser reproduced every bin and three of the four sums to the digit; the 1.784872559 program is unseeded and its sum moves with the draw. The weak tier does not clear at rate under either reading.

Sonnet tier: 23 of 25 clear. 25 valid programs of 45 at the primary 10^{-6} (9, 11 and 5 by cell), every one calling `scipy.optimize`, and **23 clear the family**, each of the 23 valid at 10^{-9} as the clearance rule requires; the two non-clearing programs, both at $N = 13$, are also valid at 10^{-9} (sums 1.5 and 1.767686476, under the argmax), so the count is 23 of 25 at either tolerance: 7 of 9 at $N = 13$, 11 of 11 at $N = 21$, 5 of 5 at $N = 31$, the best program per cell reaching 1.820699211 against the argmax 1.776142375, 2.340549845 against 2.258883476 and 2.864990189 against 2.748528137 — 2.5%, 3.6% and 4.2% above what the recipe family can express. Those are best-of-cell margins; across the 23 clearing programs the excess over the argmax runs from 0.67% to 4.24% with a median of 3.62%. P-CL2 holds and **F-CL1 fires** at 3 of 3 cells. Against the published best-known values, three-decimal lower bounds the bound table carries (1.829 at $N = 13$, 2.362 at $N = 21$, 2.889 at $N = 31$; Appendix E.1), the best programs sit 0.5%, 0.9% and 0.8% below.

Arm CCP: the isolating arm. Between the 0 of 37 cell (arm CCS, agent runtime, 10-second budget) and the 23 of 25 cell the serving path, the dispatch context and the execution budget changed with the import allowlist, so neither cell alone said which factor cleared. The arm that answers the objection two outside reads made was registered after them (`arm_ccp_preregistration.txt`, commit `ebf213e`, pushed before the first call): arm CC’s math-only prompt, byte-identical, on arm CL’s path (Sonnet tier, pinned API, temperature 1.0, no system prompt) under arm CL’s 120-second budget, with `numpy` and `scipy` blocked by the AST gate. Predictions were competing: P-CCP1, no valid output clears (the library is what clears); P-CCP2, clearance at 20% or more of valid outputs at two of three cells (path or budget explain arm CL), with the falsifier F-CCP1 rewriting contribution 1 to “capability plus budget” if P-CCP2 held. 45 invocations (one in-flight refusal at row 24, resumed with one worker, request unchanged). 35 valid at both tolerances (10, 14 and 11 by cell; 10 geometrically invalid; no timeouts, no blocked imports) and **0 clear**. None reaches the argmax: the best program per cell sums to 1.625, 2.1 and 2.672523443 against 1.776142375, 2.258883476 and 2.748528137, and at $N = 13$ and $N = 21$ that best is the template anchor $T(k^*, N)$ itself, with the anchor’s structure modal (7 of 10 programs at $k = 4$, 12 of 14 at $k = 5$). P-CCP1 holds; F-CCP1 does not fire. The three Sonnet program cells read 0 of 37 (runtime, 10 seconds), 0 of 35 (pinned, 120 seconds) and 23 of 25 (pinned, 120 seconds, libraries): within the Sonnet tier, path and budget moved nothing and the libraries moved everything. At the weak tier the same libraries move nothing at rate (arm CL-W, above), which is why §1 says leaving the family at rate took the tier *and* the optimizer.

Arm B: the optimizer alone. The blind review of the round that reached the table asked the question its title asks: does a fixed `scipy.optimize` program with no model clear the family? Arm B was registered after that read (`arm_b_preregistration.txt`, commit `1932cc0`, pushed before the first scored run; post-review motivation disclosed) with competing predictions, P-B1, the reference program clears at 20% or more of valid runs at two or more cells (the optimizer alone clears; contribution 1 is rewritten by a rule fixed in the registration), against P-B2, it never clears, and a secondary S-B1 on whether its best packing per cell exceeds arm CL’s Sonnet best. The program is what a reader would write first: random-restart SLSQP over centre and radius variables with analytic gradients, uniform random starts, a uniform radius shrink by the largest constraint violation, 50 restarts; its 94 lines are reproduced in Supplement S7. Its authorship is the manuscript’s (Claude drafting under the human author, §8): version 1 was committed at `1932cc0` and version 2 at `c486843`, both before the single scored run, and no result informed either. “No model” means no model call at proposal time: one fixed program is handed every N and seed, where every other program cell in Table 2 is written afresh by a model at each invocation. It was executed through arm CL’s pipeline unmodified (`python -I -S`, the fixed driver, 120-second wall clock, one core, arm F’s scoring, §2.4’s rule) at the three cells, 15 seeds each. One amendment is disclosed: version 1 bounded its restarts by wall clock and imported `time` and `sys`, which the registered allowlist forbids, so the AST gate blocked all 45 rows; that run is kept in the artifact, and version 2 fixes the count at 50 (one unscored timing run outside the gate fit 80 restarts in 95 seconds at $N = 31$, so 50 leaves margin inside the 120-second wall clock) and changes nothing else. Result: 45 of 45 valid at both tolerances and **45 clear**, 15 of 15 at every cell; best per cell 1.829542412, 2.359585843 and 2.883035274 against the Sonnet tier’s 1.820699211, 2.340549845 and 2.864990189, so S-B1 holds at 3 of 3 cells. Against the published best-known values the reference sits 0.1% and 0.2% below at $N = 21$ and 31, and at $N = 13$ its 1.8295 matches the published 1.829+ to the three truncated digits published. Its packings are off-family (dominant grid orders 3 to 7 across cells, no single template). **P-B1 holds**, and the integration rule applies: the library optimizer clears the family with no model in the loop.

What the arms establish. The optimizer clears the family on its own (arm B). A library-enabled program written by a middle-tier model clears at rate because it invokes that optimizer, and lands below the fixed reference at every cell; a weak-tier program invokes the same library in every one of its 19 valid programs and clears nothing at rate; no math-only program clears at either tier, on either path or at either budget. With arm CCP the attribution is no longer crossed within the Sonnet tier: the library is the factor, and with arm B the library is sufficient. What the tiers differ in is whether their programs drive the optimizer to a clearing packing, which is the execution finding of §4 in the code channel. The 10-second budget of arms CC, CC2 and CCS remains a difference from the library arm for the weak tier’s own contrast, disclosed in §7. Per the registered integration rules, the abstract and title name the library-enabled Sonnet tier as the regime

that clears, and arm CCP’s 35 valid programs are reported here by name and not pooled with the 0 of 37. Ledgers and scripts are listed in Appendix D.

4 Mechanism: choice follows the score table, execution does not

Table 2 says a weak-tier proposer leaves the family at rate in no channel and a Sonnet-tier one leaves it at rate only in the library cell. Two readings of that are open: the weak tier *prefers* the template, or it *cannot build* anything else. One registered arm separates them, and four other directions show the same split.

4.1 Arm CH: the argmax in hand does not dislodge the template

Arm CH hands the model the entire recipe family: the bare prompt plus every admissible $T(k, N)$ and $V(k, m)$ value enumerated, including the family argmax, with the instruction to use whichever scores highest or anything better ($n = 15$ per cell at $N = 13, 21, 31$; 45 invocations; registered together with arm MU before sampling, prompt hashes in `arm_mu_prompts.json`). The §7 tractability alternative says the model truncates at k^* because that is the only construction it can compute; enumerating the family removes the computation burden, so under that reading the modal output should become the stated argmax (P-CH2). Under the template reading it stays $T(k^*, N)$ despite the argmax being printed in the prompt (P-CH1).

Registered verdict: P-CH1 holds (2 of 3 cells) — but the registered metric carries a survivorship artifact we disclose rather than exploit. Under the registered scorer, which takes the modal output *among valid samples*, the template reading wins: at $N = 13$ the modal valid output is $T(4, 13) = 1.6250000$ against a stated argmax of 1.7761424, at $N = 31$ it is $T(6, 31) \approx 2.5833$ (4 of 7 valid) against 2.7485281 (3 of 7) — a one-sample margin — and only $N = 21$ goes to the argmax (2.2588835, 3 of 5 valid). P-CH2 holds at only 1 of 3 cells. This arm sets no evaluability floor of its own, so it grants $N = 13$ a per-cell verdict on two valid samples; under the GM chain’s floor of three that cell would not be scoreable and P-CH1 would hold at 1 of 2. The floors differ because each was fixed at its own arm’s registration, and we report both readings rather than harmonize them afterwards (§2.4). Read no further and the tractability alternative fails the probe. But validity is low in this arm — 2, 5 and 7 valid of 15 by cell — and the failure modes change the picture at the *attempt* level.

The attempt-level reading, post hoc and disclosed. Of the 31 invalid rows, 30 are evaluable, and **all 30 attempt the argmax at its own grid order k^*-1** : 12 of 13 at $N = 13$, 10 of 10 at $N = 21$, 8 of 8 at $N = 31$, with the order backed out of each row’s dominant emitted radius the way Appendix E.2 recovers k . The thirty-first row emits radical literals the scorer cannot evaluate and is therefore not scored as attempting anything. Attempting is not the same as carrying a family filler radius, which would not identify which k was tried; the order does. Of the 30, all misplace the fillers — typically at the container corners, where radius $(\sqrt{2} - 1)/(2k)$ overlaps the adjacent grid circle, against the correct construction’s grid interstices.

The control this reading needs. Attempt-level counts mean nothing without knowing what the model attempts when *nothing* hands it the argmax. Arm F’s bare invocations at the same three cells are that condition, and the same dominant-radius estimator applies to them. Over *all* evaluable rows, valid and invalid alike, the bare arm reaches for the argmax order in **5 of 57** (9%; these 57 are arm F’s evaluable rows at $N = 13, 21$ and 31, not the 57 valid emissions of §3.3’s four cells) against arm CH’s **36 of 44** (82%); Fisher’s exact test gives $p = 2.2 \times 10^{-14}$ one-sided (`arm_f_attempt_uncond.py`, post hoc, disclosed). We report that comparison rather than the invalid-row one because the invalid-row version conditions on a variable the manipulation itself moves: arm CH is 31 of 45 invalid, arm F 10 of 60 at the same cells, so restricting both arms to their invalid rows conditions on a collider. For the record it points the same way — 3 of 7 evaluable invalid bare rows at the rival order against arm CH’s 30 of 30, $p = 5.3 \times 10^{-4}$ (`arm_f_attempt_control.py`) — but it is a diagnostic, not the evidence. Both comparisons are post-hoc strata rather than registered outcomes.

What the arm supports. Counting attempts rather than valid outputs, the model reaches for the argmax in 36 of 45 CH invocations — the 30 evaluable invalid attempts plus 6 valid outputs that land it; the control above quotes 36 of 44 because the forty-fifth row, the radical-literal one, is not evaluable and is excluded from both arms’ denominators there. It also *builds* it: 0 of 2 valid at $N = 13$, 3 of 5 at $N = 21$, 3 of 7 at $N = 31$, 6

of 14 pooled. So the model does not *fail* to build the argmax; it builds it unreliably, in 6 of 14 valid outputs, while reaching for it in 36 of 45 invocations. The registered metric was fixed before sampling and we report its verdict as registered — and flag in the same breath that it conditions on exactly the variable, execution success, that separates the two readings. The honest summary is a decomposition rather than a clean winner: handed the argmax, the model’s *choice* follows the score table, and its *execution* succeeds less than half the time it is attempted, failing off-template in 30 of the 31 invalid attempts on the disclosed post-hoc reading — the tractability alternative is wrong about selection and right about instantiation.

N	Stated argmax	Modal valid output	Valid	Invalid rows attempting argmax	Registered verdict
13	1.7761424	$T(4, 13) = 1.6250000$	2/15	all	P-CH1 holds
21	2.2588835	argmax, 3/5	5/15	all	P-CH2 (argmax)
31	2.7485281	$T(6, 31) \approx 2.5833$ (4/7)	7/15	all	P-CH1 holds (one-sample margin)
pooled	—	template 2 of 3 cells	14/45	30/30 evaluable (all mis-place fillers; 1 further row unevaluable, radicals)	choice follows score table (36/45 attempt it); execution lands it 6/14 valid

4.2 The same asymmetry from four other directions

In code. Delegating the arithmetic to an executed math-only program is the cleanest removal of the hand-computation burden, and it recovers the argmax construction exactly twice in 78 valid weak-tier programs, one per wave, both at $N = 13$ — a program that computes $V(3, 4) = 1.7761424$ — and never passes it (§3.4). The lost probability mass sits below the anchor, not above it: greedy radius-growth heuristics converging under the lattice. With `scipy.optimize` available the weak tier still clears nothing in 19 valid programs under the registered parser (4 of 41 under a post-hoc lenient one), while the Sonnet tier clears in 23 of 25 and, without the libraries on the same path and budget, in 0 of 35 (§3.5): the weak tier’s valid library programs all call `scipy.optimize` and none drives it past the family, so the library supplies execution that the weak tier invokes but does not steer to a clearing packing, while a fixed reference program with no model in the loop steers it there at 45 of 45 (arm B).

At a held-out cell. Arm CN’s one miss, $N = 50$, is the arm-CH pattern in the wild: 10 of 11 valid outputs are the registered 7×7 grid plus exactly one filler, 6 of them with the filler pushed into a container corner instead of the interstice (Appendix E.5). The template is selected; the filler is mis-executed; the value criterion registers a miss.

Across a vendor. A second weak-tier family, gemma-4-26b under a laboratory-grade instrument, does what the original does not: at every discriminating cell its modal output is the rival argmax itself, and a post-hoc structural diagnostic finds all 27 of 27 rival-valued packings carrying the rival’s exact two-radius decomposition, 27 of 43 discriminating-cell validities against the original family’s 2 of 57 (§5.4, Supplement S2). Read against the tractability alternative, the anchor value marks where construction-by-mental-arithmetic lands, and a family with more arithmetic reach executes the harder in-family construction rather than truncating — same recipe family, other branch.

Up the ladder. Holding prompt, container and scoring fixed across three nominal tiers at the three discriminating cells, ambition rises monotonically — truncated template, perturbed hybrid, recursive gasket — while validity rises then collapses, $83\% \rightarrow 100\% \rightarrow 13\%$ on the matched cells (the ladder table and the top alias are Appendix C; the alias-to-weights binding is not attestable). “Ambition” is measurable: a post-hoc diagnostic (`diagnostics_ambition.py`, all samples with emitted geometry, invalid included) gives median distinct radii per sample $1 \rightarrow 2 \rightarrow 3$ and median off-lattice center fraction $0.08 \rightarrow 0.43 \rightarrow 0.95$ across the tiers. The top alias attempts at every cell a more ambitious construction than either other tier — quarter-circle corners with Apollonius-style fillers, coarse grids with border strips — and fails geometrically rather than numerically (24 overlaps and 2 nonpositive radii in its 26 failures; Appendix C). The tier that attempts the most builds the least.

What the four directions share is the order of failure. Selection is not the bottleneck: the weak tier picks the argmax when shown it, a family with more reach emits it unprompted, and the ambitious tier reaches for constructions the family does not contain. Execution is: the picked construction is mis-built in text, reached and not passed in math-only code, and cleared at rate only where a Sonnet-tier model hands the arithmetic to an optimizer. The template is what survives execution, and the library-enabled Sonnet cell is where what survives changes.

5 Zero-shot characterization does not transfer

The weak tier’s template is a closed form, and Appendix E reports the prediction in full: modal at all seven tested N , at four of five held-out N , under a perturbed container and under paraphrase. A characterization that sharp invites the inference that it describes the proposer inside a loop. Four registered arms each changed one thing about the setting — a parent, a generation, the serving path, the reasoning budget — and each carried registered predictions for the inference; two more changed the vendor. The persistence prediction failed in each of the first three: arm MU’s falsifier fired, arm L failed two of its four predictions (a third was registered open), and arm P’s predictions were unscorable because validity collapsed, which arm P-D traces to the budget. This section reports them as results, not as caveats, because the same inference is routinely drawn from zero-shot characterizations of LLM proposers.

5.1 Arm MU: the anchor dissolves after one conditioning step

Arm MU issues one-parent mutation calls: the bare prompt plus an existing packing and the instruction to modify it so the sum of radii increases. Three parents per cell ($N = 13, 21, 31$; $n = 15$ per condition-cell, 135 invocations): A_anchor , the model’s own unconditioned anchor $T(k^*, N)$; B_rival , the provably better in-family rival $V(k^* - 1, m)$; and $C_offfamily$, a diagonal-chain packing scoring ≈ 0.59 – 0.62 , far below either. It was registered together with arm CH before sampling, with numeric decision thresholds, a power analysis (0.83 at the registered effect size), tie-against conventions, and a falsifier (F-MU1) whose consequence — the anchoring claim is scoped to *unconditioned* calls only — was written down before the first invocation.

F-MU1 triggered; the anchoring result dissolves under conditioning. Under B_rival the pooled anchor-rate is 0 of 26 valid (0%, CI [0.00, 0.13]) against a registered DISSOLVES threshold of $\leq 20\%$; keep-or-improve is 26 of 26 against a threshold of $\geq 50\%$. Every valid B_rival sample keeps or refines the rival construction, and every one inherits the rival’s grid order (k -inheritance 26/26, disconfirming P-MU3’s $< 50\%$ prediction). The template does not reassert itself when a better in-family parent is present. A_anchor makes the same point from the other side: even with the model’s *own* anchor as parent and an explicit increase-the-sum instruction, only 7 of 28 valid samples (25%, CI [0.13, 0.43]) stay on the anchor — P-MU1’s $\geq 50\%$ prediction is not met. The one condition where the template does reassert itself is the one the registration predicted it would: given the *off-family* diagonal parent, 20 of 34 valid samples (58.8%, CI [0.42, 0.74]) abandon the parent entirely and emit a k^* -grid construction — P-MU4 holds per-cell at $N = 21$ (7/13) and $N = 31$ (12/14), though not at $N = 13$ (1/7). A good in-family parent is followed, a bad off-family parent is discarded for the template: the attractor is real but weak, capturing the trajectory only when the parent is worse than what the template supplies.

Parent	Valid	Anchor-rate (CI)	Registered threshold	Keep-or- improve	k -inher.	Reading
B_rival (in-family argmax)	26	0/26 (0%, [0.00, 0.13])	DISSOLVES if $\leq 20\%$	26/26 (bar $\geq 50\%$)	26/26	F-MU1 triggered — anchoring does not transfer to one-parent calls
A_anchor (model’s own anchor)	28	7/28 (25%, [0.13, 0.43])	P-MU1 predicted $\geq 50\%$	—	—	not met — even the anchor-as-parent does not hold the anchor
$C_offfamily$ (diagonal, ≈ 0.6)	34	20/34 abandon parent for k^* -grid (58.8%, [0.42, 0.74])	P-MU4: template reasserts	—	—	holds at $N = 21, 31$; not at $N = 13$ (1/7)

No conditioned output clears the family argmax (post hoc, disclosed). Arm MU’s rows were scored against the *anchor*, so whether any conditioned output clears what the family can express was never asked of them at registration. It is asked here. Recomputing the family argmax from the closed form — 1.776142375 at $N = 13$, 2.258883476 at $N = 21$, 2.748528137 at $N = 31$ — arm MU’s 135 invocations give 88 valid outputs at the primary tolerance and 63 at 10^{-9} . Eighteen of the 88 sit above the family argmax by more than the 5×10^{-8} granularity at which their own sums are recorded, by at most 1.5×10^{-6} , and *every one of the eighteen fails the 10^{-9} test*: they are packings whose overlap is small enough to pass the loose gate, and the excess is the overlap. Twelve further rows exceed the argmax only within that recording granularity, which is to say they are the argmax, written down. Arm CH’s 45 invocations add 14 valid outputs at both tolerances, none above. The verdict rests on the 10^{-9} leg, where none clears, and it is corroborated at structure level: each of the eighteen carries an overlap the tighter gate rejects, and the excess is bounded above by the overlap. This re-reading adds no invocations to the study total; the recomputation is `arm_mu_ceiling.py`.

5.2 Arm L: iterating the loop moves the anchor; nothing clears the family argmax

Every arm to this point samples once. The regime that discovery loops actually run in iterates: a proposal is scored, kept or discarded, and fed back as a parent. Arm L runs the minimal generational loop and reports the four registered predictions whether or not they held. The design is registered before sampling, with both prompt templates hashed and the stepper and scorer committed alongside. Generation 0 is the bare template A.1, byte-identical to arms F, M and CN; generations 1–5 are arm MU’s registered mutation template verbatim, instantiated with the parent packing and its score. Two discriminating cells, $N = 13$ and $N = 31$, are crossed with two archive regimes — GREEDY, which retains only the single best valid proposal, and DIVERSE, which retains up to five distinct-valued ones — giving four lineages of six rounds at population 5: 120 invocations, of which 2 were runtime-rejected and counted. The parent for population slot i is `archive[i mod |archive|]`, so the loop carries no sampling randomness of its own. All four lineages are evaluable under the registered floor. This is a minimal analogue and the registration says so: population 5, six rounds, one-parent mutation, no islands, no crossover, no program representation, and no evaluator signal beyond the parent’s own score. Nothing here transfers to deployed systems, and no claim about any named system follows.

L2 confirmed; L1 and L3 not met. L2 (dissolution survives iteration) holds in both cells: the on-prediction rate falls from 4 of 9 valid generation-0 proposals to 5 of 20 conditioned ones at $N = 13$, and from 7 of 8 to 6 of 29 at $N = 31$. The falsifier F-L1 did not fire. L1 (the registered prediction is modal at generation 0 in every lineage) is **not met**, on one lineage: the 13-DIVERSE seed round returned four valid proposals split two on the prediction and two elsewhere, and the registered tie rule counts the tie against it. The other three lineages pass. The failure is a power artifact of the registration — we set four separate per-lineage bars over roughly four valid proposals each rather than one pooled bar per cell — and it is reported as scored, not rescued. L3 (loop diversity is scaffold-produced, so the greedy archive should return proposals toward the family) is **not met, with the direction reversed**: the conditioned on-prediction rate is 3 of 22 under GREEDY against 8 of 27 under DIVERSE. The comparison is confounded — greedy and diverse lineages differ in parent identity as well as archive width, and 13-GREEDY produced only 4 valid proposals across five conditioned rounds — so the reversal is reported as a disconfirmation of the registered prediction and not as evidence for its converse.

L4: no lineage clears the family argmax. L4 was registered open, with no direction predicted: report best-of-run against the family argmax. **No lineage clears it.** At $N = 13$ both lineages top out exactly at the family argmax 1.7761424; at $N = 31$ GREEDY reaches 2.6104167 and DIVERSE lands exactly on the family argmax 2.7485281. Across 49 valid conditioned outputs, zero exceed the best value the recipe family can express. Iteration moves the proposer off the anchor without moving it past the family *in these arms*. One correction is disclosed: the registered configuration stored the family argmax to seven decimals, and the clearance test compares at 10^{-9} ; a proposal landing exactly on the rival at $N = 31$ therefore tested as *above* the truncated literal by 3.7×10^{-8} . The scorer now recomputes the argmax in closed form, agreeing with the stored value to 1.3×10^{-15} ; the change is post-sampling and affects only the 10^{-9} clearance test (corrections ledger item 34). Scale is the honest limit: five conditioned generations, population 5, 49 valid conditioned outputs, one tier. The arm shows that this loop, at this size, on this channel, does not clear

the family; it cannot show that a larger one would not. Raw rows in `arm_1_collect.jsonl`; frozen report `arm_1_report.json`, scored once on the complete ledger. The study corpus stood at 1068 invocations at that point.

5.3 Arm P: the registered prompt on a bare, pinned serving path

Every direct-emission arm above ran through an agent runtime whose decoding parameters are not exposed and whose conditioning context the prompt digests lock only in part (§3.1). An external referee named that the paper’s largest instrument hole. Arm P answers it the only way it can be answered: the registered prompts, byte-identical — the builder asserts every hash on file, and gives $N = 43$ its first — sent through a serving path that exposes and pins temperature (OpenRouter chat-completions, the vendor’s own `claude-haiku-4.5` alias, temperature 1.0, no system prompt), at $n = 15$ per cell across the seven square cells, the five held-out cells and the three code cells: 225 invocations, 219 returning content and 6 refused by the serving path and recorded as such. It was registered with competing predictions on mode identity (P-P1, P-P2), on the code-channel clearance (P-P3, P-P4) and a falsifier on the discriminating-cell rate (F-P1), and it is a new arm on a new instrument, reported beside arms F, CN and CC and merged into none of them.

Validity collapsed. At the seven square cells, 0 of 105 outputs are valid at either tolerance. Every failure is an overlap: the responses are well-formed lists of the requested count with no prose and no fences, emitted with no reasoning tokens, and their median maximum pairwise overlap at $N = 13$ is 0.10, fifty times the value window. At the held-out cells, 4 of 75 are valid, all four at $N = 50$ and all four summing to 2.5, a full unit below the family argmax there (3.5295867); none clears. The direct-emission predictions are therefore UNSCOREABLE at every cell under the arm’s own floor of five valid outputs, and F-P1’s denominator is empty — the registration did not anticipate a collapse, and we report that rather than read a zero either way. On the code cells, arm CC’s prompt byte-identical and executed and scored exactly as arm CC’s (10-second timeout), the picture is the paper’s: 12 valid programs of 45 (4, 6 and 2 by cell, only $N = 21$ above the floor) and none clears the family under §2.4’s rule; P-P3 holds.

Two post-hoc diagnostics, disclosed as such. Temperature is not the cause: the $N = 13$ prompt at temperature 0.0, 0.5 and 1.0 gives 0 of 6 valid at each, all overlaps. The instrument is: the same prompt with the registered dispatch wrapper, through the agent runtime’s haiku alias that arms F, CN and CC used, six subagents in parallel on the same day, gives 6 of 6 valid at both tolerances and lands the registered anchor $T(4, 13) = 1.625$ in 4 of 6, with one 3×3 grid carrying corner fillers at 1.6144 and one centre-plus-border construction at 1.7500, above the anchor and below the family argmax.

Arm P-D: which component of the instrument. Registered after an outside review asked for a diagnostic beyond temperature (`arm_pd_preregistration.txt`, commit 5554cc0, pushed before the first call): the same $N = 13$ prompt, same alias, same pinned path and decoding, under two one-line manipulations, $n = 15$ each. D1 turns the vendor’s extended thinking on (the routing layer’s default effort; median 3160 reasoning tokens per row) with no system prompt; D2 leaves thinking off and adds one system line, “You are a careful assistant. Think before you answer.” Competing predictions: P-PD1, D1 valid in at least 8 of 15 and D2 in at most 3 (the reasoning budget is the cause); P-PD2, the mirror image (the context is); P-PD3, anything else. The arm closed at 27 of 30 rows (14 D1, 13 D2) by a registered amendment after the serving path refused the last three for credit; they could not have changed the verdict, since D2 can reach at most 2 valid and D1 already holds 12. Result: **D1 is valid in 12 of 14** (11 at 10^{-9} ; the two failures overlap), **D2 in 0 of 13**, all overlaps, exactly arm P’s collapse. P-PD1 holds. And the anchor comes back with the budget: the modal valid D1 output is $T(4, 13) = 1.625$ (5 of 12, the 4×4 grid truncated), $k = 4$ is the modal structure (6 of 12), and one output is the family argmax 1.7761424 itself, a construction the runtime rows never reach at this cell (0 of 18 valid; both runtime rival hits are at $N = 21$). S-PD1 is met.

What the arm establishes. The anchoring result is a property of the instrument that produced it, and arm P-D names the component: the reasoning budget. The same prompt, sent bare to the same vendor’s alias with temperature pinned, does not yield valid packings from this model class at all; give the pinned call the vendor’s extended thinking and it yields valid packings at the runtime’s rate with the template modal, while a system line that asks for care changes nothing. The runtime’s undisclosed conditioning context — the fragment §3.1 disclosed the digests do not lock — is not what carries the template; the thinking budget

is. That is why Table 2 carries an instrument column and never pools across it, and it is a reproducible instrument, since the budget is a request parameter on the pinned path. Ledgers and scripts are listed in Appendix D.

5.4 Across vendors: a boundary that runs both ways

Two registered arms send the same byte-identical prompts across vendors; the verdict table, arm GM3 per cell and arm V in full are Supplement S2. **Arm GM3** runs the unconditioned-call protocol against a second vendor’s open-weight family (gemma-4-26b-a4b-it) through its first-party API, seven N cells at twenty samples each, predictions and falsifier fixed before sampling. The pooled prediction held — 46 of 80 valid packings (57.5%) land within 2×10^{-3} of the value predicted by §2.3’s rule, against a registered 30% bar — and the pass rides on the non-discriminating cells: on the three discriminating cells the rate is 11 of 43 (25.6%), below the bar; the cell-level prediction did not hold (modal in 2 of 5 scoreable cells, short of the registered 5) and the falsifier (failure in ≥ 4) did not trigger either, so the registration left a dead zone and the outcome landed in it. The three misses are the finding §4.2 uses: the family’s modal output at every discriminating cell is the rival argmax itself, 27 of 43 discriminating-cell validities carrying its exact two-radius decomposition. **Arm V** varies vendor and serving class at once, thirteen free-tier aliases at five cells and five samples each (877 attempt rows over 325 registered slots). Ten of thirteen aliases finished below the registered validity floor and claim nothing in either direction. Of the three that cleared it, two transfer at point estimate — 8/12 and 11/18 on-prediction, Wilson intervals spanning the 50% bar — and gemma-4-31b does not, at 0/15. Restricted to the cells that can discriminate anchoring from family search the decomposition gives 3 of 4 and 5 of 9, and neither separates from the template-shape null of Appendix E.2 (exact upper tails 0.111 and 0.145); zero rival-argmax emissions in 16 discriminating-cell validities.

So the anchor is neither a single-vendor artifact nor universal: it transfers at the pooled level to two further vendors, one weak-tier family takes the rival instead, and a sibling of that family at the same floor anchors on nothing. With arms MU, L and P this closes the boundary from every side the study reached: the template is a property of one tier, one channel, one instrument, and one conditioning regime, and it is inside those four conditions that the closed form names the modal output in advance.

6 Related work

LLM-driven discovery and the circle-packing scoreboard. Language Model Crossover (Meyerson et al., 2024), ELM (Lehman et al., 2022), EvoPrompt (Guo et al., 2024), EoH (Liu et al., 2024) and LLaMEA (van Stein & Bäck, 2025) established the LLM call as an EC operator over prompts and programs; FunSearch turned the pattern into a discovery claim, and every system in §1’s scoreboard inherits it. Two orderings in that scoreboard matter: HELIX’s figure is marginally *below* ShinkaEvolve’s relaxed-tolerance figure (and above its strict one) while being described as state of the art, and ThetaEvolve claims new best-known bounds with a printed $n = 26$ figure equal to HELIX’s, so the benchmark is saturated across the AlphaEvolve–ShinkaEvolve–HELIX cluster and the claims outrun the digits. OpenEvolve (Sharma, 2025), the open-source AlphaEvolve reproduction most practitioners run, reports 2.634292 in its published circle-packing example (at the time of access; the repository is unpinned). Appendix E.1 records the correction to our record comparison and its $N > 40$ coverage gap. Two critiques bear on our framing and are frequently conflated, so we attribute them explicitly: Gideoni et al. (2026) show simple baselines recover much of the reported advantage, and Berthold et al. (2026) show classical solvers do too.

Template convergence and diversity collapse. “Mutation Without Variation” (Gurkan et al., 2026) finds iterated LLM program mutation collapsing onto previously seen structural templates — in 87% of chains, over 93% of mutations revisit a previously seen form — which is the same bias family in a different regime: mutation loops, program space, no closed form, no preregistration. It is also the closest published neighbour to our arm L, whose registered L2 confirmed the analogous statement inside a scored loop: iterating the archive moves the anchor and no valid output clears the family argmax (§5.2). Their result is about structural revisiting; ours is about a value bound that the revisited structures cannot cross. On anchoring, Huang et al. (2026) treats anchoring as a general bias over initial information, Lou & Sun (2024) document numeric primes dragging point estimates, and Malberg et al. (2025) test anchoring on a numeric prime among

thirty biases; our modality is the difference, since a construction template has a computable value where a scalar anchor has only a measurable pull. Concurrently, “Measuring the Gap Between Human and LLM Research Ideas” (Chen et al., 2026) reports the same collapse one level up, in idea space: against 11,683 human ideas, LLM research ideas concentrate on bridge-type opportunities (47–64% versus 12% for humans; normalized entropy down to 0.55 versus >0.92), and chain-of-thought pushes the distribution *further* from the human one. That result is distributional; ours is the limiting case where the collapsed distribution’s mode is a single closed-form-predictable object — narrowness measurable there, computable here. “Artificial Hivemind” (Jiang et al., 2025) documents the same homogeneity across open-ended domains at survey scale, which makes cross-domain generality of the collapse the expected default; against that backdrop our contribution is not that collapse occurs but that, in this regime, its mode is a single closed-form-computable object, stated before sampling. Converging skepticism about what the proposer contributes comes from “Dictionaries, Not Darwin” (Li, 2026), Wang et al. (2026), BehaveSim (Zhang & Lu, 2026), Strategy Diversity (Yang et al., 2026) and the bin-packing critiques (Herrmann & Pallez, 2026; Sim et al., 2025). Two results cut the other way and belong on the other side of the ledger: Zhang et al. (2024) provides empirical grounding for the *importance* of evolutionary search, which is evidence against our substitution of an unconditioned call for the loop and is engaged in §7; and Zhang et al. (2026b) finds strong LLM optimizers act as *local refiners*, which is what a parent-conditioned call would be, so it bears on §7’s mutation arm.

Scaling inversions. Zhou et al. (2024) show larger and more instructable models becoming less reliable; the o3/o4-mini system card (OpenAI, 2025) reports a higher PersonQA hallucination rate for o3 than for o1 alongside higher accuracy, attributed to more claims made; Shojaee et al. (2025) reports collapse past a complexity threshold; GeoBuildBench (Kim et al., 2026) finds geometric construction a regime where nominal capability and executed correctness diverge, and MathConstruct (Balunović et al., 2025) benchmarks constructive proofs, where the object built is checkable. Our instantiation is *attractor family* versus executability.

Preregistration lineage. Thomas et al. (2026), Williams et al. (2026), Vaccaro (2026) and Zhang (2026) preregister recipes, outcome-blinded predictions and agent protocols; HindsightBench (Jia, 2026) releases frozen preregistrations of directional aggregate hypotheses. We adopt the standard rather than claim it. What we add is the object locked — exact closed-form point predictions of specific output values. Structural precedents: Kaplan et al. (2020) and Hoffmann et al. (2022) publish a functional form then run the confirming instance; Snell et al. (2024) is the closest structural twin.

7 Limitations and untested alternatives

Scale. §6’s evidence that evolutionary search contributes (Zhang et al., 2024) and that strong LLM optimizers act as local refiners (Zhang et al., 2026b) is answered by arm MU’s one conditioning step (§5.1) and arm L’s five generations of a registered archive-and-mutate loop (§5.2); what stays open is scale — population five over five generations at one tier is a loop, not the loops of §6. The mechanism-free reading of the concentration — the $k \times k$ grid is what survives mental arithmetic — was probed by the preregistered arm CH (§4.1) and split in two: *choice* follows the stated score table, so tractability does not explain the selection; *execution* fails off-template on that arm’s disclosed post-hoc attempt-level reading, so tractability does explain the instantiation. Typicality bias in preference-tuned models (Zhang et al., 2026a) predicts concentration on the most typical construction without reference to geometry; no observation here excludes it and it was not tested against the tractability reading.

Contamination is probed at one level only. $N = 26$ is the canonical published cell and plausibly in training data; membership and extraction probes of the kind used to detect benchmark contamination (Carlini et al., 2021; Golchin & Surdeanu, 2024) were not run. Arms CN, CP and RP (Appendix E.5, Appendix E.6) bound the recall mechanisms — held-out N , a container that is not the benchmark’s, a direct request for the published values — and a planted-canary test is inapplicable to a benchmark we did not author. Absence of the fitted cells from the pretraining corpus is established by none of it.

Vendor scope. All tiers in the main study come from one provider; §5.4 moves the pooled-level claim across vendors. One finding belongs in a limitations section: a cross-vendor null under a fixed output budget is not evidence about the model until the budget’s interaction with the vendor’s response structure is checked.

Arm GM2 returned 0 of 140 parseable at a 4,096-token budget; the same model at 16,384 tokens parses at 123 of 140, because the vendor API returns deliberation on a separate channel that the smaller budget died inside. Scope after §5.4: the anchoring law is cross-vendor at the pooled level, does not separate from Appendix E.2’s template-shape null at the discriminating cells, and has its full cell-level mode structure confirmed only within the original family.

Two instruments, and the contrast that crossed them. The agent runtime does not expose decoding parameters, so every runtime row is a distributional claim over an observed sampling regime; the pinned rows (arms P, CL, CCP and CL-W) exist to bound that, arm P shows the two paths are not interchangeable for the weak tier (§5.3), and Table 2 therefore never pools across instruments. The comparison the headline once leaned on crossed them: the Sonnet math-only cell (0 of 37, arm CCS) ran on the runtime and the Sonnet library cell (23 of 25, arm CL) on the pinned path, and they also differ in execution budget — 10 seconds for arms CC, CC2 and CCS against 120 seconds for arm CL, a $12\times$ gap registered on purpose (a 10-second gate measures the timeout, not the search) that the search programs the Sonnet tier writes can nonetheless be bound by. Arm CCP (§3.5) isolates the library at the Sonnet tier — arm CC’s prompt on the pinned path at 120 seconds, registered after those reads — and returns 0 of 35 clear, none at the argmax, so the headline attributes the Sonnet gap to the library alone. Within the weak tier the contrast is same-instrument (0 of 12 and 0 of 11, both pinned) but crosses the budget instead: its math-only cell, arm P’s code cells, ran at 10 seconds, exactly as arm CC’s did. That is what remains crossed, and no weak-tier program in any arm has passed the argmax under the registered parser, at either budget; the lenient reading of the library arms finds 4 of 41 (§3.5).

opus_alias provenance. The label denotes the serving alias addressed, not an identified model version. Without pinned weights we cannot separate a model-tier effect from a serving-path effect, and the supporting anomalies live only in the runtime transcript.

Prompt robustness. Every arm other than PP conditions on a single hash-locked prompt; locking buys reproducibility, not robustness to rewording. Arm PP samples three semantic-preserving paraphrases at the two discriminating cells and the registered prediction stays modal at all six paraphrase-cells (Appendix E.6), so the anchor is not keyed to the literal template string at those cells; a formatspread-style sweep on the scale of Sclar et al. (2024) remains untested. The cross-vendor concentration differences in §5.4 are measured but not explained; alignment-induced diversity loss (Kirk et al., 2024) is an untested account of them.

8 Reproducibility, deviations, and disclosures

Provenance in brief; Appendix D carries it in full:

- **Registration and provenance.** Every arm’s prompts were hashed and its predictions and falsifiers fixed before its sampling; no arm was selected on its result, and every arm is reported. Appendix D lists each registration with its commit, ledgers and scripts, the limits on what the prompt digests lock, a metadata correction to the shipped ledger, the excluded records, and the internal adversarial check. Deviations are tabled in Appendix B.
- **Use of AI systems.** Claude models were used to draft manuscript prose and to assist implementation; the human author designed the study, set the research questions, specified the instruments and their decision rules, approved each preregistration before sampling, made the final inclusion, stopping and submission decisions, and is solely responsible for the content. The claim a reader can check is the ordering, not the authorship: each preregistration commit is a git ancestor of the sampling it governs. Because the drafting models share a family with arms under study (§3.3), what a reader is asked to trust is the released artifact — the scripts, raw ledgers, LP oracle and frozen outputs, all of which accompany this submission as anonymized supplementary material — not the author.
- **Artifact availability.** The complete artifact — ledgers, analysis and reproduction scripts, LP oracle, every preregistration and amendment, frozen reports, corrections ledger — accompanies this submission as anonymized supplementary material, together with the supplement PDF, which is part of the same submission; every number regenerates via `HOW_TO_RUN.md` from the shipped ledgers.

Only the pinned-path arms (P, CL, CCP, CL-W and P-D) can be re-sampled: the runtime rows depend on a dispatch wrapper and inherited context the digests do not lock (§3.1), so a reader can re-score them but cannot re-draw them. The working repository is a public remote, unlinked here for anonymity.

- **Which arms are in the body.** Every registered arm has a row in Table 1 and a ledger named in the section that reports it or in Appendix D. An arm has a body subsection when its registered verdict bears on Table 2, on the mechanism (§4) or on the transfer boundaries (§5). Arms whose verdicts rest on non-discriminating cells, on fewer valid samples than their floor, or on a null that does not separate from a template-shape null are reported outside the body, and the cases differ. Arms T and G are supplement-only, in Supplement S6 and Supplement S5; arm B’s reference program is listed verbatim in Supplement S7. Arms V and GM3 are split: their registered pooled verdicts rest on non-discriminating or below-floor cells and are reported in full in Supplement S2, while their discriminating-cell findings carry a paragraph in §5.4. Arm M keeps its registered verdicts in Appendix E.4 and its per-cell decomposition is Supplement S4. The `opus_alias` post-hoc dual accounting is Supplement S3. The pooled-zero table of Supplement S1 is a pooling over body arms, not an arm of its own. The criterion moved no number.
- **Stopping rule.** The manuscript closed after implementing the internal-check findings; no further data-dependent analysis was performed. Arms P and CL (§5.3, §3.5) were registered together after an external review and pushed before either was sampled, and each was sampled and scored once; arm P’s two post-hoc diagnostics are labelled as such in Table 3. Arms CC, CC2 and CCS (§3.4), CN (Appendix E.5), CP, RP (with its control amendment) and PP (Appendix E.6), and L (§5.2) were each registered, run and scored once under their own preregistrations after that close, each analysis script committed before its sampling and not modified after, as were arms CCP and CL-W (§3.5) and P-D (§5.3), each registered in response to an outside read and pushed before its first call, and arm B (§3.5), registered the same way, run under one disclosed amendment and scored once. The analyses named but not run are recorded in Supplement S6.5 and Appendix E.3.

References

- Mislav Balunović, Jasper Dekoninck, Nikola Jovanović, Ivo Petrov, and Martin Vechev. Mathconstruct: Challenging llm reasoning with constructive proofs, 2025. URL <https://arxiv.org/abs/2502.10197>.
- Timo Berthold, Dominik Kamp, Gioni Mexi, Sebastian Pokutta, and Imre Polik. Out-of-the-box global optimization for packing problems: New models and improved solutions, 2026. URL <https://arxiv.org/abs/2605.04850>.
- Nicholas Carlini, Florian Tramèr, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Úlfar Erlingsson, Alina Oprea, and Colin Raffel. Extracting training data from large language models. In *30th USENIX Security Symposium*, 2021.
- Mert Cemri, Shubham Agrawal, Akshat Gupta, Shu Liu, Audrey Cheng, Qiuyang Mang, Ashwin Naren, Lutfi Eren Erdogan, Koushik Sen, Matei Zaharia, Alex Dimakis, and Ion Stoica. Adaevolve: Adaptive llm driven zeroth-order optimization, 2026. URL <https://arxiv.org/abs/2602.20133>.
- Ziyu Chen, Yilun Zhao, and Arman Cohan. Measuring the gap between human and llm research ideas, 2026. URL <https://arxiv.org/abs/2607.01233>.
- Erich Friedman. Packing center: circles in squares, maximizing the sum of radii. <https://erich-friedman.github.io/packing/cirRsqu/>, n.d. Best-known values transcribed into `n_sweep_forecast.py`; accessed 2026.
- Yonatan Gideoni, Sebastian Risi, and Yarin Gal. Simple baselines are competitive with code evolution, 2026. URL <https://arxiv.org/abs/2602.16805>.
- Shahriar Golchin and Mihai Surdeanu. Time travel in LLMs: Tracing data contamination in large language models. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2308.08493.

- Qingyan Guo, Rui Wang, Junliang Guo, Bei Li, Kaitao Song, Xu Tan, Guoqing Liu, Jiang Bian, and Yujiu Yang. Connecting large language models with evolutionary algorithms yields powerful prompt optimizers. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2309.08532.
- Can Gurkan, Forrest Stonedahl, and Uri Wilensky. Mutation without variation: Convergence dynamics in llm-driven program evolution, 2026. URL <https://arxiv.org/abs/2606.05408>.
- Julien Herrmann and Guillaume Pallez. An in-depth study of llm contributions to the bin packing problem, 2026. URL <https://arxiv.org/abs/2510.27353>.
- Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, Tom Hennigan, Eric Noland, Katie Millican, George van den Driessche, Bogdan Damoc, Aurelia Guy, Simon Osindero, Karen Simonyan, Erich Elsen, Jack W. Rae, Oriol Vinyals, and Laurent Sifre. Training compute-optimal large language models, 2022. URL <https://arxiv.org/abs/2203.15556>.
- Yiming Huang, Biquan Bie, Zuqiu Na, Weilin Ruan, Songxin Lei, Yutao Yue, and Xinlei He. Understanding the anchoring effect of llm with synthetic data: Existence, mechanism, and potential mitigations, 2026. URL <https://arxiv.org/abs/2505.15392>.
- Haozhe Jia. Hindsightbench: A black-box behavioral audit protocol for parametric hindsight in time-indexed llm decision tasks, 2026. URL <https://arxiv.org/abs/2607.18867>.
- Liwei Jiang, Yuanjun Chai, Margaret Li, Mickel Liu, Raymond Fok, Nouha Dziri, Yulia Tsvetkov, Maarten Sap, Alon Albalak, and Yejin Choi. Artificial hivemind: The open-ended homogeneity of language models (and beyond), 2025. URL <https://arxiv.org/abs/2510.22954>.
- Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. Scaling laws for neural language models, 2020. URL <https://arxiv.org/abs/2001.08361>.
- Valentin Khrulkov, Andrey Galichin, Denis Bashkirov, Dmitry Vinichenko, Oleg Travkin, Roman Alferov, Andrey Kuznetsov, and Ivan Oseledets. Gigaev0: An open source optimization framework powered by llms and evolution algorithms, 2025. URL <https://arxiv.org/abs/2511.17592>.
- Jinwoong Kim, Rui Yang, and Huishuai Zhang. Geobuildbench: A benchmark for interactive and executable geometry construction from natural language, 2026. URL <https://arxiv.org/abs/2605.13167>.
- Robert Kirk, Ishita Mediratta, Christoforos Nalmpantis, Jelena Luketina, Eric Hambro, Edward Grefenstette, and Roberta Raileanu. Understanding the effects of RLHF on LLM generalisation and diversity. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2310.06452.
- Robert Tjarko Lange, Yuki Imajuku, and Edoardo Cetin. Shinkaevolve: Towards open-ended and sample-efficient program evolution, 2025. URL <https://arxiv.org/abs/2509.19349>.
- Joel Lehman and Kenneth O. Stanley. Abandoning objectives: Evolution through the search for novelty alone. *Evolutionary Computation*, 19(2):189–223, 2011.
- Joel Lehman, Jonathan Gordon, Shawn Jain, Kamal Ndousse, Cathy Yeh, and Kenneth O. Stanley. Evolution through large models, 2022. URL <https://arxiv.org/abs/2206.08896>.
- Pan Li. Dictionaries, not darwin: Set-level selection beats llm evolution in scientific equation discovery, 2026. URL <https://arxiv.org/abs/2607.04108>.
- Fei Liu, Xialiang Tong, Mingxuan Yuan, Xi Lin, Fu Luo, Zhenkun Wang, Zhichao Lu, and Qingfu Zhang. Evolution of heuristics: Towards efficient automatic algorithm design using large language model. In *International Conference on Machine Learning (ICML)*, 2024. arXiv:2401.02051.
- Jiaxu Lou and Yifan Sun. Anchoring bias in large language models: An experimental study, 2024. URL <https://arxiv.org/abs/2412.06593>.

- Simon Malberg, Roman Poletukhin, Carolin M. Schuster, and Georg Groh. A comprehensive evaluation of cognitive biases in llms, 2025. URL <https://arxiv.org/abs/2410.15413>.
- Elliot Meyerson, Mark J. Nelson, Herbie Bradley, Adam Gaier, Arash Moradi, Amy K. Hoover, and Joel Lehman. Language model crossover: Variation through few-shot prompting, 2024. URL <https://arxiv.org/abs/2302.12170>.
- Jean-Baptiste Mouret and Jeff Clune. Illuminating search spaces by mapping elites. *arXiv preprint arXiv:1504.04909*, 2015.
- Alexander Novikov, Ng  n V  , Marvin Eisenberger, et al. Alphaevolve: A coding agent for scientific and algorithmic discovery, 2025. URL <https://arxiv.org/abs/2506.13131>.
- OpenAI. OpenAI o3 and o4-mini system card. <https://openai.com/index/o3-o4-mini-system-card/>, 2025.
- Bernardino Romera-Paredes, Mohammadamin Barekatain, Alexander Novikov, Matej Balog, M. Pawan Kumar, Emilien Dupont, Francisco J. R. Ruiz, Jordan S. Ellenberg, Pengming Wang, Omar Fawzi, Pushmeet Kohli, and Alhussein Fawzi. Mathematical discoveries from program search with large language models. *Nature*, 625(7995):468–475, 2024. doi: 10.1038/s41586-023-06924-6.
- Melanie Sclar, Yejin Choi, Yulia Tsvetkov, and Alane Suhr. Quantifying language models’ sensitivity to spurious features in prompt design or: How i learned to start worrying about prompt formatting. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2310.11324.
- Asankhaya Sharma. Openevolve: an open-source implementation of AlphaEvolve. GitHub repository, <https://github.com/codelion/openevolve>, 2025. Repository unpinned; circle-packing example value read at time of access, 2026.
- Parshin Shojaee, Iman Mirzadeh, Keivan Alizadeh, Maxwell Horton, Samy Bengio, and Mehrdad Farajtabar. The illusion of thinking: Understanding the strengths and limitations of reasoning models via the lens of problem complexity, 2025. URL <https://arxiv.org/abs/2506.06941>.
- Kevin Sim, Quentin Renau, and Emma Hart. Beyond the hype: Benchmarking llm-evolved heuristics for bin packing, 2025. URL <https://arxiv.org/abs/2501.11411>.
- Charlie Snell, Eric Wallace, Dan Klein, and Sergey Levine. Predicting emergent capabilities by finetuning, 2024. URL <https://arxiv.org/abs/2411.16035>.
- Chang Su, Zhongkai Hao, Zhizhou Zhang, Zeyu Xia, Youjia Wu, Hang Su, and Jun Zhu. Helix: Evolutionary reinforcement learning for open-ended scientific problem solving, 2026. URL <https://arxiv.org/abs/2603.07642>.
- Maria Thomas, Kristina Gligoric, and Nihar B. Shah. Mitigating llm-based p-hacking by preregistering for the next llm, 2026. URL <https://arxiv.org/abs/2606.27687>.
- Michelle Vaccaro. Preregistration for experiments with ai agents, 2026. URL <https://arxiv.org/abs/2606.11217>.
- Niki van Stein and Thomas B  ck. Llamea: A large language model evolutionary algorithm for automatically generating metaheuristics, 2025. URL <https://arxiv.org/abs/2405.20132>.
- Haorui Wang, Parshin Shojaee, Kazem Meidani, Kunyang Sun, Jos   Miguel Hern  ndez-Lobato, Teresa Head-Gordon, Jiajun He, Chandan K. Reddy, Chao Zhang, and Yuanqi Du. Towards diverse scientific hypothesis search with large language models, 2026. URL <https://arxiv.org/abs/2606.10587>.
- Yiping Wang, Shao-Rong Su, Zhiyuan Zeng, Eva Xu, Liliang Ren, Xinyu Yang, Zeyi Huang, Xuehai He, Luyao Ma, Baolin Peng, Hao Cheng, Pengcheng He, Weizhu Chen, Shuohang Wang, Simon Shaolei Du, and Yelong Shen. Thetaevolve: Test-time learning on open problems, 2025. URL <https://arxiv.org/abs/2511.23473>.

- Marcus Williams, Hannah Sheahan, Cameron Raymond, Tomek Korbak, Deng Pan, Peilin Yang, Leon Maksin, Ningyi Xie, Phillip Guo, Ian Kivlichan, and Micah Carroll. Predicting llm safety before release by simulating deployment, 2026. URL <https://arxiv.org/abs/2607.07184>.
- Xia Yang, Xuanyi Zhang, Hao Hu, and Feng Ji. Beyond accuracy: Evaluating strategy diversity in llm mathematical reasoning, 2026. URL <https://arxiv.org/abs/2605.09292>.
- Dong Zhang. Testing frontier large language models’ physics literacy in parallel physical worlds, 2026. URL <https://arxiv.org/abs/2607.00276>.
- Jiayi Zhang, Simon Yu, Derek Chong, Anthony Sicilia, Michael R. Tomz, Christopher D. Manning, and Weiyan Shi. Verbalized sampling: How to mitigate mode collapse and unlock llm diversity, 2026a. URL <https://arxiv.org/abs/2510.01171>.
- Rui Zhang and Zhichao Lu. Rethinking code similarity for automated algorithm design with llms, 2026. URL <https://arxiv.org/abs/2603.02787>.
- Rui Zhang, Fei Liu, Xi Lin, Zhenkun Wang, Zhichao Lu, and Qingfu Zhang. Understanding the importance of evolutionary search in automated heuristic design with large language models, 2024. URL <https://arxiv.org/abs/2407.10873>.
- Xinhao Zhang, Xi Chen, François Portet, and Maxime Peyrard. What makes an llm a good optimizer? a trajectory analysis of llm-guided evolutionary search, 2026b. URL <https://arxiv.org/abs/2604.19440>.
- Lexin Zhou, Wout Schellaert, Fernando Martínez-Plumed, Yael Moros-Daval, Cèsar Ferri, and José Hernández-Orallo. Larger and more instructable language models become less reliable. *Nature*, 634(8032):61–68, 2024. doi: 10.1038/s41586-024-07930-y.

A Prompts, verbatim

A.1 Bare template (`{n}` parameterized; the $N = 13$ instance hashes to 32db485bea625ff9f39f4723ebf1a01f337559a9e2cf567fb486928f71f7f8df):

```
Pack exactly {n} non-overlapping circles inside the unit square [0,1]x[0,1]
so that the sum of their radii is as large as possible.
MUST hold: exactly {n} circles; every circle fully inside the unit square
[0,1]x[0,1] (x-r >= 0, x+r <= 1, same for y); no two circles overlap
(distance between centers >= sum of the two radii; touching is allowed).
Do not write or execute code - construct the packing by reasoning alone.
Output ONLY the raw Python list of {n} [x, y, r] lists. No explanation, no
code fences, no other text.
```

A.2 trace_v2 template — identical to A.1 through line 3, then:

```
First write one line beginning "METHOD:" naming the construction you used.
Do not use square brackets anywhere in that line.
After the METHOD line, output ONLY the raw Python list of {n} [x, y, r]
lists. No explanation, no code fences, no other text after the list.
```

A.3 Dispatch wrapper (both arms; *not* part of the hashed task prompt, hence an instance of the fragment-hashing limitation in §3.1): Do not use any tools. Your entire final message must be the answer and nothing else.

A.4 Registered prompt hashes. bare $N=13$ 32db485b..., $N=17$ 8437df75..., $N=21$ a415425b..., $N=31$ a664d003..., $N=35$ 3b08c56e..., $N=37$ 2c0a8819...; trace_v2 $N=13$ a920f1c9..., $N=21$ 91205727..., $N=31$ bd490b7b.... Arm CN (held-out N): $N=50$ 0b858558..., $N=58$ 200c0bb4..., $N=62$ b9854a84...,

N=65 5fb5fe fd..., N=75 a5db8fc3..., registered at commit 69a1642. The N=43 bare hash 1208e7d2... was ledger-only until arm P registered it (§5.3), and the pilot prompt was never hashed (§8).

A.5 Parser. Completions are parsed with `ast.literal_eval` after stripping code fences; a completion failing to yield a list of N numeric triples is recorded `literal_eval_failed` and scored invalid, with the failure mode logged. Radii must be strictly positive; containment and pairwise non-overlap are checked at the two registered tolerances.

B Deviations and disclosure table

Table 3: Deviations from preregistration.

Prediction	Registered rule	As registered?	Outcome / reason
P1–P3	exact-value point predictions	yes	Appendix E.3, Supplement S5
P4 (N=37)	converge, 3.0345178	yes	confirmed
P5 (N=35)	truncate, 2.9166667, structural separation	yes	3/4 valid; $n = 4$, “consistent with”
P-S1...P-S4	directional tier comparisons	yes	§3.3
P-O1, P-O2, P-O4	tier comparisons on on-prediction / rival rates	no	declared not evaluable post hoc after a 13% validity collapse; the registered disconfirmation (regression toward the trap) did not occur trivially satisfied on 4/4 valid
P-O3	multi-radius fraction ≥ 0.9 of valid	yes	not confirmed, $p = 0.30$
P-T1	direction at each N and pooled $p < 0.05$	yes	
P-T2	directional, no alpha	yes	confirmed as registered (1/53 vs 2/50); no inferential weight claimed
P-T3	directional, no alpha	yes	met as registered; $p = 0.0325$ uncorrected fails Holm; carried by $N = 13$ (Supplement S6.4)
P-T4	$\geq 90\%$ of scoreable claims match	yes, plus corrected scorer	38/41 (93%) original; 54/56 (96.4%) blind, frozen rubric
P-M1–P-M3	modal output = $V(k^*, m)$ at $N = 20/30/41$	yes	disconfirmed at all three cells — modal outputs are (k^*+1) -grid values (Appendix E.4)
P-M4	$N = 57$ modal = $T(8, 57)$, rival count 0	yes	confirmed (6/10 on-prediction, 0/10 rival)
F-M1 (arm M falsifier, filler branch)	modal $\neq V(k^*, m)$ at 2+ of 3 converge cells, ties count as NOT	yes, tie-stated in registration	triggered at 3 of 3 ; $V(k, m)$ restricted to $m \leq 1$ (Appendix E.4)
F-M2 (arm M falsifier, truncate branch)	$N = 57$ modal $\neq T(8, 57)$	yes	not triggered
Arm M rejection rule	runtime deaths excluded, counted	yes (registered for this arm)	3 $N = 41$ invocations excluded, cell reported as 12 of 15
GM3 serving path	first-party API only	amendment registered before resumption (arm_gm3_amendment1_openrouter.md, 67cc4a1)	41/140 rows via routed path, provider logged; 59.0% / 57.5% dual accounting, no verdict change (§5.4)
Arm V proposer set	3 opencode aliases	expanded 3 $\rightarrow$ 13 by amendments 1–2, each registered before the sampling it governs	verdicts in §5.4; big-pickle disclosed as vendor-unattested in amendment 1
Arm V transport repair	single pass, <code>max_tokens</code> 8192	amendment 3 registered before resumption (mechanical: 24,576 budget, retry, censored-row accounting)	877 attempt rows over the 325 registered slots (314 carry rows); floor evaluated per slot (§5.4)
GM3 rival-structure check	—	post hoc , disclosed (diagnostics_gm3_rival.py)	27/27 rival-valued discriminating-cell packings carry the rival’s two-radius structure (§5.4); the registered signature is scoped to k^* and cannot see (k^*-1) structure
Falsifier	<code>trace_v2</code> validity $\leq$ bare at 2+ of 3 N	both readings reported	triggered under the registered tie-inclusive reading (2/3 cells); not triggered under the strict $<$ the code implemented (Supplement S6.3)
Ladder inclusion	—	post hoc	<code>opus_alias</code> initially excluded, later included after results known; ladder reported both ways (§4.2)
Cap exclusion	—	not registered	five records excluded; both rates reported (Appendix E.3)
Arm CN ledger transcription	verbatim rows	disclosed (arm_cn_results.txt)	eleven $N = 65$ rows reconstructed from the filler triple, flagged per row; discarded, P-CN1 reads 3 of 4 (Appendix E.5)

continued on next page

Prediction	Registered rule	As registered?	Outcome / reason
Arm M out-of-family sweep	—	post hoc , not preregistered	three converge-cell rows ($N = 20$ twice, $N = 30$) above the family best by $1.5\text{--}2.0 \times 10^{-2}$; not modal (Appendix E.4)
Arm CC2 tool use	no tools (wrapper A.3)	disclosed (arm_cc2_results.txt)	one $N = 13$ invocation used tools; scored as emitted (arm-M convention)
Uniform-template null	—	post hoc , not preregistered	discriminating-cell rates exceed the 1/3 null; binomial tails 0.043 to 2.9×10^{-4} (Appendix E.2) (<code>diagnostics_template_null.py</code>)
Arm CP (perturbed container)	P-CP1/P-CP2 + S-CP1/S-CP2 + F-CP1	registered before sampling (105e8b7); drafted at review stage, disclosed	P-CP1 confirmed (modal 4/5, $N = 75$ a registered tie against); S-CP1 not met (4/61 rival); S-CP2 5/5; F-CP1 not triggered (Appendix E.6)
Arm RP (direct recall)	P-RP1/P-RP2 + F-RP1	registered before sampling (105e8b7); draft cell $N=32 \rightarrow N=30$, documented	P-RP1 confirmed : 0/44 recalls, 87/87 scored responses UNKNOWN; F-RP1 not triggered (Appendix E.6)
Arm RP positive control	C-RP1 vs C-RP2	amendment 1 registered before sampling (15555eb), responding to a review finding	C-RP1: 14/15 state 0.5 at $N=1$; abstention-prior alternative rejected (Appendix E.6)
Arm PP (paraphrase)	P-PP1/P-PP2 + S-PP1/S-PP2	registered before sampling (2e89614); paraphrases authored results-aware, disclosed	P-PP1 confirmed 6/6 modal; S-PP1 not met (6/75 rival); S-PP2 5/6 (Appendix E.6)
Arm L (iterated loop)	L1–L4	registered before sampling (c9645fe), scored once (750bb32); one pooled bar per cell rather than four independent bars, a design choice that cost the arm power and is not repaired after the fact	L2 confirmed ; L1 not met on 13-DIVERSE, a 2–2 tie at $n = 4$ decided against the prediction by the registered tie rule; L3 not met and reversed , the prediction withdrawn rather than reframed; L4 open and answered — 0 of 49 valid conditioned outputs clear the family argmax (§5.2)
Arm P (pinned path)	P-P1–P-P4 + F-P1	registered before sampling (aa473ad); post-review motivation disclosed	validity collapsed to 0/105 at the square cells; direct-emission predictions not evaluable under the registered floor; P-P3 holds (0/12); F-P1 denominator empty, reported as such (§5.3)
Arm P temperature diagnostic	—	post hoc , not preregistered	0/6 valid at temperature 0.0, 0.5 and 1.0; temperature excluded as cause (§5.3)
Arm P runtime control	—	post hoc , not preregistered	same prompt through the agent runtime, same day: 6/6 valid, anchor 4/6 (§5.3)
Arm P-D (runtime component)	P-PD1/P-PD2/P-PD3 + S-PD1	registered before sampling (5554cc0); run interrupted by credit at 27 of 30 rows; closed there by amendment 1 (<code>arm_pd_amendment1_closed_at_27.md</code>), since no completion could change the verdict	P-PD1 holds and the last three rows could not change it: thinking on 12/14 valid, modal $T(4, 13)$; system line 0/13; S-PD1 met (§5.3)

continued on next page

Prediction	Registered rule		As registered?	Outcome / reason
Arm B (optimizer alone)	P-B1/P-B2 + S-B1		registered before the first scored run (1932cc0); post-review motivation disclosed; amendment 1 (c486843): version 1 imported time and sys, blocked 45/45 by the registered allowlist, run kept	P-B1 holds : 45/45 clear, 3/3 cells at the 20% bar; S-B1: best exceeds the Sonnet best at 3/3 cells; contribution 1 rewritten per the registered rule (§3.5)
Arm CL (library code)	P-CL1/P-CL2 + F-CL1		registered before sampling (aa473ad); one execution-mechanism correction made and disclosed in the same commit	weak tier P-CL1 holds (0/11); Sonnet tier P-CL2 holds, F-CL1 fired (23/25, 3/3 cells); registered integration rule applied (§3.5)
Arm CL lenient reparse	—		post hoc , not preregistered	the 22 weak-tier parse failures are all <code>numpy</code> scalar reprs; stripping the wrapper recovers 9 valid programs (3, 4, 2 by cell) and 1 clearance at $N = 13$ (1.812531608); registered 0 of 11 stands, 1 of 20 reported beside it (§3.5) (cl_lenient_reparse.py)
Arm CCP (isolating arm)	P-CCP1/P-CCP2 + F-CCP1		registered before sampling (ebf213e); post-review motivation disclosed; one 402 interruption at row 24, resumed with one worker, request unchanged (32fd3c7)	P-CCP1 holds : 0/35 clear, none at the argmax; F-CCP1 not triggered; registered integration rule applied (§3.5)
Arm CL-W (weak-tier top-up)	P-CLW1/P-CLW2 + F-CLW1 + S-CLW1		registered before sampling (f73222a); pooling rule and the lenient secondary reading fixed in advance; 45 rather than the drafted 90, credit-limited, disclosed	P-CLW1 holds : 0/19 pooled clear (Wilson upper 16.8%); S-CLW1 4/41 lenient, dead zone; F-CLW1 not triggered (§3.5)
Wave stratification of the bare arm	—		post hoc , not preregistered	the trace effect is carried by $N = 13$ and confounded with collection wave; a diagnosis, no confirmatory weight (Supplement S6.4)
Arm MU clearance read	—		post hoc , not preregistered	the clearance question was never registered for arm MU; 63 valid at 10^{-9} , none above the argmax (§5.1) (arm_mu_ceiling.py)
Arm CH attempt-level reading	—		post hoc , not preregistered	36/45 reach for the argmax against the bare arm's 5/57; the registered per-cell metric is reported first (§4.1) (arm_f_attempt_uncond.py)
Structural k back-out	—		post hoc , not preregistered	50/50 on-prediction samples k^* -structured, 64/69 overall (Appendix E.2) (diagnostics_kmatch.py)
Ambition diagnostic	—		post hoc , not preregistered	distinct radii and off-lattice fraction rise monotonically across tiers (§4.2) (diagnostics_ambition.py)

C The opus_alias arm: the ladder and the failure taxonomy

The tier ladder that §4.2 quotes is Table 4; the two post-hoc decisions behind it and the two corpora behind its first rung are Supplement S3.

‡ Addressed through a bare serving alias; the alias-to-weights binding is not attestable, and this row is not a statement about any released model version. Tier denominators are not matched (the weak tier spans seven N against three), which is why the matched-cell row is given separately. Sources: `arm_f_candidates_v2.jsonl`, the three preregistration files.

Table 4: Three-attractor ladder. Valid columns count samples passing containment and non-overlap at the 10^{-6} and 10^{-9} tolerances; On-pred and Rival count valid samples landing on the predicted value and on the rival argmax. The matched-cell row is the comparison to quote across tiers ($83\% \rightarrow 100\% \rightarrow 13\%$).

Tier	Cells	n	Attempted ily	fam- ily	Valid 10^{-6}	Valid 10^{-9}	On-pred	Rival	Failure mode
haiku (bare)	13, 17, 21, 31, 35, 37, 43	45	uniform truncated; when underfull	grid, fillers	35 (78%)	29 (64%)	12/16 (3 matched cells)	2	overlap 6, outside 1, parse 3
haiku, matched cells	13, 21, 31	60	as above		50 (83%)	—	35/50	2	overlap 5, outside 2, parse 3
sonnet	13, 21, 31	30	perturbed hexagonal 2–3 radii	grid, rows,	30 (100%)	27 (90%)	1/30	5 (structure)	none
opus_alias ‡	13, 21, 31	30	quarter-circle corners, Apollonius fillers, coarse grid + border strips		4 (13%)	4 (13%)	0/4	0	overlap 24, non-positive radius 2

The top tier attempts recursive constructions and fails to build them. The third arm was invoked through a bare tier alias. We name it `opus_alias` throughout and make no claim about which weights served it: the runtime accepts only the bare alias and exposes no dated identifier, so the binding is a vendor promise rather than an attestable fact. No version number appears anywhere in this paper.

That caveat matters: completion-time and token-count anomalies were observed in the unreleased session transcript for this arm, but the ledger carries no latency or token field, so they appear in no released artifact and nothing below relies on them.

Under those caveats the arm — preregistered fully blind in `arm_o_preregistration.txt` — was valid in 4 of 30 invocations (13%) at both tolerances. At every cell it attempted a more ambitious construction than either other tier: quarter-circle corners at $r = 0.25$ with Apollonius-style fillers at $N = 13$ (10/10 samples), mixed-radius 4×4 -ish grids at $N = 21$, and a coarse 3×3 grid at $r \approx 1/6$ with border strips at $N = 31$ (0/10 valid). Failures are geometric rather than numerical — edge strips at $r = 0.03$ placed 0.138 from an $r = 1/6$ grid circle needing 0.197 — and twice the arm padded to count with zero-radius circles. Valid samples score *below* the trap they were expected to fall into (1.26–1.41 at $N = 13$ against 1.625). A post-hoc diagnostic (not preregistered; labeled as such in the repository) finds no tolerance artifact behind the 13%: 24 of 26 invalid samples overlap grossly (median maximum overlap 3.3×10^{-2} ; radius-shrink repair costs a median 15% of sum-of-radii), while tolerance-scale near-misses ($< 2.5 \times 10^{-5}$) fall instead in 5 of 7 geometry-scored weak-tier failures — the exact-tangency grids, not the ambitious constructions.

D Provenance, in full

Section 8 summarises this; here it is in the form the body carried before the structural cut, so that a reader auditing ancestry finds every registration, commit, ledger and script in one place.

- **Registration.** Prompts were hashed with SHA-256 and the hashes recorded before sampling, with the predicted values they were to be tested against: `arm_f_repro.py` (header, P1–P5, square arm), `arm_s_preregistration.txt`, `arm_o_preregistration.txt`, `arm_t_preregistration.txt` (prereg hash `ab7900a8...`), `arm_m_preregistration.txt` with `arm_m_prompts.json` (committed `e528c6b` before sampling), and `arm_mu_preregistration.txt` with `arm_mu_prompts.json` (committed `2b7d202` before sampling; decision thresholds, power analysis, tie conventions and falsifier F-MU1 all registered), and the cross-vendor chain `arm_gm_preregistration.txt` (commit `37b3adb`), `arm_gm2_preregistration.txt` (`3019aab`), `arm_gm3_preregistration.txt` (`c0b5b97`) with its serving-path deviation registered in `arm_gm3_amendment1_openrouter.md` (`67cc4a1`) before

the affected cells were sampled, and `arm_v_preregistration.md` (commit 5444f10, ancestry-checked by its runners before any sampling) with its three amendments, each committed before the sampling it governs. Arm CC's registration (`arm_cc_preregistration.txt` with frozen prompts, executor and smoke test) was committed and pushed to the public remote at 285deca before its first invocation; arms CC2 and CCS were registered together (`arm_cc2_preregistration.txt`, commit 4d7b7c5, pushed before sampling) after CC's results were known, with that motivation disclosed in the registration itself. Arm CN (`arm_cn_preregistration.txt`, commit 69a1642, pushed before sampling) was registered after the rest of the manuscript was complete, in response to a pre-submission check; its motivation is post-hoc with respect to the paper's results and is disclosed in the registration, and its predictions were fixed before its first invocation. Arms CP and RP (`arm_cp_preregistration.txt`, `arm_rp_preregistration.txt`, each with its prompt-hash JSON and analysis script, commit 105e8b7, pushed before the first invocation) were drafted at review stage and registered together; RP's one draft-to-registration cell change ($N = 32 \rightarrow 30$) is documented in its registration, as is its positive-control amendment (15555eb). Arm PP (`arm_pp_preregistration.txt`, prompt hashes and scorer committed, 2e89614, pushed before the first invocation) discloses that its paraphrases were authored results-aware. The rectangle transfer was registered before collection and never refitted. Arms P and CL (`arm_p_preregistration.txt`, `arm_cl_preregistration.txt`, builders, frozen prompts, runners and analysis scripts) were committed and pushed together at aa473ad before either was sampled; ledgers `arm_p_collect.jsonl` and `arm_cl_collect.jsonl`, frozen reports `arm_p_report.json` and `arm_cl_report.json`, and the two post-hoc diagnostics `arm_p_diag_temperature.jsonl` and `arm_p_diag_runtime.jsonl`, each labelled post hoc in its header. On multiplicity: this study now spans many registered arms, and no correction across arms is applied — each arm's predictions were registered with its own bars and falsifiers before its sampling, every arm is reported regardless of outcome (three GM-chain instruments failed or stalled before one succeeded, and Appendix E.4's falsifier fired), and no arm was selected for inclusion on its result. The protection against garden-of-forking-paths here is registration and exhaustive reporting, not a family-wise error rate; a reader who wants one can count the twenty-six arms from §3's table. That defence covers the registered predictions and not the post-hoc diagnostics: which unregistered analyses we chose to run is a degree of freedom registration does not constrain. Table 3 lists them, and three cut against us — Appendix E.4's out-of-family sweep, §5.4's rival-structure check, and Supplement S6.4's wave stratification.

- **Four limits on that claim.** (i) The digests cover the task prompt only; the subagent inherits instruction files outside it, so what is locked is a *fragment* of the conditioning context, a replicator cannot reconstruct those files, and we cannot establish they were identical across the arm-F and arm-T windows — the boundary Supplement S6.4 identifies as a confound. (ii) `arm_f_prompts.json` covers $N = 13, 17, 31, 35$ and 37 ; the $N = 21$ bare hash is registered in `arm_t_preregistration.txt`; the $N = 43$ bare hash appeared in no preregistration and no prompts file until arm P's, existing only in the ledger — it recomputes exactly from the bare template, but recomputation after the fact is not registration. (iii) **No hash was ever registered for the pilot prompt**, whose drift the preregistration describes only in prose; the ten pilot rows carry `prompt_sha256: null` in the corrected ledger with an explanatory note. (iv) The rectangle ledger (`arm_g_candidates.jsonl`) carries no prompt hashes, no run dates and no proposer fields at all.
- **Ledger metadata correction.** Three internal-check findings established that the shipped ledger's provenance fields contradicted the paper: trace rows carried the *bare* prompt hash for their cell, `run_date` was uniformly 2026-07-30 on all 215 rows, and the Sonnet and `opus_alias` rows were stamped with the Haiku alias and dated id. All three are corrected in `arm_f_candidates_v2.jsonl`; the correction is metadata only — every scored quantity, geometry, raw output, arm label and sample id is byte-identical to v1, verified field by field, each corrected field carrying a sibling `*_v1` field. `arm_f_candidates.jsonl` is unmodified (sha256 02ecabcd...); v2 is 9f845f78.... Correction rules, sources, wave-boundary arithmetic and caveats — including that `run_date` has day granularity only — are in `ledger_v2_corrections.md`. Analysis scripts read v2.

- **Excluded and unreleased records.** The five concurrency-cap rejections appear in the ledger in no form, no `excluded_reason` field exists, and the exclusion rule was not preregistered; both rates are reported in Appendix E.3. The `opus_alias` latency and token-count anomalies (§4.2) come from the runtime session transcript; no latency or token field exists in any released artifact.
- **Analysis artifacts.** Raw outputs are stored verbatim (`arm_f_raw.json`, `arm_f_candidates.jsonl`, `arm_f_candidates_v2.jsonl`, `arm_g_candidates.jsonl`, `arm_m_collect.jsonl` scored by `arm_m_analysis.py`, `arm_mu_collect.jsonl` scored by `arm_mu_analysis.py` with frozen output `arm_mu_results.txt`, and the GM3 ledger `arm_gm_gm3_checkpoint.jsonl` — every row’s raw API response verbatim, serving path and provider tagged per row — scored by `arm_gm3_analysis.py` into `arm_gm3_report.json` with both serving-path accountings) and scoring is deterministic and local. Value claims are gated on an independent LP oracle rather than on the closed form being tested: `n_sweep_forecast.py` with `verify_against_lp` checks 83 configurations to within 10^{-9} , and the same oracle produced Supplement S5.2’s negative result. For arm T (Supplement S6): the faithfulness rubric (`p7_faithfulness_rubric.md`) was frozen by git commit `e181d2a` (2026-08-01 03:56:56) before any rescoring, and all 60 blind labels with justifications are in `p7_blind_labels.json`. `arm_t_analysis.py`, `n_sweep_forecast.json`, `rect_forecast.json` (arms T and G, Supplement S6 and S5), `p11_mode_baseline.json` and `STATE.md` are included; every table regenerates from raw outputs by a command documented in `HOW_TO_RUN.md`.
- **Internal checks.** Before submission the manuscript passed an internal adversarial check: language models of other families (`deepseek-v4-pro`, `deepseek-v4-flash`, Gemini), under written protocols, independently recomputed the paper’s statistics from the released ledger; thirty-two claims were corrected as a result, itemized in the supplementary `corrections_ledger.md`. This is quality assurance, not peer review, and carries no external authority.

D.1 Per-arm ledgers and scripts

The sentence each arm’s subsection used to end with, verbatim, under the subsection’s title.

- **Arms MU and CH: no transfer to one-parent calls; the argmax in hand does not dislodge the template** Raw rows verbatim in `arm_mu_collect.jsonl` (180 of 180 launched invocations collected across arms MU and CH, no runtime deaths); scoring in `arm_mu_analysis.py`, committed before sampling completed; frozen output in `arm_mu_results.txt`. This arm brings the study corpus to 468 invocations.
- **Arm CC: an executed program does not escape the family** Raw rows verbatim in `arm_cc_collect.jsonl` (45 of 45 launched invocations collected, no runtime deaths); scoring `arm_cc_analysis.py`; frozen report `arm_cc_report.json` and summary `arm_cc_results.txt`.
- **Arm P: the pinned path** Registration `arm_p_preregistration.txt`, prompts and hashes `arm_p_prompts.json`, runner `arm_p_run.py`, scorer `arm_p_analysis.py`, ledger `arm_p_collect.jsonl` (225 rows, six refused by the serving path), report `arm_p_report.json`; the two post-hoc diagnostics in `arm_p_diag_temperature.jsonl` and `arm_p_diag_runtime.jsonl`, labelled post hoc in their headers. Arm P-D: registration `arm_pd_preregistration.txt` (commit `5554cc0`), builder `arm_pd_build.py` (arm P’s $N = 13$ prompt, hash asserted), spec `arm_pd_prompts.json`, runner `arm_pd_run.py`, scorer `arm_pd_analysis.py` (arm P’s scorer imported unmodified), ledger `arm_pd_collect.jsonl` (27 rows plus the 402 rows kept for resume), report `arm_pd_report.json` marked INTERIM.
- **Arm CL: the library channel** Registration `arm_cl_preregistration.txt`, builder `arm_cl_build.py` (one clause replaced in arm CC’s prompt, hash asserted), prompts `arm_cl_prompts.json`, runner `arm_cl_run.py`, scorer `arm_cl_analysis.py` (driver and numpy/scipy versions recorded in the report), ledger `arm_cl_collect.jsonl` (90 rows), report `arm_cl_report.json`. The post-hoc re-execution and lenient reparsing are `recount_cl.py` and `cl_lenient_reparse.py` in the paper repository’s `loop/` directory, outputs `cl_recount.json` and `cl_lenient.json`.

- **Arm CCP: the isolating arm** Registration `arm_ccp_preregistration.txt` (commit `ebf213e`), builder `arm_ccp_build.py` (arm CC’s prompts byte-identical, hashes asserted), prompts `arm_ccp_prompts.json`, runner `arm_ccp_run.py` (the one-worker amendment after the row-24 interruption, `32fd3c7`), scorer `arm_ccp_analysis.py` (arm CL’s pipeline with the allowlist set to `math`), ledger `arm_ccp_collect.jsonl` (45 rows plus the one 402 row, both segments timestamped), report `arm_ccp_report.json`.
- **Arm CL-W: the top-up** Registration `arm_clw_preregistration.txt` (commit `f73222a`), builder `arm_clw_build.py` (arm CL’s weak-tier prompts byte-identical, hashes asserted), prompts `arm_clw_prompts.json`, runner `arm_clw_run.py`, scorer `arm_clw_analysis.py` (arm CL’s pipeline over the pooled 90 rows, both readings per row), ledger `arm_clw_collect.jsonl` (45 rows), report `arm_clw_report.json` with every row’s registered and lenient result.
- **Arm B: the optimizer alone** Registration `arm_b_preregistration.txt` (commit `1932cc0`), amendment `arm_b_amendment1_fixed_restarts.md` (`c486843`), program `arm_b_baseline.py` (sha asserted by the runner), runner `arm_b_run.py` (arm CL’s pipeline, no model), ledger `arm_b_collect.jsonl` (45 substituted sources), report `arm_b_report.json`; the blocked version-1 run is kept as `arm_b_v1_blocked_collect.jsonl` and its report and log. No model invocation; the study corpus stays at 1500.
- **Arms CC2 and CCS: the replication and the tier probe** Ledgers `arm_cc2_collect.jsonl` and `arm_ccs_collect.jsonl` (45/45 each, no runtime deaths); runner `arm_cc2_analysis.py` reuses CC’s registered pipeline unmodified. These arms bring the study corpus to 603 invocations.

E The mode prediction and the arms that bound it

The closed form that names the weak tier’s modal output in advance, reported in full; the body keeps only what Table 2 and §4 need from it.

E.1 Trap zones and the bound table

Substituting $k^*(N)$ into the branch condition: TRAP $N \in [k^2 - k + 1, k^2 - 1]$, CONVERGE $N \in [k^2, k^2 + k]$. Over $N = 10 \dots 60$: `[13, 15]`, `[21, 24]`, `[31, 35]`, `[43, 48]` and `[57, 63]` clipped to `[57, 60]` by the sweep bound.

Inside a zone the rule generally loses value against the best construction available *within the recipe family*. Worst-in-zone penalties fall as N grows: 8.51% at $N = 13$, 7.03% at 21, 6.01% at 31, 5.25% at 43, 4.66% at 57. The penalty reaches zero at $N = 35$, 48, 14 and 15, but **not** at $N = 24$, where truncation gives $T(5, 24) = 2.4000000$ against $V(4, 8) = 2.4142136$, a residual 0.59%. Branch label and penalty are therefore not coextensive: 4 of the 22 trap-branch N in range cost nothing, so “trap” names the branch, not a guaranteed loss. Where the penalty is zero, truncation is separable from convergence only by structure; Appendix E.3 uses $N = 35$ as that control.

The linear-programming verification summarized in §2.2 originally covered $k = 2 \dots 7$ and stopped short of two registered cells, arm M’s $T(8, 57)$ and arm CN’s predictions at $k = 8$ and 9; the sweep was extended rather than the claim narrowed (`lp_extend_k89.py`), and every registered closed-form value is inside the verified range.

Our bound table (`n_sweep_forecast.json`, transcribed from Friedman (n.d.)) covers $N = 10 \dots 40$; through the round before this one it stopped at $N = 30$, although the source page has listed $N = 31$ since 2011 and, as fetched on 2026-09-02, runs to $N = 40$ (corrections ledger, row 35). Its $N = 26$ entry, 2.63598, matches ShinkaEvolve’s figure truncated. That entry has since been re-credited on the source page to a July 2026 submission, so we withdraw the reading that LLM-driven systems hold the record here and claim only what the table supports. What survives is that the recipe family is never competitive with that record: the family’s best value sits 0.02–0.16 below it across $N = 10 \dots 40$ (the rule’s own value 0.02–0.31), and 0.0946 at $N = 26$ ($2.63598 - 2.5414214 = 0.0945586$). Above $N = 40$ there is no bound table, so both the deficit claim and the LP gate’s “never exceeds a published bound” abort are unchecked there, covering two of five trap zones and one of our seven square cells ($N = 43$).




Figure 2: *Predicted versus optimal value*, $N = 10 \dots 60$. (i) the rule’s prediction; (ii) the best value in the recipe family; (iii, dotted) published best-known values (Friedman, n.d.), terminating at $N = 40$ because the source page stops there. Shaded bands are the trap zones, the last being $[57, 63]$ clipped by `sweep(10, 60)`. Worst-in-zone penalties 8.51%, 7.03%, 6.01%, 5.25%, 4.66%; these are family-internal arithmetic and exact at every N , while the published-bound comparison (curve iii) is checkable only up to $N = 40$, where the bound table stops. The gap closes to zero at $N = 35$ and $N = 48$ but **not** at $N = 24$, where 0.59% remains — curve (ii) sits strictly above (i) across $[21, 24]$.

E.2 The mode prediction: seven cells, all modal

An exact point prediction invites a fair objection: a per-sample hit rate well short of 100% is not “predicting the exact output”. The number should be read against the sampling entropy of the distribution, computed in `p11_mode_baseline.json` (bare arm, 2×10^{-3} buckets). At **all seven tested N the predicted value is the empirical modal output** — four discriminating cells plus three where the prediction and the family argmax coincide, so the anchoring-specific evidence sits in the four — and the on-prediction rate equals the modal frequency to within one sample. The two columns are counted on different partitions: the modal-frequency column groups at four decimals, the on-prediction column uses the registered 2×10^{-3} window. They come apart at exactly one cell, $N = 31$, where the window admits one packing the four-decimal bucket separates out, which is why that row reads 12/17 against 13/17:

N	predicted	empirical mode	modal freq	on-prediction / valid	k^* -structure / valid [†]
13	1.6250000	1.6250	10/18	10/18 (56%)	17/18
17	2.0517767	2.0518	3/4	3/4 (75%)	3/4
21	2.1000000	2.1000	12/15	12/15 (80%)	13/15
31	2.5833333	2.5833	12/17	13/17 (76%)	17/17
35	2.9166667	2.9167	3/4	3/4 (75%)	3/4
37	3.0345178	3.0345	3/4	3/4 (75%)	4/4
43	3.0714286	3.0714	6/7	6/7 (86%)	7/7

[†] Post-hoc structural check, not preregistered, run after the square arm closed (`diagnostics_kmatch.py`): the grid order backed out of each sample’s dominant radius ($k = \text{round}(1/2r)$) equals $k^*(N)$. All 50 of 50 on-prediction samples are k^* -structured — the value match is not an accounting coincidence — and 64 of 69

valid samples overall (93%) sit on the k^* grid, so most value misses are k^* -grid variants (perturbed radii or filler counts), not different constructions.

No predictor of a single value can exceed the modal frequency, and this one attains it at 7 of 7 cells: the residual is dispersion, not misprediction.

The modal-frequency argument holds only if the mode and the on-prediction count are read on the same partition, so the two partitions in play have to be separated. On-prediction is the registered 2×10^{-3} value window. The modal-bucket column is grouped at four decimals, a reporting choice, and the two differ at exactly one cell: at $N = 31$ on-prediction is 13 of 17 against a 4-dp modal bucket of 12 of 17. Read on the registered window — the partition the modal-frequency claim is about — the mode is 13 and the predictor attains it. The 4-dp column is kept because it is what the per-cell tables display, not because any verdict rests on it. Three cells ($N = 17, 35, 37$), all non-discriminating, carry four valid samples each, where a 3/4 mode is thin evidence on its own; the well-populated cells of the held-out probe (Appendix E.5) are the stronger test of the same rule. As a floor, a naive “the model emits a round number” baseline hits 2 of the same 69 valid samples (3%). A stronger baseline, disclosed post hoc: a proposer choosing uniformly among the family’s plausible template shapes — one branch value per order $k \in \{k^*-1, k^*, k^*+1\}$ — lands on-prediction 1/3 of the time; the observed discriminating-cell rates sit above that null with exact binomial upper tails $0.043, 2.9 \times 10^{-4}, 3.4 \times 10^{-4}$ and 6.9×10^{-3} at $N = 13, 21, 31$ and 43 (`diagnostics_template_null.py`). All four survive Holm correction at $m = 4$, whose thresholds are 0.0125, 0.0167, 0.025 and 0.05 — including $N = 13$, the weakest, which an earlier draft of this section reported as uncorrected. The 7-of-7 mode identity does not depend on this null. Two scope conditions: this is the bare weak-tier arm — pooled over the four discriminating square cells (full ledger), the two higher tiers and the rectangle arm (arm G, Supplement S5), the on-prediction rate is 46% (47/102, pooling the counts rather than the rates: 41 + 1 + 0 + 5 on-prediction over 57 + 30 + 4 + 11 valid; the three non-discriminating square cells, 9/12, are excluded because prediction equals the family argmax there), because §3.3’s higher tiers are almost never on-prediction — and “modal” is a property of the observed sampling regime, not of a pinned decoding configuration (§7).

Which on-prediction number is which. The rate is reported on several corpora and cell sets, each for a reason; every one is tabulated here against its own denominator, so that no later mention has to be read twice.

figure	counts	corpus	cells
56–86%	on-prediction	bare weak-tier arm	each of the seven square cells
18/23	on-prediction	original corpus	the four discriminating square cells
41/57	on-prediction	full shipped ledger	the same four
2/23, 2/57	rival emissions	those two corpora	the same four
9/12	on-prediction	full shipped ledger	the three non-discriminating square cells
12/16	on-prediction	original 45-row arm	the three matched cells (13, 21, 31)
35/50	on-prediction	full matched-cell ledger	the same three
47/102	on-prediction	full ledger plus both higher tiers and the rectangle arm (Supplement S5)	the four discriminating square cells, plus those tiers and that container

The headline figure for the anchoring claim is **41/57**: the full shipped ledger at the four cells where prediction and family argmax differ. Everything else is a subset, a superset, or a different question.

One label belongs on that figure at the point of use rather than only in §3.1. Of its four cells, two — $N = 21$ and $N = 43$ — carry no individually registered point prediction; their values follow deterministically from a formula fixed before sampling, but a reader auditing what was locked should know which is which. Dropping $N = 43$, whose prompt hash is also unregistered (§8), gives 35/50. Restricting to the two cells carrying both a registered hash and a registered prediction, $N = 13$ and $N = 31$, gives **23/35**. The rate degrades from 72% to 70% to 66% across those three readings, so nothing turns on the choice; we quote 41/57 and state 23/35 beside it.

E.3 Square container, original corpus and full ledger

The square arm sampled weak-tier proposers at $N = 13, 17, 21, 31, 35, 37$ and 43 . The rule was fitted on $N = 23/26/27$ only, so every cell is out of sample. Four cells are *discriminating*: rule and family optimum disagree, so a proposer that searched the family would return a higher number. Across those cells ($N = 13, 21, 31, 43$), **18 of 23 valid invocations landed on the predicted construction** and the rival-argmax value was reached **2 times in 23** — both at $N = 21$, both the 4×4 grid plus five fillers at 2.2588835.

Those figures cover the **original 45-invocation bare arm only** (ids ≤ 5 at $N = 13/31$, ≤ 10 at $N = 21$, all rows at $N = 17/35/37/43$), the corpus that existed when §3 was written. The trace study (Supplement S6) later added 40 bare rows at $N = 13/21/31$ under the same arm label; over the full shipped ledger the same four cells give **41/57 on-prediction and 2/57 rival**. Both are reported. §3 and the trace study are computed on nested subsets of one arm, not on independent samples, and Supplement S5 carries the split. The 40 added rows are the bare arm of the trace study (Supplement S6): 60 bare samples were collected there, of which 20 are the pre-existing arm-F rows — $N = 13$ ids 1–5, $N = 21$ ids 1–10, $N = 31$ ids 1–5 — and 40 are new, so the 57 valid rows behind 41/57 are the original arm plus those 40 and nothing else.

That the anchor first breaks at $N = 21$ is suggestive — it is the bottom of a zone carrying a 7.03% penalty, so the anchor looks weak where obeying it costs — but that is a hypothesis, not a finding. **P4** registered $N = 37$ as converging on 3.0345178, a prime predicted clean, contradicting the prior work’s guess; it confirmed. **P5** registered $N = 35$ at 2.9166667, the top-of-zone control; three of four valid samples landed there, separable from convergence only because the structure classifier reads radii rather than totals; the fourth used a 7×7 lattice truncated to 35 (sum 2.5), outside the rule. At $n = 4$ this is “consistent with” rather than confirmed, and deserves running to $n = 20$.

Bookkeeping. Five invocations were rejected by the runtime’s 20-subagent concurrency cap *before reaching a model*. We excluded them under a rule that was not preregistered, so the arm is reported both ways — $35/45 = 77.8\%$, or $35/50 = 70.0\%$ had the rejections been scored as model failures — and the five records appear in the ledger in no form (§8). The arm also contains three parse failures, and one $N = 37$ proposer derived $r = (\sqrt{2} - 1)/12$ correctly in prose, then transcribed 0.03571429.

E.4 Arm M: the filler branch fails at $m > 1$; the fifth k confirms

Arm M was registered before sampling to close two gaps: $V(k, m)$ ’s filler branch was supported only at $m \in \{0, 1\}$, and the $k = 8$ trap zone had never been sampled. $n = 15$ bare weak-tier invocations at each of $N = 20$ ($k^* = 4, m = 4$, predicted $V(4, 4) = 2.2071068$), $N = 30$ ($k^* = 5, m = 5$, $V(5, 5) = 2.7071068$), $N = 41$ ($k^* = 6, m = 5$, $V(6, 5) = 3.1725890$) and $N = 57$ ($k^* = 8$, trap, $T(8, 57) = 3.5625000$, rival $V(7, 8) = 3.7366935$), with a tie-stated falsifier for each half.

F-M1 triggered at 3 of 3 cells; P-M4 confirmed. Not one of the three converge cells has $V(k^*, m)$ as its modal output: the model moves *up* to the (k^*+1) grid and truncates or exactly fills it, and exactly one sample in 36 valid converge-cell rows emitted the registered construction. Per the registered wording, $V(k, m)$ is restricted to the $m \leq 1$ support previously observed and the extend arm is disconfirmed at $m \geq 4$. At $N = 57$ the modal valid output is $T(8, 57)$ (6 of 10 valid), 0 of 10 reach the rival, and the truncate arm has five confirmed k (4, 5, 6, 7, 8). Validity by cell: 15/15, 14/15, 7/12 and 10/15 at 10^{-6} ; three $N = 41$ invocations died in the runtime before reaching a model and are excluded under the rejection rule registered for this arm.

Table M — Arm M, forecast versus outcome (all four registered cells).

N	k^*	Branch	Registered prediction	predic- tion	Modal valid output	Freq	Valid / sampled	Verdict
20	4	extend, $m = 4$	$V(4, 4)$ 2.2071068	$= T(5, 20)$ $(5 \times 4 \text{ exact})$	$= 2.0000000$	12/15	15/15	P-M1 discon- firmed
30	5	extend, $m = 5$	$V(5, 5)$ 2.7071068	$= T(6, 30)$ $(6 \times 5 \text{ exact})$	$= 2.5000000$	11/14	14/15	P-M2 discon- firmed
41	6	extend, $m = 5$	$V(6, 5)$ 3.1725890	$= T(7, 41)$ $(7 \times 6 - 1)$	$= 2.9285714$	6/7	7/12	P-M3 discon- firmed
57	8	truncate	$T(8, 57)$ 3.5625000, rival count 0	$= T(8, 57)$ ri-	$= 3.5625000$	6/10	10/15	P-M4 confirmed; rival 0/10

Two $N = 20$ rows clear the family under §2.4’s rule: staggered rows of four at $r = 1/9$, sum 2.2222222, against $V(4, 4) = 2.2071068$, both valid at 10^{-9} ; a third converge-cell row at $N = 30$ ($r = 1/11$, sum 2.7272727) is valid at 10^{-6} only and does not clear. All three are hexagonal layouts the family does not contain, found by a post-hoc sweep disclosed as such, not modal at either cell. The per-cell decomposition, the arithmetic reading of the (k^*+1) rectangles and the bookkeeping are Supplement S4. Regenerates from `arm_m_collect.jsonl` via `arm_m_analysis.py`.

E.5 Arm CN: the rule predicts the mode where there is nothing to recall

The rule was fitted on $N = 23/26/27$, with $N = 26$ the canonical published cell, so a rival reading survives every square-arm result above: the proposer recalls packings it saw rather than constructing a round($\sqrt{N}$) grid. Arm CN samples the bare template A.1 at five N absent from the sum-of-radii scoreboard cited in §6 and from every prior arm — $N = 50, 58, 62, 65$ and 75 , mirroring the square arm’s branch mix (three truncation cells, two $m = 1$ extend cells; the $m \geq 4$ branch is already disconfirmed by arm M). Absence from the pretraining corpus cannot be established and is not claimed. Two predictions were registered against each other before sampling (pushed first; motivation disclosed as post-review): P-CN1, construction — the registered value is modal at ≥ 4 of 5 cells including ≥ 2 of the 3 discriminating cells; P-CN2, recall — modal at ≤ 2 of 5. Fewer than five valid outputs is UNDERPOWERED; ties count against. $n = 15$ weak-tier invocations per cell, scored once by the arm-F pipeline unchanged.

N	k^* , prediction	disc.	valid [Wilson 95%]	on-pred.	mode (count, margin)	verdict	k^* -struct.
50	7, $V(7, 1) = 3.5295867$	no	11/15 [48–89%]	4	3.5122554 (6, +2)	MISS	10/11
58	8, $T(8, 58) = 3.6250000$	yes	14/15 [70–99%]	12	3.6250000 (12, +11)	HIT	14/14
62	8, $T(8, 62) = 3.8750000$	yes	13/15 [62–96%]	13	3.8750000 (13, +13)	HIT	13/13
65	8, $V(8, 1) = 4.0258883$	no	13/15 [62–96%]	9	4.0258871 (9, +7)	HIT	12/13
75	9, $T(9, 75) = 4.1666667$	yes	12/15 [55–93%]	8	4.1666667 (8, +6)	HIT	8/12

P-CN1 holds. The registered value is modal at 4 of 5 cells, at all three discriminating cells, with margins of 6 to 13 samples; every cell clears the evaluability floor; the family rival is emitted 0 times in 39 discriminating-cell validities and no valid output exceeds the family argmax (0/63). The registered structural secondary holds at 5 of 5 cells: the dominant radius backs out $k = k^*$ in 57 of 63 valid outputs. The one miss, $N = 50$, is the arm-CH pattern: 10 of 11 valid outputs are the registered 7×7 grid plus one filler, 4 at the registered interstitial radius $(\sqrt{2} - 1)/14$ and 6 with it in a container corner at $r \approx 0.0123$ — the cell’s mode, 0.017 below the prediction. The template is selected, the filler mis-executed, and the miss is reported as one. At $N = 75$ the four off-mode outputs are alternative constructions, none of them the rival $V(8, 11)$. Validity is 63/75 (84%) with no collapse at large N ; the invalid rows are nine overlaps, two count errors and one output of fraction literals the registered parser does not evaluate (the CH “radical literal” mode).

The $N = 65$ reconstruction. Raw rows are in `arm_cn_collect.jsonl` (75/75 launched, none lost). Eleven $N = 65$ rows were transcribed by reconstructing the standard 8×8 grid string from the emitted filler triple, flagged per row (`reconstructed: true`) and disclosed in `arm_cn_results.txt`; the reconstruction is character-identical for these standard grids. One of the eleven is invalid, so discarding them removes ten of the cell’s thirteen valid rows and leaves three — below the five-row floor, which makes the cell unscoreable

and P-CN1 3 hits of 4 evaluable cells. The three discriminating cells are untouched either way. The arm brings the study corpus to 678 invocations.

What this does and does not show. The anchor survived the move to N with nothing to recall at four of five cells and structurally at all five, so the closed form describes what the proposer builds, not what it has seen, at the branches where it was already supported. This does not establish that the fitted cells are absent from training text and does not extend to the $m \geq 4$ branch; the perturbed-container and direct-recall probes are now arms CP and RP (Appendix E.6), and the iteration question is now arm L (§5.2).

E.6 Arms CP, RP and PP: the anchor without its words, recall asked directly, and paraphrase

Arms CP and RP were drafted at review stage (draft files committed 2026-08-22, disclosed as such) and registered together; arm PP followed. Each registration carries its prompt hashes and analysis script, pushed before the first invocation, and each arm was sampled once and scored once under its own stopping rule, on the same weak-tier subagent channel as arms F, M and CN.

Arm CP destroys every lexical handle while preserving the geometry: the container becomes $[3, 5] \times [3, 5]$, so the model never sees “unit square”, “[0,1]”, or any published magnitude, while the similarity map $x \mapsto 2x + 3$ leaves the optimum, the family and the branch rule invariant (predictions are exactly $2V$ or $2T$; circles are mapped back and scored under the arm-F conventions unchanged). Five cells, all discriminating, 15 each. P-CP1 (modal valid output equals the mapped prediction at ≥ 4 of 5 evaluable cells) **confirmed**: all five cells clear the five-valid floor and the mapped prediction is modal at four — $N = 13$ (valid 14/15, on-prediction 9), 21 (13/15, 7), 31 (9/15, 8) and 58 (12/15, 11). At $N = 75$ (13/15 valid) the mode ties, five outputs at the mapped $T(9, 75) = 8.3333333$ against five at the mapped $T(15, 75) = 5.0$, a deeper truncation of the *same* family; the registered tie rule scores the cell a miss. Pooled validity is 81%, far above the 40% collapse branch of the lexical-keying alternative P-CP2, and the structural secondary S-CP2 holds at 5 of 5 cells ($k = k^*$ majority). The strong form S-CP1 (rival-argmax emissions ≤ 1) is **not met**: 4 of 61 discriminating-cell validities emit the mapped rival, two at $N = 13$ and two at $N = 21$ — 6.6%, against 9 of 261 = 3.4% over every other weak-tier arm in the study pool, defined next. The falsifier F-CP1 (lexical keying) did not fire. One handle survives by construction: N itself, the index any memorized table would be keyed by.

The pooled rival rate, both ways (post hoc, disclosed). Over the weak-tier discriminating-cell validities of the original prompt lineage and the cross-vendor zero-shot arm, the provably better rival is emitted **13 times in 322**, 4.0%: arm F 2/57, trace (arm T, Supplement S6) 1/53, PP 6/75, CP 4/61, CN 0/39, arm V (Supplement S2) 0/16, rectangle (arm G, Supplement S5) 0/11, arm M 0/10. That rule excludes arm GM3, which is weak-tier and direct-emission too and emits the rival 27 times in its 43 discriminating-cell validities (§5.4); folding it in gives **40 of 365**, **11.0%** — nearly triple. We report both figures and no longer describe the narrower pool as covering every arm in the study. The inclusion rule is the instrument: the pooled arms share the arm-F prompt lineage, and GM3 runs the GM-chain instrument at twenty samples per cell against this pool’s five to twenty.

Three further disclosures. The pool is an unregistered aggregate — each constituent verdict is registered, the aggregate is not — over arms with different bars, containers and cell sets, so it is quoted descriptively and never tested. The two arms carrying most of the narrower count, CP and PP, are exactly the two whose registered strong-form secondaries are not met; a pool dropping them reads 3 in 186, which we do not headline because the arms it drops are the ones that failed. The code-channel arms CC and CC2 sit outside both pools; folding them in gives 15 of 400, 3.75%, slightly *below* the rate we report.

Arm RP asks the recall question outright: “What is the best-known maximum sum of radii for $\{n\}$ non-overlapping circles packed inside the unit square?”, answer as a number or UNKNOWN, at three cells with a published value in the paper’s own citations ($N = 13, 26, 30$; the draft named $N = 32$, replaced before registration by $N = 30$ because no cited value exists at 32, documented in the registration) and three held-out cells ($N = 50, 62, 75$). **Every one of the 87 scored responses answers UNKNOWN** — 0 of 44 scoreboard-cell responses recall a published value (bar for the recall branch P-RP2: ≥ 10), zero state a family value, zero state anything else; three invocations were runtime-rejected and counted. P-RP1 **confirmed**, the family-rate asymmetry trivially zero, and the registered falsifier F-RP1 (P-RP2 satisfied) did not fire. A second-round review noted the design carried no positive control, so uniform UNKNOWN could reflect

an abstention prior induced by offering the option; a registered amendment closed this (pushed before its sampling): at $N = 1$, where the optimum 0.5 is elementary, the same prompt draws the correct number in **14 of 15** responses (one UNKNOWN) — C-RP1. The UNKNOWNs therefore reflect the absence of a statable value under this prompt, not the format. The registered caveat stands: a model that fails to state a value may still have been shaped by it, which bounds the simplest mechanism only.

Arm PP holds the container and varies the words: three semantic-preserving paraphrases of template A.1 (clause reorder with synonyms; a declarative numbered-list rewrite; compact mathematical phrasing), fixed and hashed, registered before sampling; the paraphrases were authored by the pipeline that had seen all prior results, disclosed there. Two discriminating cells, $N = 13$ and 31, 15 invocations each. P-PP1 **confirmed at 6 of 6**: the registered prediction is the modal valid output at every paraphrase-cell (valid 11–14 of 15 per cell, pooled validity 83%, on-prediction 4–12), so the anchor is not keyed to the hash-locked prompt string — §8’s (model, prompt string) scoping weakens to the task level at these cells. The strong form again does not carry: 6 of 75 valid outputs emit the rival against the registered S-PP1 bar of ≤ 2 ; $k = k^*$ majority holds at 5 of 6 cells. A formatspread-style sweep (Sclar et al., 2024) remains out of scope.