

CONTENTS

A Closed Form for What the Model Emits: Template Anchoring in Unconditioned Zero-Shot Circle Packing	
Abstract	
1. Introduction	
1.1 Trap zones	
2. Task and recipe family	
2.1 The benchmark	
2.2 The recipe family	
2.3 The order rule is a definition; the branch rule is the hypothesis	
3. Preregistered forecast, out of sample	
3.1 Registration protocol	
3.2 The formula sits at the mode ceiling	
3.3 Square container	
3.4 Preregistered extension (arm M): the filler branch fails at $m > 1$; the fifth k confirms	
3.5 Preregistered extension (arms MU and CH): the anchor does not transfer to one-parent calls, and handing the model the argmax does r	
4. Rectangle transfer	
4.1 Transfer of the rule	
4.2 Negative result: the closed form does not survive the move to rectangles	
5. The tier ladder: three attractor families	
6. Elicitation and faithfulness	
6.1 Design and bundling disclosure	
6.2 Registered outcomes	
6.3 The registered falsifier, both readings	
6.4 Limits of P-T3: what the concentration result will and will not carry	
6.5 Faithfulness with ground truth	
7. Related work	
8. Limitations and untested alternatives	
9. Reproducibility, deviations, and disclosures	
Appendix A — Prompts, verbatim	
Appendix B — Deviations and disclosure table	
Appendix C — Corrections ledger	

A Closed Form for What the Model Emits: Template Anchoring in Unconditioned Zero-Shot Circle Packing

Soham Shailesh Gugale Independent Researcher 28gugales@gmail.com

(Revision 2, 2026-08-01. Implements `p4_review_stats.md`, `p4_review_reviewer2.md` and `p4_review_gecco.md`, plus the arbitration in `p6_cruxes.md`. Every claim that changed relative to revision 1 is itemized in Table 5 (Appendix C); the deviations table is Table 4 (Appendix B), and the ledger correction, the review-process disclosure and the stopping rule are in §9.)

Abstract

Language models are increasingly placed in the proposal role of discovery loops such as FunSearch and AlphaEvolve, on the assumption that their outputs explore a diverse solution space. We characterize one component of that assumption: the *unconditioned* proposal — a single zero-shot, code-free call with no parent program, no fitness feedback and no evaluator in context. On a classic constructive-geometry benchmark, maximizing the sum of radii of N circles in a unit square, a weak-tier proposer does not search. It emits a grid-with-corner-fillers template and truncates it, even when a provably better construction is one parameter away. The behavior admits a closed form: a nearest-square order $k^* = \text{round}(\sqrt{N})$ with a value function $V(k, m)$ identifies the *empirical modal output* — which template, and hence its computed value — at all seven tested N , stated before sampling. Per-sample agreement equals that modal frequency — 56–86% by cell — while a round-number baseline hits 2 of 69 valid samples: the formula captures everything short of sampling entropy. We preregistered these point predictions with prompt hashes before sampling and tested them out of sample on two containers, a square and a rectangle, for which the rule was restated but never refitted; on the rectangle the informative result is negative — 0 of 11 valid samples reached the provably higher-scoring rival — while on-prediction support is partial (5 of 11, not separable from a uniform-template null). Across three tiers we find three attractor families and an inversion: constructive ambition rises monotonically with nominal tier while execution validity does not — it rises then collapses (78% $\rightarrow$ 100% $\rightarrow$ 13% at the primary 10^{-6} tolerance; the third arm is an unattributable serving alias, reported with that caveat throughout). Trace elicitation is an intervention rather than an observation: requesting a method line concentrates outputs onto the anchor (87% vs 70% on-prediction, $p = 0.03$ uncorrected — failing Holm over the registered family, carried by one of three cells and confounded with a collection-wave split of comparable size), with no detectable validity change ($p = 0.30$, $n = 60/\text{arm}$; not powered to exclude an effect of the observed size). Method lines are checkable against emitted coordinates: 54 of 56 scoreable claims (96.4%) describe the object actually built. Two scope conditions are load-bearing: the closed form is a weak-tier regularity, established so far within a single vendor lineage (§8), that does not describe the two higher tiers sampled; and it describes what an *unconditioned* call emits, not what a loop converges to. A preregistered extension run after external review (arm M, §3.4) sharpens the scope in both directions: the branch rule's fifth k confirmed out of sample at $N = 57$ (modal output $T(8, 57) = 3.5625000$, 0 of 10 valid samples on the rival), while the registered falsifier for the filler branch **triggered at 3 of 3 cells** — at $N = 20, 30$ and 41 the model does not extend-and-fill but moves to the (k^*+1) grid, so $V(k, m)$ describes behavior only on the $m \leq 1$ support previously tested, and the extend branch of the rule is disconfirmed at $m \geq 4$. A second preregistered extension (arms MU and CH, §3.5) tested whether the result transfers to parent-conditioned calls, and its registered falsifier **also triggered**: given a provably better in-family parent, 0 of 26 valid mutation samples return to the anchor (keep-or-improve 26/26, grid-order inheritance 26/26), so the anchoring result is a property of unconditioned calls only in the in-family-parent regime — the one-parent regime that discovery loops actually run in is not described by it, and the title of this paper is scoped accordingly. The template reasserts itself only against a far-worse off-family parent (58.8% snap-back). A companion probe that enumerates the full recipe family in the prompt (arm CH) decomposes the mechanism: choice follows the stated score table, but execution fails off-template in 30 of 31 invalid attempts — the template is less what the model prefers than what it can reliably build.

1. Introduction

Ask a language model to place 21 circles in a unit square so as to maximize the sum of their radii, and it will not search. It reaches for the nearest square grid — five by five, radius $1/(2k)$ — then truncates it to 21 circles, leaving four cells empty and their area unclaimed. Those cells are where the value went. A better construction is one parameter away: drop to a four-by-four grid and add corner fillers of radius $(\sqrt{2}-1)/(2k)$, and the sum rises. The model does neither, and behaves the same way on the next sample.

The behavior can be written down. Writing $k^*(N) = \text{round}(\sqrt{N})$ for the nearest-square order, a value function $V(k, m)$ over grid order k and filler count m identifies the modal output of the proposal distribution — which template the model emits most often at each N , and hence its value, from the problem parameters alone, before sampling. The empirical mode equals the predicted value at all seven N tested, and per-sample agreement is bounded by the modal frequency itself, 56–86% by cell (§3.2). The rule also predicts where the behavior costs value: when $k^{*2} \leq N$ the model extends the grid with fillers (a branch arm §3.4 later restricts to $m \leq 1$); when $k^{*2} > N$ it truncates instead of dropping to k^*-1 and filling. These *trap zones* — $N \in [13,15], [21,24], [31,35], [43,48], [57,63]$ (clipped to $[57,60]$ by our sweep bound) — are a property of the branch behavior, not of the packing problem.

Circle packing is the showcase benchmark of the LLM-driven discovery literature, the lineage FunSearch (Romera-Paredes et al., *Nature* 625(7995):468–475, 2024) opened. It is the task on which AlphaEvolve (Google DeepMind technical report, 2025) reported 2.63586276 for $n = 26$ and on which later systems report a narrow band — ShinkaEvolve (Sakana AI, 2025) 2.635983283, HELIX 2.63598308 (2603.07642), GigaEvo 2.636 (2511.17592), AdaEvolve 2.636 (2602.20133) — with ThetaEvolve (2511.23473) claiming new best-known bounds rather than matching it.

What we measure, and what we do not. Those systems query $p(\text{program} \mid \text{parent programs} + \text{fitness scores} + \text{evaluator feedback} + \text{archive context, at generation } t > 0)$. We query $p(\text{coordinates} \mid \text{task description, no parent, no code channel, no feedback, generation } 0)$. These differ on four axes at once — parent conditioning, output modality, feedback, generation index — and we do not claim to have measured the in-loop distribution. What we characterize is the *unconditioned* call. A source audit finds no loop stage in the cited systems that makes precisely this call: FunSearch reseeds emptied islands by cloning a best surviving program, ShinkaEvolve’s stagnation restart re-seeds from its own archive (“initial”, “best” or “archive_random”), and OpenEvolve islands copy the user-provided seed — every in-loop call is program-conditioned. The nearest deployed relatives are initializations conditioned on a seed that “can be trivial” (FunSearch) or “rudimentary”, single-line functions returning constants (AlphaEvolve), plus AlphaEvolve’s “No evolution” ablation, which re-prompts from the same rudimentary program throughout. The unconditioned call is the limiting case these approach as seed information goes to zero, and finding that this limit is a template lookup, its value computable in advance, is actionable for the people who build these loops — seed diversification, forced- k initialization, trap- N avoidance in benchmark selection. Anything about what the loop samples at $t > 0$ is a hypothesis here. The one published result bearing directly on the substitution — “Dictionaries, Not Darwin” (2607.04108, Pan Li), reporting parent-conditioned evolution indistinguishable from fresh independent sampling at equal budget — is warrant from equation discovery, not geometry. §8 names the experiment that would close the gap.

Three framing commitments. Behavioral anchoring is established territory (2505.15392; 2412.06593; 2410.15413) and we inherit its vocabulary; what differs is modality — the anchor is a *construction template*, not a number. The benchmark choice is strategic rather than novel. Preregistration is an adopted standard rather than a contribution (2606.27687; 2607.07184; 2606.11217; 2607.00276); our twist is what is locked — exact-output point predictions from a closed form, on a held-out container.

Contributions. (1) *A closed form that names the modal model output in advance.* $k^*(N)$ with $V(k, m)$ and $T(k, N)$ identify the value a weak-tier proposer emits most often at a given N , verified against a linear-programming oracle on 83 configurations to within 10^{-9} and tested out of sample on two containers, the rectangle rule restated ($q^* = \text{round}(\sqrt{N/a})$, $p^* = \text{round}(\sqrt{N \cdot a})$) but never refitted. Prior work publishes functional forms predicting *aggregate* metrics ahead of a run (2001.08361; 2203.15556; 2411.16035); to our knowledge none predicts a specific multi-decimal *individual emitted output value* ahead of sampling. The claim is scoped to the weak tier (§5). (2) *A tier boundary condition, as three attractor families.* Ambition rises monotonically with nominal tier; validity does not — 78% → 100% → 13% at the primary tolerance, 64% → 90% → 13% at 10^{-9} . The most ambitious tier attempts recursive gaskets and mostly fails to produce a valid packing; it was addressed through a bare serving alias and appears as `opus_alias` with that caveat everywhere. (3) *Checkable faithfulness, and trace requests as interventions.* Method lines are checkable against emitted coordinates, and 54 of 56 scoreable claims (96.4%) describe the object actually built — a clean, well-powered finding (§6.5). Requesting the line at all is an intervention rather than an observation: the concentration shift it produced (87% vs 70%, $p = 0.0325$ uncorrected) fails multiplicity correction, is carried by one of three cells and is wave-confounded (§6.4), so it is reported as met-as-registered exploratory material, not a demonstrated effect. Studies collecting process descriptors *by request* are nonetheless measuring a perturbed distribution.

Non-claim guard. Everything here is a behavioral regularity over emitted outputs. We make no claim about mechanism inside the weights, and the paper uses that vocabulary throughout: the model *emits*, the distribution *concentrates*, the formula *identifies the modal output* (storage-and-retrieval vocabulary corrected; Table 5, item 1). §8 states the leading mechanism-free alternative, arithmetic tractability, and records that we did not test it.

1.1 TRAP ZONES

Substituting $k^*(N)$ into the branch condition: TRAP $N \in [k^2 - k + 1, k^2 - 1]$, CONVERGE $N \in [k^2, k^2 + k]$. Over $N = 10 \dots 60$: [13,15], [21,24], [31,35], [43,48] and [57,63] clipped to [57,60] by the sweep bound.

Inside a zone the rule generally loses value against the best construction available *within the recipe family*. Worst-in-zone penalties fall as N grows: 8.51% at $N = 13$, 7.03% at 21, 6.01% at 31, 5.25% at 43, 4.66% at 57. The penalty reaches zero at $N = 35$, 48, 14 and 15, but **not** at $N = 24$, where truncation gives $T(5,24) = 2.4000000$ against $V(4,8) = 2.4142136$, a residual 0.59% (Table 5, item 3). Branch label and penalty are therefore not coextensive: 4 of the 22 trap-branch N in range cost nothing, so “trap” names the branch, not a guaranteed loss. Where the penalty is zero, truncation is separable from convergence only by structure; §3.3 uses $N = 35$ as that control.

Every closed-form value is recomputed over the constructed coordinates by an independent linear program that knows nothing about the recipe; the script aborts on disagreement. 83 configurations, both branches, every k in $2 \dots 7$, drift below 10^{-9} .

Our bound table (`n_sweep_forecast.json`) covers $N = 10 \dots 30$ only, and its $N = 26$ entry, 2.63598, is ShinkaEvolve’s figure truncated, so the LLM-driven systems are the record on this problem (Table 5, item 4). What survives is that the recipe family is never competitive with that record: deficit 0.02–0.26 across $N = 10 \dots 30$, and 0.0946 at $N = 26$ ($2.63598 - 2.5414214 = 0.0945586$). Above $N = 30$ there is no bound table, so both the deficit claim and the LP gate’s “never exceeds a published bound” abort are unchecked there, covering three of five trap zones and four of our seven square cells.




Figure 1. Predicted versus optimal value, $N = 10 \dots 50$. (i) the rule’s prediction; (ii) the best value in the recipe family; (iii, dotted) published best-known values, terminating at $N = 30$ because the bound table stops there. Shaded bands are the trap zones, the last being [57,63] clipped by `sweep(10,60)`. Worst-in-zone penalties 8.51%, 7.03%, 6.01%, 5.25%; these are family-internal arithmetic and exact at every N , while the published-bound comparison (curve iii) is checkable only up to $N = 30$, where the bound table stops. The gap closes to zero at $N = 35$ and $N = 48$ but **not** at $N = 24$, where 0.59% remains — curve (ii) sits strictly above (i) across [21,24].

2. Task and recipe family

2.1 THE BENCHMARK

Maximize the sum of radii of N non-overlapping circles in the unit square. Feasibility is decidable exactly from emitted coordinates, so proposals score with no human in the loop.

All invocations are zero-shot and code-free. The prompt (Appendix A) fixes the count, states containment and non-overlap, forbids writing or executing code, and demands only a raw Python list of `[x, y, r]` triples. Our reason for the code-free design is that a code channel routes the task to a numerical optimizer, so the distribution would reflect the optimizer. That reason is asserted, not measured — we ran no code-enabled arm (§8) — and it also means the measured object is not the object the code-channel systems query, which is why §1 scopes the claim to unconditioned calls.

2.2 THE RECIPE FAMILY

Three terms, used consistently: the **recipe family** is the parametric set of grid-plus-filler constructions; a **template** is one member; the **anchor** is the template the distribution concentrates on at a given N .

A $k \times k$ grid places one circle per cell center with $r_{\text{grid}} = 1/(2k)$; m fillers sit on interior grid vertices, tangent to their four surrounding grid circles, with $r_{\text{filler}} = (\sqrt{2} - 1)/(2k)$, $0 \leq m \leq (k-1)^2$. Hence $V(k, m) = k/2 + m(\sqrt{2} - 1)/(2k)$. When $N < k^2$ the observed behavior is not to drop to a smaller grid and fill it but to *truncate* — lay out the $k \times k$ lattice, occupy N cells, leave the radius unchanged — giving $T(k, N) = N/(2k)$. These reproduce every previously reported anchor with no fitting: $V(5,1) = 2.5414214$, $V(5,2) = 2.5828427$, $V(5,0) = 2.500$, $T(5,23) = 2.300$, $V(4,7) = 2.3624369$.

The finding that 94 of 95 valid proposals in prior arms were grid-plus-filler is retained only as motivation, since those arms are not reported here; the self-contained version is §5's observation that across the 155 weak-tier rows collected before arm M only two exceeded the family best, both by $\sim 10^{-7}$ (arm M's 57 rows are scored separately in §3.4 and were not swept for this claim).

2.3 THE ORDER RULE IS A DEFINITION; THE BRANCH RULE IS THE HYPOTHESIS

We write $k^*(N) = \lfloor \sqrt{N + 1/2} \rfloor = \text{round}(\sqrt{N})$ for the nearest-square order. This is a *definition* (Table 5, item 2). Given that the model emits some $k \times k$ grid from this family, “the nearest square lattice” yields $\text{round}(\sqrt{N})$ by construction; and $\text{argmin}_k |k^2 - N|$ switches at $N = k^2 + k + 1/2$ where $\text{round}(\sqrt{N})$ switches at $k^2 + k + 1/4$, so on the integers the two coincide exactly — two formalizations of “nearest” agreeing is a symptom that no substantive hypothesis is being selected. Scoring four candidate orders against the three fitting anchors gave `nearest` 3/3, `floor` 2/3, `argmax` 2/3, `ceil` 1/3; under a null where each candidate matches each anchor with probability $1/2$, some candidate among four goes 3/3 about 40% of the time, so this is a weakly identified discrete selection, not a zero-parameter law.

The falsifiable content is the **branch rule**: $k^{*2} \leq N \rightarrow \text{extend with } m = N - k^{*2} \text{ fillers (converge)}$; $k^{*2} > N \rightarrow \text{truncate the } k^* \text{-grid (trap)}$. The extend arm survived testing only at $m \leq 1$ — a preregistered extension disconfirmed it at $m \geq 4$, where the model prefers the (k^*+1) rectangle grid (§3.4). Nothing about “nearest” forces truncation — a proposer with any local search would drop to k^*-1 and fill, which is why the trap zones exist and why a higher-valued rival is reachable inside them. Truncate-versus-drop-and-fill is a genuine dichotomy, and it is what §3 and §4 test. The governing quantity is the signed distance $N - k^{*2}$. Primality is not the variable — 13, 23, 31, 43, 47 and 59 trap while 11, 17, 19, 29, 37, 41 and 53 converge — refuting the prior work's conjecture about primes near 30 before any model is queried.

3. Preregistered forecast, out of sample

3.1 REGISTRATION PROTOCOL

For each cell the prompt is pinned and its SHA-256 digest written to disk first — the $N = 13$ square prompt hashes to `32db485b...`. Predictions P1–P5 are registered in the header of `arm_f_repro.py`, with standalone files for the parallel arms (`arm_s_preregistration.txt`, `arm_o_preregistration.txt`, `arm_t_preregistration.txt`). Raw outputs are stored verbatim, failures included; scoring is deterministic and local.

Four properties are disclosed rather than repaired. Temperature, top-p and top-k are not exposed by the runtime. The alias-to-weights binding is a promise, not a hash. The subagent inherits instruction files outside the task prompt, so the digests lock a *fragment* of the conditioning context — they are prompt-fragment hashes, and a replicator cannot reconstruct the inherited files from what we release. Finally, hash coverage is incomplete (§9).

Two scoring conventions were fixed in advance. Validity is reported at 10^{-9} and 10^{-6} , both logged, 10^{-6} primary: proposers print six to eight decimals, so an eight-decimal tangency misses contact by $\sim 5 \times 10^{-9}$, while 10^{-6} sits far below the $\sim 10^{-2}$ gap between rival constructions. Both tolerances appear in every validity figure below, as registered. Value matching uses a 2×10^{-3} window; §5 records where that window is too loose for the distinction it was asked to carry.

Reproducibility posture. Prompt hashes were registered before sampling in the preregistration files for arms F, S, O and T (`arm_f_repro.py` header, `arm_s_preregistration.txt`, `arm_o_preregistration.txt`, `arm_t_preregistration.txt`); what is hashed is the task prompt only, not the harness, so the dispatch wrapper and the inherited instruction files sit outside the digest; what is not attestable at all is the `opus_alias` alias-to-weights binding, a vendor promise rather than a hash; hash-coverage gaps, the unhashed pilot prompt and the ledger provenance corrections are itemized in §9 and Appendix A.

3.2 THE FORMULA SITS AT THE MODE CEILING

An exact point prediction invites a fair objection: a per-sample hit rate well short of 100% is not “predicting the exact output”, which is what revision 1’s abstract claimed (Table 5, item 7). The number should be read against the sampling entropy of the distribution, computed in `p11_mode_baseline.json` (bare arm, 2×10^{-3} buckets). **At all seven tested N the predicted value is the empirical modal output**, and the on-prediction rate equals the modal frequency to within one sample:

N	predicted	empirical mode	modal freq	on-prediction / valid	k*-structure / valid†
13	1.6250000	1.6250	10/18	10/18 (56%)	17/18
17	2.0517767	2.0518	3/4	3/4 (75%)	3/4
21	2.1000000	2.1000	12/15	12/15 (80%)	13/15
31	2.5833333	2.5833	12/17	13/17 (76%)	17/17
35	2.9166667	2.9167	3/4	3/4 (75%)	3/4
37	3.0345178	3.0345	3/4	3/4 (75%)	4/4
43	3.0714286	3.0714	6/7	6/7 (86%)	7/7

† Post-hoc structural check, not preregistered, prompted by an external review of this revision (`diagnostics_kmatch.py`): the grid order backed out of each sample’s dominant radius ($k = \text{round}(1/2r)$) equals $k^*(N)$. All 50 of 50 on-prediction samples are k^* -structured — the value match is not an accounting coincidence — and 64 of 69 valid samples overall (93%) sit on the k^* grid, so most value misses are k^* -grid variants (perturbed radii or filler counts), not different constructions.

No predictor of a single value can exceed the modal frequency, and this one attains it at 7 of 7 cells (at $N = 31$ the on-prediction count exceeds the 4-dp modal bucket by one sample — 13/17 against 12/17 — a bucketing convention, not a violation of the ceiling): the residual is dispersion, not misprediction. As a floor, a naive “the model emits a round number” baseline hits 2 of the same 69 valid samples (3%). Two scope conditions: this is the bare weak-tier arm — pooled over the four discriminating square cells (full ledger), the two higher tiers and the rectangle, the on-prediction rate is 46% ($47/102 = 41/57 + 1/30 + 0/4 + 5/11$; the three non-discriminating square cells, 9/12, are excluded because prediction equals the family argmax there), because §5’s higher tiers are almost never on-prediction — and “modal” is a property of the observed sampling regime, not of a pinned decoding configuration (§8).

3.3 SQUARE CONTAINER

The square arm sampled Haiku-tier proposers at $N = 13, 17, 21, 31, 35, 37$ and 43. The rule was fitted on $N = 23/26/27$ only, so every cell is out of sample. Four cells are *discriminating*: rule and family optimum disagree, so a proposer that searched the family would return a higher number. Across those cells ($N = 13, 21, 31, 43$), **18 of 23 valid invocations landed on the predicted construction** and the rival-argmax value was reached **2 times in 23** — both at $N = 21$, both the 4×4 grid plus five fillers at 2.2588835.

Those figures cover the **original 45-invocation bare arm only** ($\text{ids} \leq 5$ at $N = 13/31$, ≤ 10 at $N = 21$, all rows at $N = 17/35/37/43$), the corpus that existed when §3 was written. §6 later added 40 bare rows at $N = 13/21/31$ under the same arm label; over the full shipped ledger the same four cells give **41/57 on-prediction and 2/57 rival**. Both are reported. §3 and §6 are computed on nested

subsets of one arm, not on independent samples (Table 5, item 28), and Table 1 carries the split.

That the anchor first breaks at $N = 21$ is suggestive — it is the bottom of a zone carrying a 7.03% penalty, so the anchor looks weak where obeying it costs — but that is a hypothesis, not a finding. **P4** registered $N = 37$ as converging on 3.0345178, a prime predicted clean, contradicting the prior work's guess; it confirmed. **P5** registered $N = 35$ at 2.9166667, the top-of-zone control; three of four valid samples landed there, separable from convergence only because the structure classifier reads radii rather than totals; the fourth used a 7×7 lattice truncated to 35 (sum 2.5), outside the rule. At $n = 4$ this is “consistent with” rather than confirmed, and deserves running to $n = 20$.

Bookkeeping. Five invocations were rejected by the runtime's 20-subagent concurrency cap *before reaching a model*. We excluded them under a rule that was not preregistered, so the arm is reported both ways — $35/45 = 77.8\%$, or $35/50 = 70.0\%$ had the rejections been scored as model failures (Table 5, item 5) — and the five records appear in the ledger in no form (§9). The arm also contains three parse failures (Table 5, item 6), and one $N = 37$ proposer derived $r = (\sqrt{2}-1)/12$ correctly in prose, then transcribed 0.03571429.

3.4 PREREGISTERED EXTENSION (ARM M): THE FILLER BRANCH FAILS AT $M > 1$; THE FIFTH K CONFIRMS

Two external reviews of this revision observed, correctly, that (i) $V(k, m)$'s filler branch was empirically supported only at $m \in \{0, 1\}$ — the converge cells $N = 17$ and 37 both have $m = 1$ — and (ii) the $k = 8$ trap zone appeared in Figure 1 but was never sampled. Arm M closes both gaps and was registered before sampling (`arm_m_preregistration.txt`, commit `e528c6b`; prompt hashes in `arm_m_prompts.json`): $n = 15$ bare weak-tier invocations at each of $N = 20$ ($k^* = 4, m = 4$, predicted $V(4,4) = 2.2071068$), $N = 30$ ($k^* = 5, m = 5$, $V(5,5) = 2.7071068$), $N = 41$ ($k^* = 6, m = 5$, $V(6,5) = 3.1725890$) and $N = 57$ ($k^* = 8$, trap, $T(8,57) = 3.5625000$, rival $V(7,8) = 3.7366935$), with an explicit tie-stated falsifier for each half.

The falsifier for the filler branch triggered at 3 of 3 cells. The registered predictions P-M1–P-M3 are disconfirmed: not one of the three converge cells has $V(k^*, m)$ as its modal output. The model does not extend a k^* -grid with fillers; it moves *up* to the (k^*+1) grid and truncates or exactly fills it — modal outputs $2.0 = T(5,20)$ (12 of 15 valid; $N = 20$ is the 5×4 rectangle exactly), $2.5 = T(6,30)$ (11 of 14; 6×5 exactly) and $2.9285714 = T(7,41)$ (6 of 7; 7×6 minus one). Exactly one sample in 36 valid converge-cell rows emitted the registered construction — a textbook $V(4,4)$, all four fillers at $(\sqrt{2}-1)/8$. Per the registered falsifier wording, **$V(k, m)$ is restricted to the $m \leq 1$ support previously observed**, and the branch rule's extend arm is disconfirmed at $m \geq 4$; the intermediate “places some fillers” outcome named in the registration did not occur (filler count 0 in 35 of 36 valid converge rows). The emergent regularity is arithmetic, not geometric: where N factors as $a \times b$ with a, b near $\sqrt{N}$ the model emits that exact rectangle of uniform circles ($20 = 5 \times 4$, $30 = 6 \times 5$), which is further evidence for §8's arithmetic-tractability alternative — rectangle factorization is mentally computable where filler tangency is not.

P-M4 confirmed. At $N = 57$ the modal valid output is $T(8,57) = 3.5625000$ (6 of 10 valid at 10^{-6}), the $k = 8$ grid truncated to 57; 0 of 10 valid samples reached the rival $V(7,8)$, against a registered rival prediction of 0. The branch rule's truncate arm now has five confirmed k (4, 5, 6, 7, 8) and Figure 1's $k = 8$ zone is sampled. Validity by cell: 15/15, 14/15 (one parse failure — fraction literals), 7/12 and 10/15 at 10^{-6} . Three $N = 41$ invocations died in the runtime before reaching a model and are excluded under the rejection rule registered for this arm in `arm_m_preregistration.txt` — registered precisely because the analogous arm F exclusion was not (§9) — and counted here: the cell is $n = 12$ sampled of 15 launched. Raw rows verbatim in `arm_m_collect.jsonl`; scoring in `arm_m_analysis.py`, conventions identical to arm F.

The net effect on the paper's claim is a sharpening, not a rescue: the closed form now reads “ $k^* = \text{round}(\sqrt{N})$ names the grid family; at $k^{*2} > N$ the model truncates that grid (five k confirmed); at $k^{*2} \leq N$ it extends only to $m \leq 1$ — beyond that it prefers the (k^*+1) rectangle, and $V(k, m)$ does not describe it.”

3.5 PREREGISTERED EXTENSION (ARMS MU AND CH): THE ANCHOR DOES NOT TRANSFER TO ONE-PARENT CALLS, AND HANDING THE MODEL THE ARGMAX DOES NOT DISLodge THE TEMPLATE

Two questions remained open after arm M, and both were pressed by the council review of this revision: does the unconditioned anchoring result say anything about calls that carry a parent program — the regime evolutionary loops actually run in — and does the §8 arithmetic-tractability alternative survive a direct probe? Arms MU and CH were registered together before sampling

(`arm_mu_preregistration.txt`, commit `2b7d202`; prompt hashes in `arm_mu_prompts.json`), with numeric decision thresholds, a power analysis (0.83 at the registered effect size), tie-against conventions, and a falsifier (F-MU1) whose consequence — the title and abstract of this paper are scoped to *unconditioned* calls only — was written down before the first invocation.

Arm MU issues one-parent mutation calls: the bare prompt plus an existing packing and the instruction to modify it so the sum of radii increases. Three parents per cell ($N = 13, 21, 31$; $n = 15$ per condition-cell, 135 invocations): *A_anchor*, the model's own unconditioned anchor $T(k^*, N)$; *B_rival*, the provably better in-family rival $V(k^*-1, m)$; and *C_offfamily*, a diagonal-chain packing scoring ≈ 0.59 – 0.62 , far below either.

F-MU1 triggered; the anchoring result dissolves under conditioning. Under *B_rival* the pooled anchor-rate is 0 of 26 valid (0%, CI [0.00, 0.13]) against a registered DISSOLVES threshold of $\leq 20\%$; keep-or-improve is 26 of 26 against a threshold of $\geq 50\%$. Every valid *B_rival* sample keeps or refines the rival construction, and every one inherits the rival's grid order (k -inheritance 26/26, disconfirming P-MU3's $< 50\%$ prediction). The template does not reassert itself when a better in-family parent is present. *A_anchor* makes the same point from the other side: even with the model's *own* anchor as parent and an explicit increase-the-sum instruction, only 7 of 28 valid samples (25%, CI [0.13, 0.43]) stay on the anchor — P-MU1's $\geq 50\%$ prediction is not met. Per the registered falsifier wording: the unconditioned-call anchoring result does not transfer to parent-conditioned calls, the loop-relevance framing is disconfirmed for the one-parent regime, and this paper's claims are scoped to unconditioned calls throughout.

The one condition where the template does reassert itself is the one the registration predicted it would: given the *off-family* diagonal parent, 20 of 34 valid samples (58.8%, CI [0.42, 0.74]) abandon the parent entirely and emit a k^* -grid construction — P-MU4 holds per-cell at $N = 21$ (7/13) and $N = 31$ (12/14, every valid sample on the $k = 6$ grid), though not at $N = 13$ (1/7). The asymmetry is the finding: a good in-family parent is followed, a bad off-family parent is discarded for the template. The attractor is real but weak — it captures the trajectory only when the parent is worse than what the template supplies.

Arm CH hands the model the entire recipe family: the bare prompt plus every admissible $T(k, N)$ and $V(k, m)$ value enumerated, including the family argmax, with the instruction to use whichever scores highest or anything better ($n = 15$ per cell, 45 invocations). The §8 tractability alternative says the model truncates at k^* because that is the only construction it can compute; enumerating the family removes the computation burden, so under that reading the modal output should become the stated argmax (P-CH2). Under the template reading it stays $T(k^*, N)$ despite the argmax being printed in the prompt (P-CH1).

Registered verdict: P-CH1 holds (2 of 3 cells) — but the registered metric carries a survivorship artifact we disclose rather than exploit. Under the registered scorer, which takes the modal output *among valid samples*, the template reading wins: at $N = 13$ the modal valid output is $T(4, 13) = 1.6250000$ against a stated argmax of 1.7761424, at $N = 31$ it is $T(6, 31) \approx 2.5833$ (4 of 7 valid) against 2.7485281 (3 of 7) — a one-sample margin — and only $N = 21$ goes to the argmax (2.2588835, 3 of 5 valid). P-CH2 is disconfirmed at 1 of 3 cells. Read no further and the tractability alternative fails the probe. But validity is low in this arm — 2, 5 and 7 valid of 15 by cell — and the failure modes reverse the picture at the *attempt* level: **all 31 invalid rows attempt the stated argmax** (every one carries the family's filler radius). Of those, 30 misplace the fillers — typically at the container corners, where radius $(\sqrt{2}-1)/(2k)$ overlaps the adjacent grid circle; the correct construction places them at the grid interstices — and 1 emits radical literals the registered scorer cannot evaluate. Counting attempts rather than valid outputs, the model *chooses* the argmax in a large majority of CH invocations and simply fails to execute it; the valid-conditioned modal is dominated by the template because the template is the only construction the model executes reliably. The registered metric was fixed before sampling, so we report its verdict as registered — and flag in the same breath that it conditions on exactly the variable (execution success) that separates the two readings. The honest summary is a decomposition, not a winner: handed the argmax, the model's *choice* follows the score table, but its *execution* fails off-template in 30 of the 31 invalid attempts — the §8 tractability alternative is wrong about selection and right about instantiation. This is the same asymmetry arm MU found from the other direction, and it narrows the paper's central claim to its defensible core: the template is not what the model prefers; it is what the model can build.

Raw rows verbatim in `arm_mu_collect.jsonl` (180 of 180 launched invocations collected, no runtime deaths); scoring in `arm_mu_analysis.py`, committed before sampling completed; frozen output in `arm_mu_results.txt`. This arm brings the study corpus to 468 invocations.

4. Rectangle transfer

4.1 TRANSFER OF THE RULE

For a $1 \times a$ rectangle the template has two parameters: $q^* = \text{round}(\sqrt{N/a})$ columns across the width, $p^* = \text{round}(\sqrt{N \cdot a})$ rows across the height. At $a = 1$ both collapse to $\text{round}(\sqrt{N})$ — the same rule with the aspect ratio put back, not a new rule fitted to new data, and no rectangle model output existed when it was written. It was verified against an independent LP at 213 configurations, drift below 10^{-9} , including a cross-domain check: at $a = 1$ with $p = q = k$ the square file's closed form must equal this file's LP. Shape mismatch grows as a leaves 1 — over $N = 10 \dots 45$ the count of N whose predicted shape differs from the optimal one rises from 12 at $a = 1.0$ to 23 at $a = 1.5$, 20 at $a = 2.0$ and 23 at $a = 3.0$.

We probed the two sharpest cells with sixteen proposers in a container none had been given before: $N = 19$ at $a = 3$ (predicted 3.1666667, rival 3.5749194) and $N = 25$ at $a = 2$ (predicted 3.1250000, rival 3.4832492). With all sixteen scored, **5 of 11 valid proposals landed on the predicted value and 0 of 11 reached the rival**.

We characterize all eleven here, since a referee objected that the majority were left undescribed. On-prediction: two at $N = 19/a = 3$ (3.166666667 and 3.16666673, the latter agreeing to six decimals not seven) and three at $N = 25/a = 2$ (3.125 exactly). Off-prediction at $N = 25/a = 2$: two samples emitted a uniform 5×5 grid at $r = 0.1$ summing to 2.5000000 — the *square* rule with no aspect correction, $k = \text{round}(\sqrt{25}) = 5$ — plus one at 3.0 and one at 3.151875. At $N = 19/a = 3$, **two** samples exceeded the prediction from outside the family, 3.45 and 3.5, both below the rival (Table 5, item 8).

We therefore report the rectangle as **partial out-of-sample support, not confirmation**: $5/11 = 45\%$, Wilson 95% CI [21%, 72%], and a proposer choosing uniformly among the three or so plausible template shapes would land on-prediction about a third of the time, so 5/11 does not separate cleanly from that null. Two qualifications: validity degrades in the tall container (4/8 at $a = 3$ against 7/8 at $a = 2$, three of four $a = 3$ failures being overlaps), a separate finding and a confound on that cell at $n = 8$; and the rectangle ledger carries no prompt hashes, no run dates and no proposer fields, so the arm with the strongest transfer claim has the least provenance (§9).

4.2 NEGATIVE RESULT: THE CLOSED FORM DOES NOT SURVIVE THE MOVE TO RECTANGLES

The rectangle generalizes the *rule* but not the *formula*, and we keep the failure because it is what a verification gate is for. In the square every interior vertex has four identical neighbors, so one expression covers every filler. In a $1 \times a$ rectangle, with half-spacings $h_x = 1/(2q)$ and $h_y = a/(2p)$, a filler is also capped by horizontal and vertical spacing to *adjacent fillers*. Those constraints are inactive when $h_x = h_y$ — why the square case could not have revealed them — and bind only when neighboring vertices are occupied, so the cap depends on m and on which vertices a construction uses, and no expression in (p, q, m, a) reproduces it. The LP gate caught this on the first run: at $a = 1$, $p = 2$, $q = 4$, $m = 1$ the natural generalization $r_f = \min(\text{diag}, h_x, h_y)$ returns **1.125** against a true **1.1545085**, capping against a neighboring filler that does not exist. We retain closed forms only where provably exact — full grid and truncated grid — and use the LP as the value oracle for the extend branch, so what §4.1 tests is the LP-backed prediction.

Table 1 — Forecast versus outcome, both containers.

Cell	Branch	Predicted	Rival argmax	Valid / sampled	On-pred / valid	Rival
N=13 (sq)	truncate	1.6250000	1.7761424	4/5 · 18/20	3/4 · 10/18	0
N=17 (sq)	extend	2.0517767	†	4/5	3/4	†
N=21 (sq)	truncate	2.1000000	2.2588835	7/10 · 15/20	4/7 · 12/15	2
N=31 (sq)	truncate	2.5833333	2.7485281	5/5 · 17/20	5/5 · 13/17	0
N=35 (sq)	truncate	2.9166667	†	4/5	3/4	†
N=37 (sq)	extend	3.0345178	†	4/5	3/4	†
N=43 (sq)	truncate	3.0714286	3.2416246	7/10	6/7	0
N=19, a=3	truncate	3.1666667	3.5749194	4/8	2/4	0
N=25, a=2	truncate	3.1250000	3.4832492	7/8	3/7	0

Plain entries are the original 45-invocation bare arm (§3.3's corpus); **bold** the same cells over the full 85-row bare ledger after §6 added 40 rows. † marks non-discriminating cells (prediction equals family argmax), excluded from rival denominators. Square discriminating cells: 18/23 and 2/23 on the original corpus; 41/57 and 2/57 on the full ledger. Rectangle: 5/11 and 0/11. Sources: `arm_f_repro.py` over `arm_f_candidates_v2.jsonl`, `arm_g_rect.py` over `arm_g_candidates.jsonl`.

5. The tier ladder: three attractor families

Holding prompt, container and scoring fixed, we varied only the nominal proposer tier across the three cells that discriminate hardest between prediction and family optimum — $N = 13, 21$ and 31 , each inside a trap zone. The tiers do not merely differ in success rate: they attempt qualitatively different constructions.

The weak tier truncates templates. The Haiku-tier proposer produces uniform grids of radius $1/(2k)$ truncated to N circles, with corner fillers only when the grid underfills. Across the 45 original bare invocations spanning $N \in \{13, 17, 21, 31, 35, 37, 43\}$, **35 were geometrically valid at the primary 10^{-6} tolerance (77.8%), 29 at 10^{-9} (64.4%)** — recomputed from the ledger, correcting the headline validity figure and the inversion triple in the abstract and Contribution 2 (Table 5, item 9). At the three discriminating cells 12 of 16 valid samples landed on the predicted value and the rival was reached twice, both at $N = 21$.

The middle tier perturbs and mixes. The Sonnet-tier arm was preregistered in `arm_s_preregistration.txt`, disclosing that 5 of 20 samples at $N = 13/21$ had been seen before registration and that $N = 31$ was fully blind. It was valid 30/30 at 10^{-6} and 27/30 (90%) at 10^{-9} — so the “100%” leg is tolerance-dependent and both numbers are stated. It was almost never on-prediction: 0/10, 0/10, 1/10 — 1/30 pooled against the weak tier's 12/16 at the same cells. Instead of truncating a uniform grid it perturbs one, with enlarged edge rows, hexagonal interior rows and two or three distinct radii in the same packing (29/30 multi-radius).

One Sonnet sample at $N = 31$ emitted 27 circles at $r = 1/12$ with 4 corner circles at $r = 1/8$, summing to 2.749999991 — above 2.7485281, the best the family reaches at that N . Recomputed from its stored coordinates, the construction is tangency-tight: minimum pairwise and wall slack both exactly zero at tolerance zero, the exactly-tangent contacts involving the binary-exact $r = 1/8$ corner circles while the $r = 1/12$ circles carry small positive slack. It is the only sample in the study to leave the recipe family upward.

That sample also exposes a scoring defect. At $N = 31$ the family rival (2.7485281) and this out-of-family value (2.75) differ by 1.47×10^{-3} , *inside* the registered 2×10^{-3} window, so the value classifier counted the escape as a hit on the construction it exceeded. The window was chosen for the anchor comparison, where the nearest competing value is ~ 0.15 away, and reused where it does not fit. Classified by structure instead, Sonnet's rival count at $N = 31$ is 2/10 and pooled 5/30 (Table 5, item 10). Every categorical claim in this paper about “reached the rival” or “left the family” is made at structure level or at 10^{-6} , never at 2×10^{-3} .

The top tier attempts recursive constructions and fails to build them. The third arm was invoked through a bare tier alias. We name it `opus_alias` throughout and make no claim about which weights served it: the runtime accepts only the bare alias and exposes no dated identifier, so the binding is a vendor promise rather than an attestable fact. No version number appears anywhere in this paper.

Two anomalies make that caveat load-bearing. Completion times ran 2.8–9 s across all 30 invocations against 75–250 s for Haiku and 150–1170 s for Sonnet, and the reported token count was uniform at 49,906 across the first 20 completions, $\approx 49.9k$ thereafter — a signal of metadata that is not per-completion, or of a cache, router-fallback or truncation path. **Neither anomaly appears in any released artifact** — we therefore treat both as anecdotal context from an unreleased session transcript, not as data-supported properties of the tier: the ledger has no latency and no token field, and both figures come from the runtime session transcript, so a referee cannot check them.

Under those caveats the arm — preregistered fully blind in `arm_o_preregistration.txt` — was valid in 4 of 30 invocations (13%) at both tolerances. At every cell it attempted a more ambitious construction than either other tier: quarter-circle corners at $r = 0.25$ with Apollonius-style fillers at $N = 13$ (10/10 samples), mixed-radius 4×4-ish grids at $N = 21$, and a coarse 3×3 grid at $r \approx 1/6$ with border strips at $N = 31$ (0/10 valid). Failures are geometric rather than numerical — edge strips at $r = 0.03$ placed 0.138 from an $r = 1/6$ grid circle needing 0.197 — and twice the arm padded to count with zero-radius circles. Valid samples score *below* the trap they were expected to fall into (1.26–1.41 at $N = 13$ against 1.625). A post-hoc diagnostic (not preregistered; labeled as such in the repository) finds no evidence of a tolerance artifact behind the 13% in this sample: 24 of 26 invalid samples overlap grossly (median maximum overlap 3.3×10^{-2} ; radius-shrink repair costs a median 15% of sum-of-radii), while tolerance-scale near-misses ($< 2.5 \times 10^{-5}$) occur instead in 5 of 7 geometry-scored weak-tier failures — the exact-tangency grids, not the ambitious constructions.

Two decisions here were made after seeing the data and are recorded as deviations (Table 4). P-O1, P-O2 and P-O4 are reported **not evaluable**, on the grounds that a validity collapse of this size makes an on-prediction tier comparison uninterpretable; P-O3 is trivially satisfied on 4/4 valid samples; the registered disconfirmation — regression toward the trap — did not occur. De-registering three of four predictions in the one arm that disconfirmed is exactly the move preregistration exists to constrain, which is why it is tabled rather than narrated. And the arm was initially excluded from the ladder, then included once its results were known; we follow the later decision and report the ladder both ways — with the arm 78% → 100% → 13%, without it 78% → 100% plus the qualitative shift from truncation to perturbed multi-radius hybrids, which is already the publishable contrast.

Table 2 — Three-attractor ladder.

Tier	Cells	n	Attempted family	Valid 10 ⁻⁶	Valid 10 ⁻⁹	On-pred	Rival	Failure mode
haiku (bare)	13,17,21,31,35,37,43	45	uniform grid, truncated; fillers when underfull	35 (78%)	29 (64%)	12/16 (3 matched cells)	2	overlap 6, outside 1, parse 3
haiku, matched cells	13,21,31	60	as above	50 (83%)	—	35/50	2	overlap 5, outside 2, parse 3
sonnet	13,21,31	30	perturbed grid, hexagonal rows, 2–3 radii	30 (100%)	27 (90%)	1/30	5 (structure)	none
<code>opus_alias</code> ‡	13,21,31	30	quarter-circle corners, Apollonius fillers, coarse grid + border strips	4 (13%)	4 (13%)	0/4	0	overlap 24, nonpositive radius 2

‡ Addressed through a bare serving alias; the alias-to-weights binding is not attestable, and this row is not a statement about any released model version. Tier denominators are not matched — the weak tier spans seven N against three — which is why the matched-cell row is given separately, and it is the matched comparison (83% → 100% → 13%) that should be quoted across tiers. Sources: `arm_f_candidates_v2.jsonl`, the three preregistration files.




Figure 2. Three attractor families at a single cell, $N = 31$. One representative packing per tier, ledger sample id in each panel title (weak tier `bare` id 1, on-prediction truncation 2.5833333; middle tier `sonnet_bare` id 1, mixed-radius rival 2.7485281; `opus_alias` id 1, invalid, with only the offending circles highlighted). Source: `arm_f_candidates_v2.jsonl` via `fig_scripts.py`, deterministic (the exemplar rows' geometry is byte-identical in v1; v2 differs only in provenance metadata).

"Monotone" applies to one axis only (Table 5, item 11): ambition rises monotonically — truncated template, perturbed hybrid, recursive gasket — while validity rises then collapses. "Ambition" is measurable, not merely a qualitative read: a post-hoc diagnostic (disclosed; `diagnostics_ambition.py`, all samples with emitted geometry, invalid included) gives median distinct radii per sample 1 $\rightarrow$ 2 $\rightarrow$ 3 (means 1.33 $\rightarrow$ 2.40 $\rightarrow$ 3.37) and median off-lattice center fraction 0.08 $\rightarrow$ 0.43 $\rightarrow$ 0.95 across the three tiers — monotone on both operationalizations. At the canonical and plausibly contaminated cell $N = 26$ all tiers converge on the same 2.5414 attractor; they diverge only at withheld trap cells. The branch rule of §2 is therefore a weak-tier regularity — single-vendor evidence until a cross-vendor arm completes (§8) — and this is a **boundary condition on the main result**, not an independent second finding. It also cuts against §1's framing: the discovery systems cited there run frontier or mid-tier proposers, so the tier at which our regularity holds is the tier they do not use, while the tier closest to what they use escapes the trap in our own data.

6. Elicitation and faithfulness

This section is explicitly secondary material and the result leans negative: the registered validity prediction fails, the concentration effect survives only as a fragile directional finding, and the faithfulness audit checks description against artifact, not against process.

6.1 DESIGN AND BUNDLING DISCLOSURE

A ten-versus-ten pilot at $N = 21$ compared the bare prompt against a variant asking the proposer to name its construction on a leading `METHOD:` line: validity rose from 7/10 to 10/10 and the higher-scoring rival disappeared, from 2 of 7 valid bare samples to 0 of 10 trace samples. Before any scaled sample was drawn we registered four predictions and an explicit falsifier in `arm_t_preregistration.txt` (sha256 `ab7900a8...`), disclosing that the pilot's trace prompt had drifted from the bare template beyond the method line — it omitted the `[0,1]x[0,1]` tokens and rewrote the output-format line — so the pilot's intervention was method-line-plus-rewording, bundled. Pilot samples are never pooled with `trace_v2`.

The scaled prompt, `trace_v2`, is itself a **bundled prompt-format-and-trace-request** intervention, though a near-minimal diff: one inserted `METHOD:` line, plus "After the METHOD line, " prepended to the output line and "no other text" changed to "no other text after the list". That second change is not cosmetic — §6.2 shows it doing measurable work — so the measured effect is *method-line-plus-output-line-rewording*, not the method line alone: a weaker version of the confound for which the pilot was discarded, held to the same standard.

The scaled arm ran 100 new invocations, bringing both arms to 20 per cell at $N \in \{13, 21, 31\}$ and the corpus to 215 logged square invocations (85 bare = 45 original + 40 new; 70 trace = 10 pilot + 60 trace_v2; 30 sonnet; 30 `opus_alias`), plus 16 rectangle invocations, 231 total at the close of revision 2 (Table 5, item 27); the post-review arm M (§3.4) later added 57 sampled weak-tier invocations (60 launched), bringing the study corpus to 288, and the post-council arms MU and CH (§3.5) added 180 more (180

launched, none lost), for a final corpus of 468. One disclosure revision 1 did not carry: **20 of the 60 bare samples are pre-existing arm-F rows** (N = 13 ids 1–5, N = 21 ids 1–10, N = 31 ids 1–5) whose results were known when P-T1–P-T3 were written; only trace_v2 is fully fresh (§6.4). Scoring used `arm_f_repro.py` unchanged with the registered 2×10^{-3} window; the Fisher exact test comes from the hypergeometric tail in `arm_t_analysis.py`, checked against a reference.

6.2 REGISTERED OUTCOMES

Preregistered criteria are evaluated exactly as registered, and unregistered alpha thresholds are applied nowhere; revision 1 applied one asymmetrically (Table 5, item 12).

P-T1, validity: NOT confirmed. P-T1 is the one prediction that registered an alpha (“validity $\geq$ bare at EACH N, and pooled one-sided Fisher $p < 0.05$ ”). Direction held at all three cells, but the pooled test gives $p = 0.30$: 53/60 (88%) trace_v2 against 50/60 (83%) bare, a difference whose detection needs several hundred samples per arm at 80% power — a failure to detect, not a demonstration of no effect (Table 5, item 13). The direction that did hold is not geometric: bare had **3 parse failures and 7 geometric failures**, trace_v2 **0 parse and 7 geometric**, so among parsed samples geometric validity is 50/57 (87.7%) versus 53/60 (88.3%). The P-T1 direction is format compliance, exactly what the reworded output-format line manipulates; validity is reported split by cause from here on.

P-T2, rival suppression: CONFIRMED as registered (directional, no alpha): 1 of 53 valid trace_v2 against 2 of 50 valid bare. Counts this near zero carry no inferential weight and we claim none (Fisher $p = 0.48$, carrying no decision). One trace_v2 sample at N = 31 hit the rival 2.7485281 to within 3.6×10^{-8} — the first weak-tier rival hit at N = 31 in any arm, and formerly one of the scorer’s mismatches (§6.5).

P-T3, anchor concentration: met as registered (directional), inferentially fragile. Among valid samples, 46 of 53 trace_v2 (87%) landed on the registered prediction against 35 of 50 (70%) bare. The p-value it never carried is 0.0325, one-sided uncorrected (§6.4).

P-T4, faithfulness: confirmed under a corrected scorer. The registered scorer gives 38 of 41 scoreable claims matching (93%) against a 90% threshold; blind hand-adjudication under a rubric frozen before rescoring gives 54 of 56 (96.4%). §6.5 reports both.

Table 3 — Paired trace grid (scaled arm only).

N	Arm	n	Valid (parse / geom. fail)	On-pred / valid	Rival
13	bare	20	18 (0 / 2)	10 (56%)	0
13	trace_v2	20	18 (0 / 2)	16 (89%)	0
21	bare	20	15 (1 / 4)	12 (80%)	2
21	trace_v2	20	18 (0 / 2)	16 (89%)	0
31	bare	20	17 (2 / 1)	13 (76%)	0
31	trace_v2	20	17 (0 / 3)	14 (82%)	1
pooled	bare	60	50 (3 / 7)	35 (70%)	2
pooled	trace_v2	60	53 (0 / 7)	46 (87%)	1

Per-cell one-sided Fisher on on-prediction: N = 13 $p = 0.030$; N = 21 $p = 0.41$; N = 31 $p = 0.50$. Pooled $p = 0.0325$; Holm threshold over the registered family ($m = 3$ tested with p-values) is 0.0167. The pilot (10 invocations at N = 21, 10/10 valid, 9/10 on-prediction) is reported separately and never summed into any total. Regenerated by `arm_t_analysis.py` over `arm_f_candidates_v2.jsonl`.




Figure 3. Bare versus `trace_v2`, by cell and pooled. Caption caveat: the bars pool the bare arm's two collection waves, which §6.4 shows drift materially, so the pooled bare bar is a mixture; the same-wave comparison in §6.4 is the wave-controlled picture.

6.3 THE REGISTERED FALSIFIER, BOTH READINGS

The registered falsifier reads: "if `trace_v2` validity $\leq$ bare at 2+ of 3 N, the pilot effect was the bundled rewording, not the method-line request — reported as such, not reframed." Observed validity: N = 13, 18/20 both arms — **tie**; N = 21, `trace_v2` 18/20 against bare 15/20 — trace higher; N = 31, 17/20 both arms — **tie**. A tie satisfies $\leq$, so under the registered wording the falsifier is **triggered at 2 of 3 N**, and we report it as triggered; the code instead evaluated direction as `t_rate >= b_rate`, strict $<$, under which it fires at 0 of 3 cells, which is what revision 1 reported (Table 5, item 14). That convention appears nowhere in the preregistration, was chosen at analysis time, and read the same tie favorably twice — as "direction held" for P-T1 and as "not a falsifier cell" here.

We name the failure mode **operationalization drift between preregistration text and analysis code**: registered prose leaves boundary behavior — ties, inclusive versus exclusive bounds, rounding — implicit, and the scorer must choose after the data exist; the preregistration literature (§7) discusses what is locked, not this tie-handling layer. Remedies, free before sampling and impossible afterward: register the comparison as executable code, or state the tie convention in the registration text.

`arm_t_analysis.py` now prints both readings and names the registered tie-inclusive one authoritative. Under that reading **the pilot's validity effect is attributed to the bundled rewording, not to the method-line request**, as the falsifier clause instructed; the §6.2 parse-versus-geometric split points the same way, and P-T1 fails under both readings regardless.

6.4 LIMITS OF P-T3: WHAT THE CONCENTRATION RESULT WILL AND WILL NOT CARRY

P-T3 met its registered directional criterion. It is not a flagship result.

No registered alpha. P-T3 registered direction only; the p-value is a post-hoc addition.

Multiplicity. Of four registered predictions, three are reported with p-values. Holm sets the first threshold at $0.05/3 = 0.0167$; $p = 0.0325$ fails it, as it fails Bonferroni at any $m \geq 2$ and Benjamini–Hochberg FDR at $q = 0.05$. Two-sided Fisher is $p = 0.0537$. P-T3 is nominally significant uncorrected and is not under family-wise error control.

Concentration in one cell. Per-cell tests give $p = 0.030, 0.41, 0.50$ (Table 3); the pooled 17-point gap is a 33-point gap at N = 13 averaged with 9- and 6-point gaps elsewhere, driven by a *low bare rate at N = 13* (10/18 = 56%) rather than a high trace rate — the trace rate there (89%) equals that at N = 21. Dropping N = 13, the comparison is 30/35 vs 25/32, $p = 0.31$.

Wave confound. The bare arm mixes two collection waves; trace_v2 is one. Splitting bare on the registered id boundary, with no intervention applied to either half, on-prediction is: N = 13, 3/4 old wave vs 7/14 new; N = 21, 4/7 vs 8/8; N = 31, 5/5 vs 8/12 — the control arm's own between-wave drift is comparable to the effect attributed to the manipulation. (An earlier draft printed 2/4 vs 10/11 at N = 21 and a 25/37 pooled figure — arithmetically impossible against §6.1's id boundary; corrected here from the ledger.) Restricting to one wave, new-wave bare is 23/34 (68%) against trace_v2's 46/53 (87%), but **this stratified comparison is post hoc, not preregistered, no confirmatory weight** — a diagnosis, not a rescue: wave is a confound *in the registered analysis*, which cannot separate it from the manipulation. The design fix, interleaved collection in one session, was not run.

What we therefore claim. A minimal method-naming request does not make the proposer better at geometry and does not measurably change which rare constructions it reaches; in the registered direction it concentrates output onto the anchor, but that concentration fails correction, is carried by one cell, is confounded with collection wave, and is bundled with an output-format rewording that demonstrably affects parse compliance. Studies collecting process descriptors *by asking for them*, including descriptor-driven quality-diversity pipelines where the descriptor is a requested self-report, should treat this as a reason to check, not a coefficient to apply. The pilot also showed a *third* effect in the P-T3 direction at the same cell (trace 9/10 on-prediction against bare N = 21 at 4/7, one-sided $p = 0.16$), so P-T3 replicated a pilot signal rather than testing a blind prediction — a different status, and the prereg's pilot disclosure covered the prompt drift only (Table 5, item 15).

6.5 FAITHFULNESS WITH GROUND TRUTH

Claim and artifact sit in the same completion and the artifact is fully determined, so a method line is checkable against the emitted coordinates. `arm_f_repro.trace_faithfulness()` matches numeric dimensions from the method line against the emitted layout signature, returning 38 of 41 scoreable claims matching (93%) against the registered 90% threshold. Its 12 exclusions are regex coverage gaps, not claims without numeric content (Table 5, item 16): `DIMS_RE` matched only the literal `NxN` form and `ROWSCOLS_RE` required rows before columns, so these fell through despite carrying explicit checkable dimensions — “Regular 3 by 4 grid with one additional circle in row 4”; “3 complete rows of 4 circles and 1 centered circle in the top row”; “rectangular grid 4-4-4-1 with uniform radius 1/8”; “Four-row grid with 4 circles in the first row and 3 in each of the remaining”; “Rectangular grid arrangement with 5 columns and 4 rows, plus one additional circle on top edge”; “Square grid 5 columns 4 rows with 1 additional circle at top”; “Rectangular grid six columns by six rows radius one-twelfth with five circles removed”; “Regular 5-column by 6-row grid plus one additional circle in the right margin”; “Rectangular grid packing with five columns and six rows plus one additional circle”. Three further exclusions — “Uniform horizontal strips with tight row spacing”, “Regular rectangular grid packing with corner circle”, “Equal-radius rectangular grid with 4 rows stacked vertically” — are closer to genuinely unscorable. The excluded set is enriched for non-canonical phrasings, where a mismatch is a priori *more* likely, so the exclusion was not conservative in a known direction; worst case, had all 12 been unfaithful, the rate would have been $38/53 = 72\%$.

Blind rescore under a frozen rubric. The rubric (`p7_faithfulness_rubric.md`) — checkable assertions in any surface form, plus match rules including a base-grid convention for “grid plus additions” claims and a truncation convention for “with k removed” claims — was committed to git as `e181d2a` at 03:56:56 on 2026-08-01, **before** any rescoring run, so the freeze is checkable, and it pre-committed the reporting rule: below 90%, P-T4 reported not confirmed, full stop, no consolation framing. Scoring was blind — claim text and circles only, no predicted or rival values, no running tally, no threshold and no sight of the existing 38/41 result — with tallying afterward; all 60 labels are released verbatim in `p7_blind_labels.json`.

Result. Over all 60 trace_v2 rows: **54 MATCH, 2 MISMATCH, 4 UNSCOREABLE** — a corrected rate of $54/56 = 96.4\%$ of scoreable claims, Wilson 95% CI [88%, 99%], passed, and on 56 scoreable claims rather than 41. The CI's lower bound sits just below the registered threshold, so “clears 90%” is a statement about the point estimate at this n. The two mismatches share a form: row 11 claims “alternating rows of 4 and 3 circles” over observed row sizes 4,3,3,3, row 33 claims “5 rows alternating between 5 and 4 circles” over observed 5,4,4,4,4 — both *mis-descriptions of alternation*, a repeating pattern named and built for the first period only. The original scorer's three mismatches were three different claim forms, not one (Table 5, item 17); the filler-adds-rows mechanism revision 1 described was real, and the rubric's base-grid convention handles it, so all three score MATCH under hand adjudication.

Scope, three ways. The blind pass covers *invalid* rows too: the original 93% conditioned on completions that produced a correct packing — the wrong conditioning, since a method line attached to an overlapping layout is where description and artifact are likeliest to diverge — while the blind pass scored all 60 rows including the 7 invalid ones, 4 of the 60 unscorable. Second, the check verifies the description against the *artifact*, not the process: the non-claim guard applies in full, and for a model already emitting grids, “5×5 grid” matching a 5×5 grid is closer to a self-consistency check than to the causal question the chain-of-thought literature asks (§7). Third, the audit is confounded with the elicitation arm, computed on the arm §6.2 shows concentrates 87% of

valid outputs onto one template, where correct description is close to free; faithfulness is informative precisely on off-template outputs, such as the pilot’s triangular-hexagonal $6+5+4+3+2+1$ sample, described exactly while summing to 1.75. The split by on-versus off-prediction is not run here (§9, stopping rule). What survives is narrow: in a domain where claims are checkable, the method lines this proposer emits are, at 54 of 56 scoreable claims, true of the object it produced — and both failures are a model describing a pattern it started and did not finish.

7. Related work

Citation errors in revision 1 are corrected here and itemized in Table 5 (items 18–26).

LLM-driven discovery and the circle-packing scoreboard. Language Model Crossover (Meyerson et al., arXiv 2023, ACM TELO 2024), ELM (2206.08896), EvoPrompt and LLaMEA (2405.20132) established the LLM call as an EC operator over prompts and programs; FunSearch turned the pattern into a discovery claim, and every system in §1’s scoreboard inherits it. Two orderings in that scoreboard matter: HELIX’s figure is marginally *below* ShinkaEvolve’s while being described as state of the art, and ThetaEvolve claims new best-known bounds, so the benchmark is saturated only across the AlphaEvolve–ShinkaEvolve–HELIX cluster and the number is still moving. OpenEvolve, the open-source AlphaEvolve reproduction most practitioners run, reports ≈ 2.635977 . §1.1 records the correction to our record comparison and its $N > 30$ coverage gap. Two critiques bear on our framing and are frequently conflated, so we attribute them explicitly: Gideoni, Risi and Gal (2602.16805) show simple baselines recover much of the reported advantage, and Berthold et al. (2605.04850) show classical solvers do too.

Template convergence and diversity collapse. “Mutation Without Variation” (2606.05408, Gurkan, Stonedahl and Wilensky) finds iterated LLM program mutation collapsing onto previously seen structural templates — in 87% of chains, over 93% of mutations revisit a previously seen form — which is the same bias family in a different regime: mutation loops, program space, no closed form, no preregistration. On anchoring, 2505.15392 treats anchoring as a general bias over initial information while 2412.06593 and 2410.15413 document numeric primes dragging point estimates; our modality is the difference, since a construction template has a computable value where a scalar anchor has only a measurable pull. Concurrently, “Measuring the Gap Between Human and LLM Research Ideas” (2607.01233) reports the same collapse one level up, in idea space: across 11,683 paired proposals, LLM research ideas concentrate on bridge-and-synthesis framings (47–64% versus 12% for humans; normalized entropy down to 0.55 versus > 0.92), and chain-of-thought pushes the distribution *further* from the human one. That result is distributional; ours is the limiting case where the collapsed distribution’s mode is a single closed-form-predictable object — narrowness measurable there, computable here. “Artificial Hivemind” (2510.22954) documents the same homogeneity across open-ended domains at survey scale, which makes cross-domain generality of the collapse the expected default; against that backdrop our contribution is not that collapse occurs but that, in this regime, its mode is a single closed-form-computable object, stated before sampling. Converging skepticism about what the proposer contributes comes from “Dictionaries, Not Darwin” (2607.04108), 2606.10587, BehaveSim (2603.02787), Strategy Diversity (2605.09292), the bin-packing critiques (2510.27353; 2501.11411) and MathConstruct (2502.10197). Two results cut the other way and belong on the other side of the ledger: 2407.10873 (Zhang et al.) provides empirical grounding for the *importance* of evolutionary search, which is evidence against our substitution of an unconditioned call for the loop and is engaged in §8; and 2604.19440 finds strong LLM optimizers act as *local refiners*, which is what a parent-conditioned call would be, so it bears on §8’s mutation arm.

Scaling inversions. Zhou et al. (*Nature*, 2024) show larger and more instructable models becoming less reliable; the o3/o4-mini system card shows the same on PersonQA; 2506.06941 reports collapse past a complexity threshold; GeoBuildBench (2605.13167) finds geometric construction a regime where nominal capability and executed correctness diverge. Our instantiation is *attractor family* versus executability.

Trace faithfulness, and what our result is not. One cluster intervenes *on trace content* and finds answers do not move — Reasoning Theater (2603.05488), Project Ariadne (2601.02314), “Beyond the Commitment Boundary” (2606.13603) and “The Chain Holds, the Answer Folds” (2605.29087) all report causal decoupling. A second cluster *estimates* faithfulness without observing the process, represented here by 2503.08679; §6.5 reports the case where ground truth for the described artifact exists. Our result sits at a third intervention point: we vary whether a trace is *requested* rather than editing its content, which is measurement reactivity, not causal faithfulness. The closest cousin is “The Price of Format” (2505.18949), showing format constraints collapse generation diversity — and as §6.1 concedes, our manipulation includes a small output-format change, so we commit a miniature version of the confound we cite. Format-restriction performance costs are documented at scale in 2408.02442. Related: 2505.14617, 2510.01171, 2506.17630.

Preregistration lineage. 2606.27687, 2607.07184, 2606.11217 and 2607.00276 preregister recipes, outcome-blinded predictions and agent protocols; HindsightBench (2607.18867) releases frozen preregistrations of directional aggregate hypotheses. We adopt the standard rather than claim it. What we add is the object locked — exact closed-form point predictions of specific output values — and, from §6.3, the observation that locked prose and executing scorer can disagree at the boundary. Structural precedents: 2001.08361 and 2203.15556 publish a functional form then run the confirming instance; 2411.16035 is the closest structural twin. QDAIF (2310.13032) is the QD-descriptor setting most affected by our §6 result, and GigaEvo (2511.17592) the concrete in-scope system running MAP-Elites with LLM-driven mutation, where a self-reported behavioral descriptor would be perturbed in practice.

8. Limitations and untested alternatives

The measured distribution is not the in-loop distribution. We sample generation-0, parent-free, code-free, feedback-free calls; the systems in §7 sample parent-conditioned program proposals under selection. The claim that survives is about unconditioned calls, and 2407.10873 is published evidence that evolutionary search contributes materially, cutting against any stronger reading. The decisive cheap experiment is a one-parent mutation arm: same cells, prompt = bare template plus one parent packing plus its score plus “propose a modification that increases the sum of radii”, under three parent conditions (the on-prediction truncated grid, the higher-value family rival, an off-family parent), measuring whether the branch rule still predicts and whether the model inherits the parent’s k . That experiment has now been run as the preregistered arm MU (§3.5), and its registered falsifier triggered: 0 of 26 valid rival-parent samples return to the anchor, keep-or-improve and k -inheritance are both 26/26, and the anchoring result is scoped to unconditioned calls as a stated negative transfer result. The concern this paragraph raised in revision 2 was correct.

A mechanism-free alternative, now directly probed. Under an explicit “construct the packing by reasoning alone” constraint, a $k \times k$ grid at $r = 1/(2k)$ is close to the only construction whose coordinates can be emitted from mental arithmetic without error accumulation, and $\text{round}(\sqrt{N})$ is the arithmetically nearest such grid. That hypothesis predicts every observation in §2–§3 without appeal to stored templates, and predicts §5’s inversion too: the more ambitious tier attempts constructions whose coordinates it cannot compute by hand and overlaps — precisely the observed failure mode (24 overlaps and 2 nonpositive radii out of 26 `opus_alias` failures). The separating experiment — hand the model the recipe family explicitly, state $V(k, m)$ and the admissible k , and ask it to *choose* k ; if it picks the argmax, the tractability reading stands — has now been run as the preregistered arm CH (§3.5). The outcome splits the hypothesis in two: the model’s *choice* follows the stated score table (so tractability does not explain the selection), but its *execution* fails off-template in 30 of 31 invalid attempts (so tractability does explain the instantiation — the grid is what survives mental arithmetic). §1’s vocabulary remains behavioral and no memorization claim is made.

Contamination is not probed. $N = 26$ is the canonical published cell and plausibly in training data. We ran no canary-string test, no comparison against N values absent from the literature, and no perturbed-container test preserving difficulty while destroying lexical overlap. The rectangle arm is the nearest thing we have, at $n = 11$ valid.

Single vendor. All tiers in the main study come from one provider. Two preregistered cross-vendor extensions are in progress at submission time, both registered before sampling and disclosed regardless of outcome. Arm GM (gemini-2.5-flash-lite, direct vendor API, prereg commit `37b3adb`) is quota-throttled and incomplete; no collected output has been analyzed. Arm GM2 (gemma-4-26b-a4b-it, prereg `3019aab`) completed 140/140 invocations with a null-compliance outcome: 0 of 140 responses were parseable under the registered pipeline — the model consumed its entire 4,096-token output budget on visible deliberation and never emitted the required coordinate list, under a prompt with which the weak-tier arm F model complied at 100% parse rate. Per the registered rule, all cells are UNSCOREABLE and no confirmatory claim is made in either direction; a follow-up with an enlarged output budget (arm GM3, registered before sampling) is running. The anchoring law therefore remains single-vendor evidence at this time — but the format-compliance cliff between families under an identical prompt is itself a preregistered observation, and it is consistent with §6.1’s point that output-format constraints are load-bearing. We flag one reading the design cannot yet exclude: a 0/140 parse rate at a 4,096-token cap is also the signature of budget truncation — a verbose model cut off before its answer — rather than format inability. Arm GM3’s enlarged budget separates the two, and we defer any characterization of the cliff’s cause until it reports.

Sampling parameters unpinned. The runtime does not expose pinned decoding parameters, so every effect is a distributional claim over the observed sampling regime, and §3.2’s mode-ceiling result inherits that scope. This is the subject of the companion paper.

`opus_alias` **provenance.** The label denotes the serving alias addressed, not an identified model version. Without pinned weights we cannot separate a model-tier effect from a serving-path effect, and the supporting anomalies live only in the runtime transcript.

Confirmatory base for the branch rule, updated by arm M. floor and round disagree exactly on $N \in [k^2-k+1, k^2-1]$, so every discriminating cell is a trap cell by construction. The truncate arm now has five confirmed k in the square (4, 5, 6, 7, 8 — the last added by arm M's $N = 57$ cell, §3.4) plus two rectangle cells; the extend arm is confirmed only at $m \leq 1$ and **disconfirmed at $m \geq 4$** (§3.4), which is the sharper scope the abstract now carries. The empirical- k back-out named as future work in earlier revisions has since been run twice as disclosed post-hoc diagnostics: `diagnostics_kmatch.py` (§3.2 — 50/50 on-prediction samples k^* -structured, 64/69 valid overall) and arm M's per-cell k_{emp} distributions (§3.4 — the converge cells sit on k^*+1 , which is how the filler-branch failure was diagnosed).

9. Reproducibility, deviations, and disclosures

Deviations from preregistration are tabled in Appendix B; every claim that changed relative to revision 1 is tabled in Appendix C. Provenance, in list form:

- **Registration.** Prompts were hashed with SHA-256 and the hashes recorded before sampling, with the predicted values they were to be tested against: `arm_f_repro.py` (header, P1–P5, square arm), `arm_s_preregistration.txt`, `arm_o_preregistration.txt`, `arm_t_preregistration.txt` (prereg hash `ab7900a8...`), `arm_m_preregistration.txt` with `arm_m_prompts.json` (committed `e528c6b` before sampling), and `arm_mu_preregistration.txt` with `arm_mu_prompts.json` (committed `2b7d202` before sampling; decision thresholds, power analysis, tie conventions and falsifier F-MU1 all registered). The rectangle transfer was registered before collection and never refitted.
- **Four limits on that claim.** (i) The digests cover the task prompt only; the subagent inherits instruction files outside it, so what is locked is a *fragment* of the conditioning context, a replicator cannot reconstruct those files, and we cannot establish they were identical across the arm-F and arm-T windows — the boundary §6.4 identifies as a confound. (ii) `arm_f_prompts.json` covers $N = 13, 17, 31, 35$ and 37 ; the $N = 21$ bare hash is registered in `arm_t_preregistration.txt`; the **$N = 43$ bare hash appears in no preregistration and no prompts file**, existing only in the ledger — it recomputes exactly from the bare template, but recomputation after the fact is not registration. (iii) **No hash was ever registered for the pilot prompt**, whose drift the preregistration describes only in prose; the ten pilot rows carry `prompt_sha256: null` in the corrected ledger with an explanatory note. (iv) The rectangle ledger (`arm_g_candidates.jsonl`) carries no prompt hashes, no run dates and no proposer fields at all.
- **Ledger metadata correction.** Three referee findings established that the shipped ledger's provenance fields contradicted the paper: trace rows carried the *bare* prompt hash for their cell, `run_date` was uniformly `2026-07-30` on all 215 rows, and the Sonnet and `opus_alias` rows were stamped with the Haiku alias and dated id. All three are corrected in `arm_f_candidates_v2.jsonl`; the correction is metadata only — every scored quantity, geometry, raw output, arm label and sample id is byte-identical to v1, verified field by field, each corrected field carrying a sibling `*_v1` field. `arm_f_candidates.jsonl` is unmodified (sha256 `02ecabcd...`); v2 is `9f845f78...`. Correction rules, sources, wave-boundary arithmetic and caveats — including that `run_date` has day granularity only — are in `ledger_v2_corrections.md`. Analysis scripts read v2.
- **Excluded and unreleased records.** The five concurrency-cap rejections appear in the ledger in no form, no `excluded_reason` field exists, and the exclusion rule was not preregistered; both rates are reported in §3.3. The `opus_alias` latency and token-count anomalies (§5) come from the runtime session transcript; no latency or token field exists in any released artifact.
- **Analysis artifacts.** Raw outputs are stored verbatim (`arm_f_raw.json`, `arm_f_candidates.jsonl`, `arm_f_candidates_v2.jsonl`, `arm_g_candidates.jsonl`, `arm_m_collect.jsonl` scored by `arm_m_analysis.py`, `arm_mu_collect.jsonl` scored by `arm_mu_analysis.py` with frozen output `arm_mu_results.txt`) and scoring is deterministic and local. Value claims are gated on an independent LP oracle rather than on the closed form being tested: `n_sweep_forecast.py` with `verify_against_lp` checks 83 configurations to within 10^{-9} , and the same oracle produced §4.2's negative result. The faithfulness rubric (`p7_faithfulness_rubric.md`) was frozen by git commit `e181d2a` (2026-08-01 03:56:56) before any rescoreing, and all 60 blind labels with justifications are in `p7_blind_labels.json`. `arm_t_analysis.py`, `n_sweep_forecast.json`, `rect_forecast.json`, `p11_mode_baseline.json` and `STATE.md` are included; every table regenerates from raw outputs by a command documented in `HOW_TO_RUN.md`.

- **Use of AI systems.** This paper studies language-model behaviour and was written with language-model assistance; both facts are stated here so a reader can weigh them together. Claude models (principally the Opus, Fable and Sonnet tiers) wrote the collection harnesses, the scoring, analysis and verification scripts, and the structural-check tooling, and drafted the manuscript prose; the external referee reports came from `deepseek-v4-pro`, `deepseek-v4-flash` and Gemini under written protocols (next bullet). The human author directed the programme, set the research questions, approved each preregistration before sampling, made the final inclusion, stopping and submission decisions, and is solely responsible for the content; no language model is an author. The claim a reader can check is the ordering, not the authorship: each preregistration commit is a git ancestor of the sampling it governs. The drafting of those preregistration texts, like the rest of this repository, was model-assisted. Two consequences are load-bearing rather than cosmetic. First, the models that wrote the analysis are of the same family as arms under study (§5), so the blind-spot risk flagged below for the referee panel applies to the analysis code as well — the LP oracle, the frozen scorer output and the referee panel's independent recomputations exist in part to bound it. Second, every number here was produced by a script a model wrote, which is why the scripts, the raw ledgers and the frozen outputs are all released: what a reader is asked to trust is the artifact, not the author.
- **Review process.** The three referee reports that drove this revision (`p4_review_stats.md`, `p4_review_reviewer2.md`, `p4_review_gecco.md`), the arbitration of their conflicting recommendations (`p6_cruxes.md`), and the blind faithfulness adjudication of §6.5 were produced by language models under written protocols, not by human referees. They independently recomputed the paper's statistics from the released ledger and found eight recomputation mismatches, all itemized in Appendix C, and their arithmetic was checkable and was checked. But the panel is not peer review, its verdicts carry no external authority, and it shares a family resemblance with the system under study — a blind-spot risk disclosed here rather than left implicit.
- **Stopping rule.** This revision closed after implementing the three named referee reports and the arbitration verdict; no further data-dependent analysis was performed. Analyses named but not run — the one-parent mutation arm, the recipe-family choice probe, $N = 57$, the empirical-k back-out, faithfulness split by on- versus off-prediction, a code-enabled probe, interleaved re-collection of the arm-T pair, and P5 at $n = 20$ — are recorded in §6.5 and §8 as future work, so that the reported statistics remain the ones the preregistrations governed.

Appendix A — Prompts, verbatim

A.1 Bare template (`{n}` parameterized; the $N = 13$ instance hashes to

`32db485bea625ff9f39f4723ebf1a01f337559a9e2cf567fb486928f71f7f8df`):

```
Pack exactly {n} non-overlapping circles inside the unit square [0,1]x[0,1] so that the sum of their radii is as large as possible.
MUST hold: exactly {n} circles; every circle fully inside the unit square [0,1]x[0,1] ( $x-r \geq 0$ ,  $x+r \leq 1$ , same for  $y$ ); no two circles overlap (distance between centers  $\geq$  sum of the two radii; touching is allowed).
Do not write or execute code - construct the packing by reasoning alone.
Output ONLY the raw Python list of {n}  $[x, y, r]$  lists. No explanation, no code fences, no other text.
```

A.2 `trace_v2` template — identical to A.1 through line 3, then:

```
First write one line beginning "METHOD:" naming the construction you used. Do not use square brackets anywhere in that line.
After the METHOD line, output ONLY the raw Python list of {n}  $[x, y, r]$  lists. No explanation, no code fences, no other text after the list.
```

A.3 Dispatch wrapper (both arms; *not* part of the hashed task prompt, hence an instance of the fragment-hashing limitation in §3.1): `Do not use any tools. Your entire final message must be the answer and nothing else.`

A.4 Registered prompt hashes. bare $N=13$ `32db485b...`, $N=17$ `8437df75...`, $N=21$ `a415425b...`, $N=31$ `a664d003...`, $N=35$ `3b08c56e...`, $N=37$ `2c0a8819...`; `trace_v2` $N=13$ `a920f1c9...`, $N=21$ `91205727...`, $N=31$ `bd490b7b...`. The $N=43$ bare hash `1208e7d2...` is ledger-only, and the pilot prompt was never hashed (§9).

A.5 Parser. Completions are parsed with `ast.literal_eval` after stripping code fences; a completion failing to yield a list of N numeric triples is recorded `literal_eval_failed` and scored invalid, with the failure mode logged (§6.2 reports parse and geometric failures separately). Radii must be strictly positive; containment and pairwise non-overlap are checked at the two registered tolerances.

Appendix B — Deviations and disclosure table

Table 4 — Deviations from preregistration.

Prediction	Registered rule	As registered?	Outcome / reason
P1–P3	exact-value point predictions	yes	§3.3, Table 1
P4 (N=37)	converge, 3.0345178	yes	confirmed
P5 (N=35)	truncate, 2.9166667, structural separation	yes	3/4 valid; n = 4, “consistent with”
P-S1...P-S4	directional tier comparisons	yes	§5
P-O1, P-O2, P-O4	tier comparisons on on-prediction / rival rates	no	declared not evaluable post hoc after a 13% validity collapse; the registered disconfirmation (regression toward the trap) did not occur
P-O3	multi-radius fraction ≥ 0.9 of valid	yes	trivially satisfied on 4/4 valid
P-T1	direction at each N and pooled $p < 0.05$	yes	not confirmed, $p = 0.30$
P-T2	directional, no alpha	yes	confirmed as registered (1/53 vs 2/50); no inferential weight claimed
P-T3	directional, no alpha	yes	met as registered; $p = 0.0325$ uncorrected fails Holm; carried by N = 13 (§6.4)
P-T4	$\geq 90\%$ of scoreable claims match	yes, plus corrected scorer	38/41 (93%) original; 54/56 (96.4%) blind, frozen rubric
P-M1–P-M3	modal output = $V(k^*, m)$ at N = 20/30/41	yes	disconfirmed at all three cells — modal outputs are (k^*+1) -grid values (§3.4)
P-M4	N = 57 modal = T(8,57), rival count 0	yes	confirmed (6/10 on-prediction, 0/10 rival)
F-M1 (arm M falsifier, filler branch)	modal $\neq V(k^*, m)$ at 2+ of 3 converge cells, ties count as NOT	yes, tie-stated in registration	triggered at 3 of 3 ; $V(k, m)$ restricted to $m \leq 1$ (§3.4)
F-M2 (arm M falsifier, truncate branch)	N = 57 modal $\neq T(8,57)$	yes	not triggered
Arm M rejection rule	runtime deaths excluded, counted	yes (registered for this arm)	3 N = 41 invocations excluded, cell reported as 12 of 15
Falsifier	trace_v2 validity ☐ $\leq$ bare at 2+ of 3 N	both readings reported	triggered under the registered tie-inclusive reading (2/3 cells); not triggered under the strict ☐ $<$ the code implemented (§6.3)
Ladder inclusion	—	post hoc	<code>opus_alias</code> initially excluded, later included after results known; ladder reported both ways (§5)
Cap exclusion	—	not registered	five records excluded; both rates reported (§3.3)

Appendix C — Corrections ledger

Table 5 — Corrections to revision 1. Every claim that changed between the reviewed draft and this one, with the finding that forced it. stats = p4_review_stats.md , R2 = p4_review_reviewer2.md , GECCO = p4_review_gecco.md , council = p6_cruxes.md .

#	§	Revision-1 claim	Corrected	Source
1	Abs, 1	"recall" a template; "template memorizer"	behavioral vocabulary only: emits / concentrates / modal output	R2 #2; R-E
2	2.3	round($\sqrt{N}$) an empirical discovery; "zero free parameters"	a definition — two formalizations of "nearest" coincide on the integers; weakly identified selection among four	GECCO M8; R2 #24
3	1.1, F1	penalty "hits exactly zero at the top of each zone"	zero at N = 14, 15, 35, 48; 0.59% residual at N = 24 ; 4 of 22 trap N cost nothing, so branch $\neq$ penalty	GECCO M14; R2 #19
4	1.1, 7	"published bounds sit above all of them"	withdrawn: the N = 26 bound (2.63598) <i>is</i> ShinkaEvolve's, so the LLM systems are the record; table stops at N = 30, leaving the LP abort gate unchecked above it	GECCO M3; R2 #20–21
5	3.3	cap exclusions "understated validity by 17%"	77.8% vs 70.0% — 7.8 points, 10% relative; exclusion unregistered, both rates reported	R2 #13; GECCO m15
6	3.3, 6	"two parse failures"; "two fraction-literal samples"	three in each case	stats F14, F15
7	Abs, 1, 3.2	"predicts the exact sum-of-radii the model will emit"	predicts the empirical modal output at 7/7 N; per-sample rate = modal frequency, 56–86% by cell, 46% pooled; round-number baseline 2/69	R2 #15; R-E
8	4.1	one a = 3 sample left the family; rectangle "confirmed"	two (3.45, 3.5); all 11 valid samples characterized; partial support, CI [21%, 72%], null stated	stats F16; R2 #17; GECCO M9c
9	Abs, 1, 5	32/45 = 71%; inversion 71 $\rightarrow$ 100 $\rightarrow$ 13	35/45 = 77.8% at 10 ⁻⁶ , 64.4% at 10 ⁻⁸ ; inversion 78 $\rightarrow$ 100 $\rightarrow$ 13 (64 $\rightarrow$ 90 $\rightarrow$ 13 strict); matched-cell 83 $\rightarrow$ 100 $\rightarrow$ 13 given separately	stats F1; R2 #3; R-F
10	5	Sonnet rival 6/30, 3/10 at N = 31	5/30 and 2/10 — the 2.75 escape sits 1.47 $\times$ 10 ⁻³ from the rival, inside the window; categorical claims now structural or at 10 ⁻⁶	stats F12; R2 #14
11	1	"a monotone inversion"	monotone in ambition only; validity rises then collapses	GECCO m21
12	6.2	P-T2 "not confirmed" (p = 0.48); P-T3 "confirmed" (p = 0.0325)	one rule applied symmetrically: P-T2 confirmed as registered (1/53 vs 2/50, no inferential weight); P-T3 met as registered but fragile	stats F7; R-C
13	Abs, 1, 6.2	"leaving validity unchanged"	no <i>detectable</i> change (p = 0.30, n = 60/arm, underpowered); the P-T1 direction is entirely parse compliance — 3 vs 0 parse failures, 7 vs 7 geometric	R2 #7–8
14	6.3	"the registered falsifier was not triggered"	triggered at 2 of 3 N under the registered $\leq$; not under the code's strict $<$; both reported, registered reading authoritative; pilot validity effect attributed to the bundled rewording	stats F8; R-B
15	6.4	pilot produced two effects, both died at scale	a third ran in the P-T3 direction at the same cell (9/10 vs 4/7, p = 0.16): P-T3 replicated a pilot signal	stats F18
16	6.5	12 claims excluded as "no numeric content"	regex coverage gaps; 9 of 12 carry checkable dimensions, quoted verbatim in §6.5; worst case 38/53 = 72%	stats F9
17	6.5	the three mismatches "are all the same case"	one of three; all three score MATCH under blind adjudication; the two remaining mismatches are alternation mis-descriptions (rows 11, 33)	stats F10; R-D

#	§	Revision-1 claim	Corrected	Source
18	1, 7	SeaEvo (2604.24372) reports a ~2.636 packing value	removed: SeaEvo does not evaluate circle packing	GECCO M4
19	7	2605.29268 reports asymmetric proposal mass in program space	deleted — that paper is on bandit compute allocation; removing it leaves no cited evidence that the anchoring survives the code channel	GECCO M5
20	7	2606.13603, 2605.29087 estimate faithfulness by perturbation / probing	both are causal-decoupling results; only 2503.08679 is an estimator	GECCO M6
21	7, 8	2407.10873 and 2604.19440 filed as loop skepticism	2407.10873 grounds the <i>importance</i> of evolutionary search — evidence against our substitution, engaged in §8; 2604.19440 recites as “local refiner”	GECCO M7, m5
22	1, 7	AlphaEvolve “2.635”; FunSearch <i>Nature</i> 2023; LMX TELO 2023; AlphaEvolve/ShinkaEvolve unidentified	2.63586276; <i>Nature</i> 625(7995):468–475, 2024; arXiv 2023 / TELO 2024; both systems identified	GECCO M3c, m2, m3
23	1, 7	ThetaEvolve “in the same band”; HELIX ordering unremarked	ThetaEvolve claims new best-known bounds; HELIX sits below ShinkaEvolve while claiming SOTA	GECCO m10, m11
24	7	“EvoDiverse (2606.10587)”; 2505.15392 as numeric anchoring	cited by title, method name unverified; 2505.15392 is general anchoring	GECCO m4, m9
25	7	HindsightBench freezes “under SHA-256”; 2607.07184 “files OSF preregistrations”	both unsupported by the sources; restated as frozen preregistrations and registered outcome-blinded predictions	GECCO m7, m8
26	7	“Mutation Without Variation”: 87% of chains and 93% of mutations	nested, not parallel: in 87% of chains, over 93% of mutations revisit a prior form	GECCO m1
27	Abs, 6.1	“hundreds of invocations”; corpus arithmetic unreconcilable	231 invocations, itemized in §6.1	R2 #30; GECCO M11c
28	3.3, 6	18/23 and 35/50 given without stating they are nested subsets	subsets defined in §3.3 and Table 1; full-ledger 41/57 and 2/57 alongside	stats F13; R2 #4; GECCO M11b
29	5	Sonnet multi-radius “against a Haiku baseline of 13/35”	dropped; that baseline reproduces at no denominator	R2 #23a
30	5	“no Haiku sample did [this] in 101 invocations”	155 weak-tier rows; two exceeded the family best, both by $\sim 10^{-7}$	stats F17; R2 #23b
31	3.1, 9	dual-tolerance reporting promised, never delivered; “cryptographic prompt hashes”	both tolerances in every validity figure; hashes restated as prompt-fragment hashes with coverage gaps itemized above	stats F21; GECCO M12
32	all	five HTML working comments and a merge-provenance header shipped	removed; the two load-bearing ones ($k = 8$ clipping, p -value convention) promoted into §1.1, Fig 1 and §6	stats F22; R2 #29; GECCO NIT-1–2